

Running Programs on a System

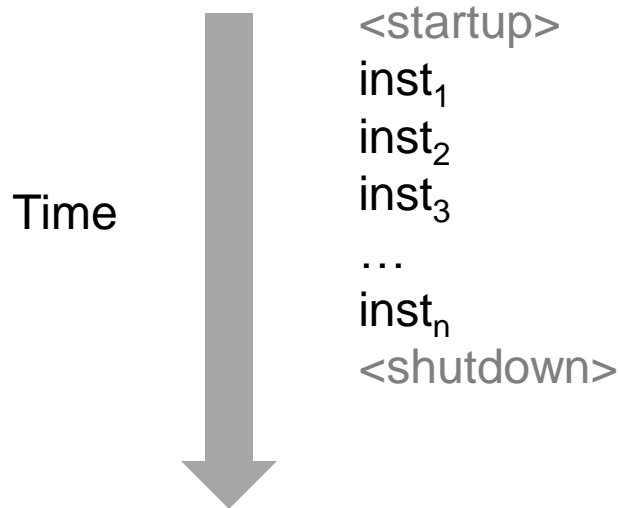
Exceptional Control Flow



Control Flow

- Processors do only one thing:
 - From startup to shutdown, a CPU simply reads and executes (interprets) a sequence of instructions, one at a time
 - This sequence is the CPU's control flow (or flow of control)

Physical control flow



Altering the Control Flow

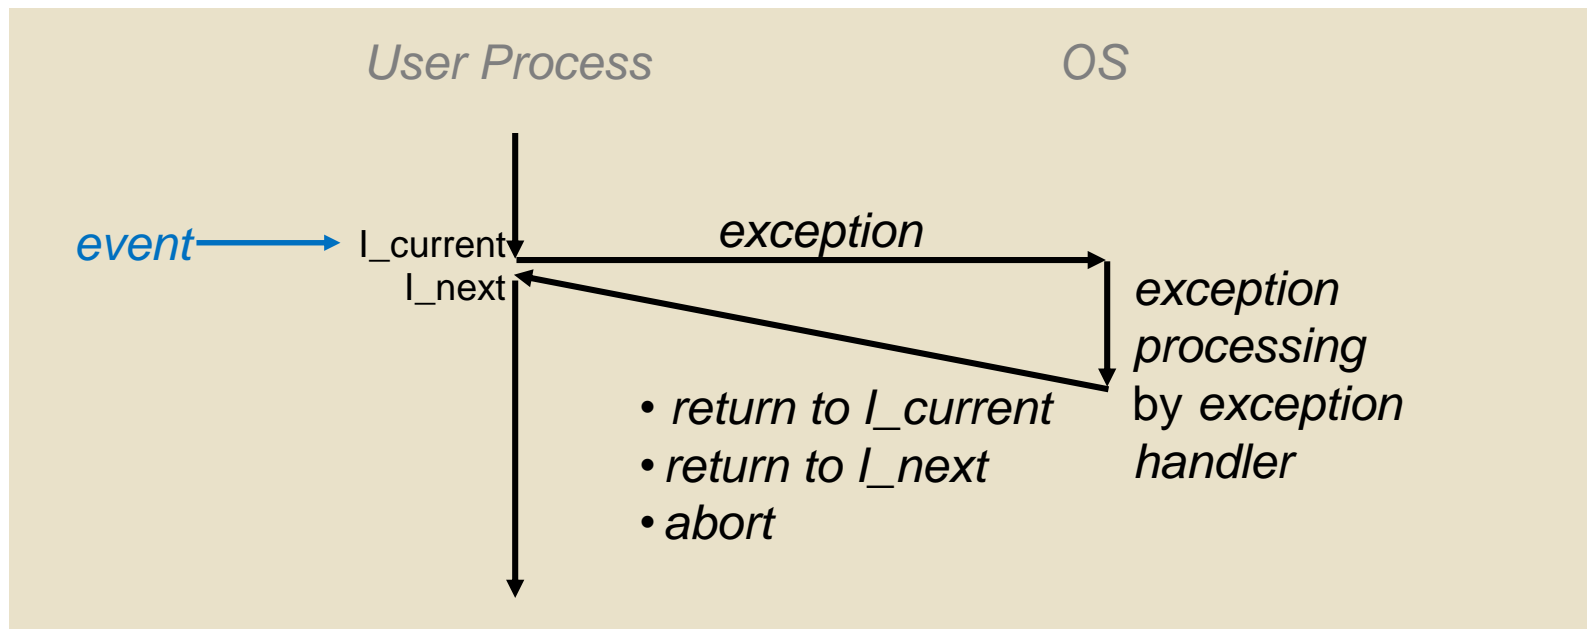
- Up to now: few mechanisms for changing control flow:
 - Jumps and branches, call and return
 - ▶ react to changes in program state
 - Signals
 - ▶ sent by the kernel
- But how exactly do we react to changes in system state?
 - data arrives from a disk or a network adapter
 - instruction divides by zero
 - user hits Ctrl-C at the keyboard
 - system timer expires
- System needs mechanisms for “exceptional control flow”

Exceptional Control Flow

- Exists at all levels of a computer system
- Low level mechanisms
 - Exceptions
 - ▶ change in control flow in response to a system event (i.e., change in system state)
 - Combination of hardware and OS software
- Higher level mechanisms
 - Process context switch
 - Signals
 - Nonlocal jumps: `setjmp()/longjmp()`
 - Implemented by either:
 - ▶ OS software (context switch and signals)
 - ▶ C language runtime library (nonlocal jumps)

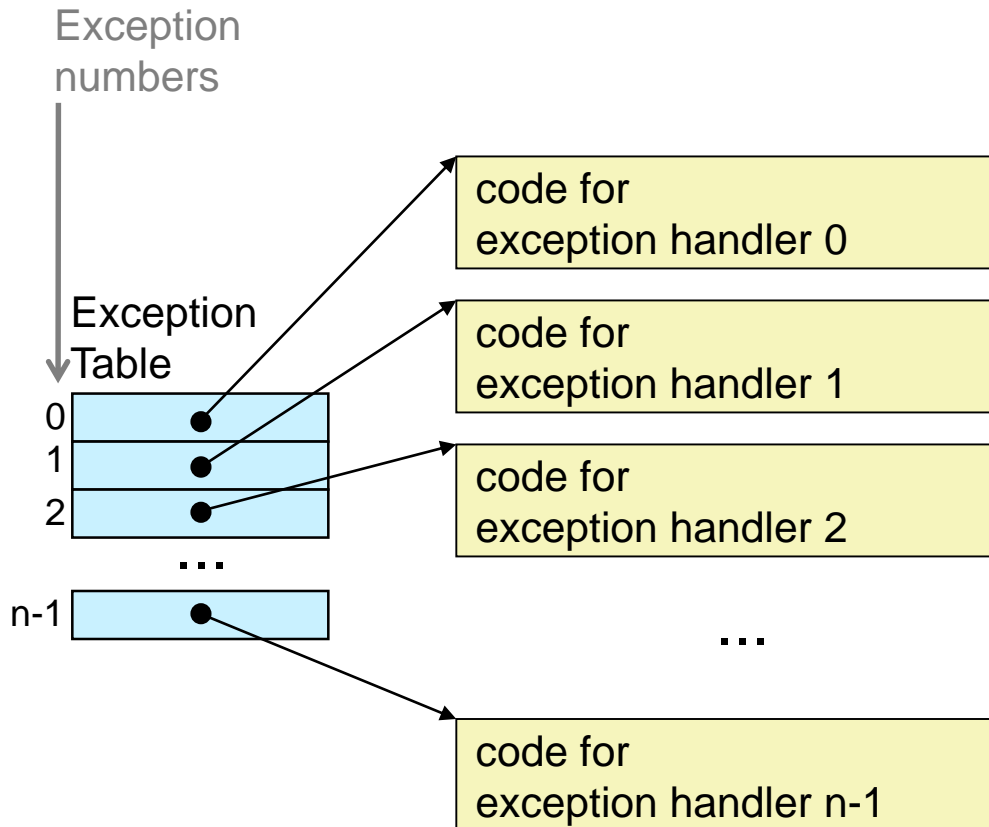
Exceptions

- Low-level exceptions provide the “glue” between hardware events and the higher-level mechanisms
 - transfer of control to the OS in response to some event (i.e., change in processor state)



- Examples:
div by 0, arithmetic overflow, page fault, I/O request completes, Ctrl-C

Interrupt Vectors



- Each type of event has a unique exception number k
- k = index into exception table (a.k.a. interrupt vector)
- Handler k is called each time exception k occurs

Intel x86 Exceptions and Interrupts

Table 6-1. Exceptions and Interrupts

Vector No.	Mnemonic	Description	Source
0	#DE	Divide Error	DIV and IDIV instructions.
1	#DB	Debug	Any code or data reference.
2		NMI Interrupt	Non-maskable external interrupt.
3	#BP	Breakpoint	INT 3 instruction.
4	#OF	Overflow	INTO instruction.
5	#BR	BOUND Range Exceeded	BOUND instruction.
6	#UD	Invalid Opcode (UnDefined Opcode)	UD2 instruction or reserved opcode. ¹
7	#NM	Device Not Available (No Math Coprocessor)	Floating-point or WAIT/FWAIT instruction.
8	#DF	Double Fault	Any instruction that can generate an exception, an NMI, or an INTR.
9	#MF	CoProcessor Segment Overrun (reserved)	Floating-point instruction. ²
10	#TS	Invalid TSS	Task switch or TSS access.
11	#NP	Segment Not Present	Loading segment registers or accessing system segments.
12	#SS	Stack Segment Fault	Stack operations and SS register loads.
13	#GP	General Protection	Any memory reference and other
14	#PF	Page Fault	Any memory reference.
15		Reserved	
16	#MF	Floating-Point Error (Math Fault)	Floating-point or WAIT/FWAIT instruction.
17	#AC	Alignment Check	Any data reference in memory. ³
18	#MC	Machine Check	Error codes (if any) and source are model dependent. ⁴
19	#XM	SIMD Floating-Point Exception	SIMD Floating-Point Instruction ⁵
20-31		Reserved	
32-255		Maskable Interrupts	External interrupt from INTR pin or INT <i>n</i> instruction.

source:

Intel 64 and IA-32 Architectures
Software Developer's Manual
<http://download.intel.com/design/processor/manuals/253665.pdf>

Exceptions and Processor Mode

- User programs run in **user mode**

- restricted access to memory, devices, etc
- to serve system calls and talk to the hardware, the process may need to run in kernel mode
- Problem: how can a user process invoke a kernel function?

→ **Exceptions cause a switch to kernel mode**


Synchronous Exceptions

- Caused by events that occur as a result of executing an instruction:
 - Traps
 - ▶ Intentional
 - ▶ Examples: system calls, breakpoint traps, special instructions
 - ▶ Returns control to “next” instruction
 - Faults
 - ▶ Unintentional but possibly recoverable
 - ▶ Examples: page faults (recoverable), protection faults (unrecoverable), floating point exceptions
 - ▶ Either re-executes faulting (“current”) instruction or aborts
 - Aborts
 - ▶ unintentional and unrecoverable
 - ▶ Examples: parity error, machine check
 - ▶ Aborts current program

Trap Example: Opening File

- User calls: `open(filename, options)`
- Function `open` executes system call instruction `int`

```
0804d070 <__libc_open>:  
. . .  
804d082:      cd 80                int     $0x80  
804d084:      5b                    pop     %ebx  
. . .
```



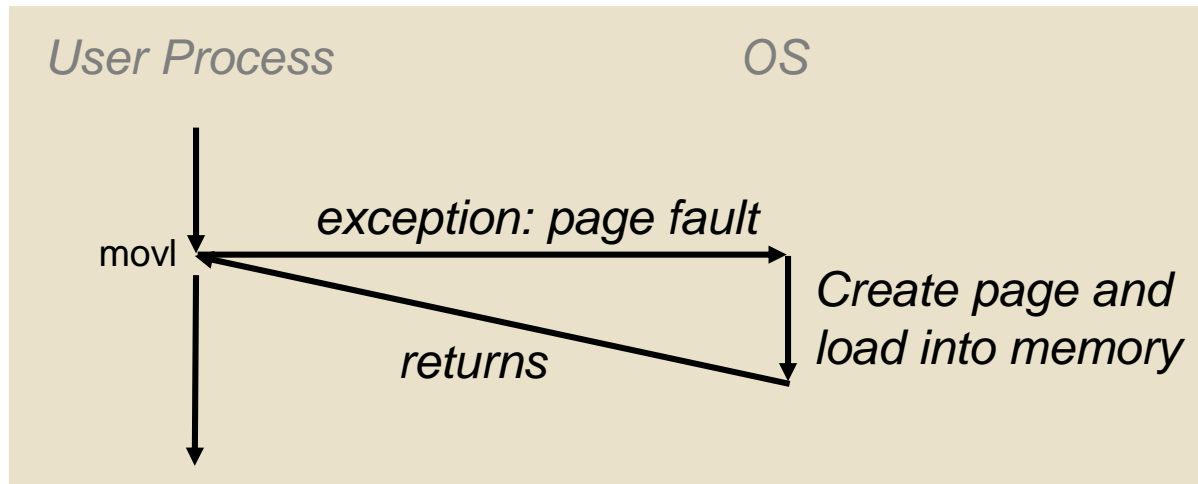
- OS must find or create file, get it ready for reading or writing
- Returns integer file descriptor

Fault Example: Page Fault

- User writes to memory location
- That portion (page) of user's memory is currently on disk

```
int a[1000];  
main ()  
{  
    a[500] = 13;  
}
```

```
80483b7:      c7 05 10 9d 04 08 0d  movl    $0xd,0x8049d10
```

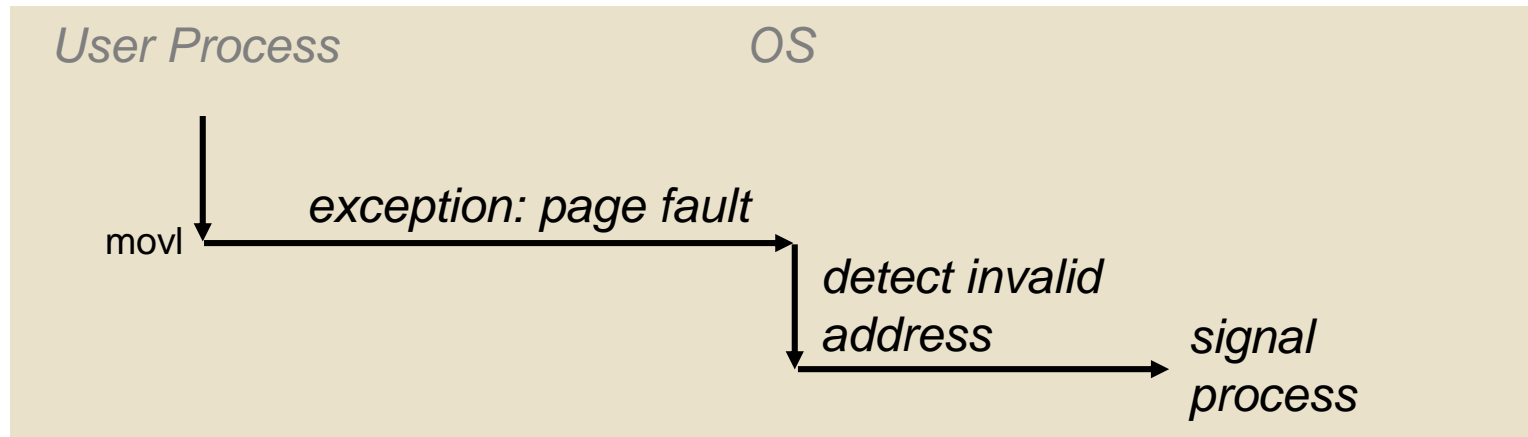


- Page handler must load page into physical memory
- Returns to faulting instruction
- Successful on second try

Fault Example: Invalid Memory Reference

```
int a[1000];  
main ()  
{  
    a[5000] = 13;  
}
```

80483b7: c7 05 60 e3 04 08 0d movl \$0xd,0x804e360



- Page handler detects invalid address
- Sends SIGSEGV signal to user process
- User process exits with “segmentation fault”

Asynchronous Exceptions (Interrupts)

- Caused by events external to the processor
 - Indicated by setting the processor's interrupt pin
 - Handler returns to “next” instruction
- Examples:
 - I/O interrupts
 - ▶ hitting Ctrl-C at the keyboard
 - ▶ arrival of a packet from a network
 - ▶ arrival of data from a disk
 - Hard reset interrupt
 - ▶ hitting the reset button
 - Soft reset interrupt
 - ▶ hitting Ctrl-Alt-Delete on a PC

I/O Interrupts


- CPU interrupt-request line triggered by I/O device
- Interrupt handler receives interrupt
- AND with *interrupt mask register* to ignore or delay some interrupts
- Interrupt vector routine determines the source of the interrupt and dispatches it to the correct handler
 - priority based
 - some interrupts are non-maskable (NMI)
- The same mechanism is used for exceptions

Basic x86 Interrupts

■ PIC: Intel 8259(A)


- introduced in 1976, survived for 30 years, now being phased out in favor of the Intel APIC Architecture.
- 8259A interface still provided by the southbridge chipset on x86 motherboards

■ x86 CPUs




- INTR: interrupt request
- INTA: interrupt acknowledge
- IRQ0~7: interrupt request lines
- cascading possible
- NMI: non-maskable interrupt

x86 Interrupts: NMI



- NMI asserted
 - CPU completes current instruction
 - issues exception #2
- Cannot be disabled
- Reserved for
 - major hardware faults
 - ▶ memory parity error
 - ▶ watchdog timers
- Multiple sources: who caused the NMI?
 - OS handler must poll potential sources
 - for normal use: too slow, second NMI might occur while first is still being handled

x86 Interrupts: IRQ



- Interrupt request
 - can be disabled with the `CLI/STI` instruction (clears/sets IE status flag)
 - if enabled, complete current instruction, then
- CPU acknowledges interrupt using `INTA` pin
- Interrupt vector is then supplied on the data bus
- CPU issues exception # from the vector

Implementation of Linux/IA32 System Calls

- Application programs use system calls to request a service from the kernel
- System calls are invoked via the `int 0x80` instruction
 - parameters passed through registers (not the stack)
 - ▶ `eax`: system call number
 - ▶ `ebx, ecx, edx, esi, edi, ebp`: (up to) 6 arbitrary arguments (interpretation depends on the invoked system call)
 - ▶ `esp` cannot be used (overwritten when the CPU enters kernel mode)
- `libc` provides convenient wrappers for most system calls

Implementation of Linux/IA32 System Calls

■ Directly invoking Linux/IA32 System calls

```
#include <stdlib.h>

int main(void)
{
    write(1, "hello, world\n", 13);
    exit(0);
}
```

```
$ gcc -o hello hello.c
$ ./hello
hello, world
$
```

```
int main()
{
    my_write(1, "hello, world\n", 13);
    my_exit(0);
}
```

myhello.c

```
.section .text
.globl my_write
.globl my_exit

# void my_write(int fd, char *string,
                unsigned int length);
my_write:

    push    %ebp
    mov     %esp, %ebp
    push    %ebx

    mov     $4, %eax           # eax: 4 (write syscall)
    mov     8(%ebp), %ebx      # ebx: fd (ebp+8)
    mov     12(%ebp), %ecx     # ecx: string (ebp+12)
    mov     16(%ebp), %edx     # edx: length (ebp+16)
    shr     $1, %edx          # ecx: string (ebp+12)

    int     $0x80             # invoke system call

    pop     %ebx
    pop     %ebp

    ret

...
```

mylib.c

Linux/IA32 System Calls

■ Directly invoking Linux/IA32 System calls

```
int main()
{
    my_write(1, "hello, world\n", 13);
    my_exit(0);
}                                     myhello.c
```

```
...
# void my_exit(int nr)
my_exit:

    push    %ebp
    mov     %esp, %ebp
    push    %ebx

    mov     $1, %eax                # eax: 1 (exit syscall)
    mov     8(%ebp), %ebx           # ebx: nr (ebp+8)

    int     $0x80                   # invoke system call

    pop     %ebx
    pop     %ebp

    ret                                     mylib.c
```

```
$ gcc -o my_hello myhello.c mylib.s
$ ./my_hello
hello, world
$
```