

Introduction to System Programming



Introduction to System Programming

- **Why system programming?**
- Basic operation of a computer system
- Summary

Acknowledgement: slides based on the cs:app2e material

Why System Programming?

“I program in Java (JavaScript, Ruby, Perl, Python, PHP, Go, ...). I don't need to know this stuff.”

Why System Programming?

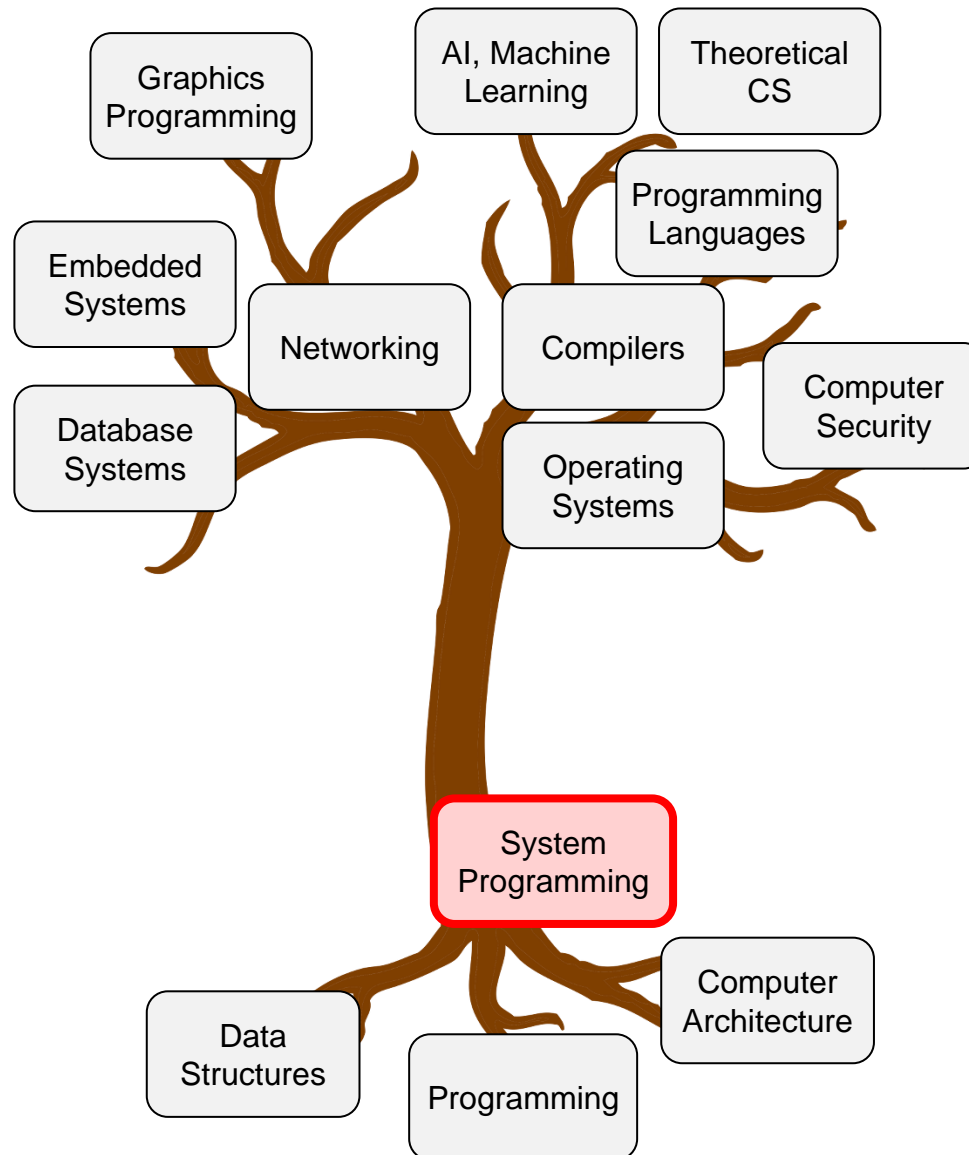
■ Abstractions only partially reflect reality

- To be a good computer engineer you have to know that abstractions in computer science are only abstractions; the reality may look very different, especially in special cases
- To understand the limitations of these abstractions and write safe and efficient code you need to know
 - ▶ computer architecture
 - ▶ assembly programming
 - ▶ system programming

■ Programming, data structures, computer architecture, and system programming prepare you for more advanced classes in computer science & engineering

- operating systems, compilers, embedded systems, networks, data base systems, ...

Why System Programming?



Reality Check #1: Integer and Real Numbers

- Integer: a number with no fractional part

..., -4, -3, -2, -1, 0, 1, 2, 3, 4, ...

- Real: rational (integers, fractions) and irrational numbers ($\sqrt{2}$, π)

..., -4, -3/2, - $\sqrt{2}$, 0, 1, 1.1, 1.11, π , 100, ...

- Abstraction: in a computer program

- `integer = int` (or `long`)
- `real = float` (or `double`)

Integer Numbers

■ integer.c

```
#include <stdio.h>

void main(void)
{
    int i = 1000000;
    int k;

    for (k = 0; k<20; k++) {
        printf("%2d:  %d\n", k, i);
        i = i * 2;
    }
}
```

■ Compile and run

```
$ gcc -o integer integer.c
$ ./integer
```

Integer Numbers

■ Result

```
$ gcc -o integer integer.c
$ ./integer
0: 1000000
1: 2000000
2: 4000000
3: 8000000
4: 16000000
5: 32000000
6: 64000000
7: 128000000
8: 256000000
9: 512000000
10: 1024000000
11: 2048000000
12: -198967296
13: -397934592
14: -795869184
15: -1591738368
16: 1111490560
17: -2071986176
18: 150994944
19: 301989888
$
```

```
#include <stdio.h>

void main(void)
{
    int i = 1000000;
    int k;

    for (k = 0; k<20; k++) {
        printf("%2d:  %d\n", k, i);
        i = i * 2;
    }
}
```


Real Numbers

■ real.c

```
#include <stdio.h>

void main(void)
{
    float a = 1e30;
    float b = 3.14159265358979;

    printf(" 1. %f + %f - %f = %f\n", a,b,a,a+b-a);
    printf(" 2. %f - %f + %f = %f\n", a,a,b,a-a+b);
}
```

■ Compile and run

```
$ gcc -o real real.c
$ ./real
```

Real Numbers

■ Result

```
#include <stdio.h>

void main(void)
{
    float a = 1e30;
    float b = 3.14159265358979;

    printf(" 1. %f + %f - %f = %f\n", a,b,a,a+b-a);
    printf(" 2. %f - %f + %f = %f\n", a,a,b,a-a+b);
}
```

```
$ gcc -o real real.c
$ ./real
 1. 1000000015047466219876688855040.000000 + 3.141593 - 1000000015047466219876688855040.000000 = 0.000000
 2. 1000000015047466219876688855040.000000 - 1000000015047466219876688855040.000000 + 3.141593 = 3.141593
$
```

Code Security Example

```
/* Kernel memory region holding user-accessible data */
#define KSIZE 1024
char kbuf[KSIZE];

/* Copy at most maxlen bytes from kernel region to user buffer */
int copy_from_kernel(void *user_dest, int maxlen) {
    /* Byte count len is minimum of buffer size and maxlen */
    int len = KSIZE < maxlen ? KSIZE : maxlen;
    memcpy(user_dest, kbuf, len);
    return len;
}

/* memcpy definition */
void *memcpy(void *dest, const void *src, size_t n);
```

- Similar to code found in FreeBSD's implementation of getpeername
- There are legions of smart people trying to find vulnerabilities in programs

Typical Usage

```
/* Kernel memory region holding user-accessible data */
#define KSIZE 1024
char kbuf[KSIZE];

/* Copy at most maxlen bytes from kernel region to user buffer */
int copy_from_kernel(void *user_dest, int maxlen) {
    /* Byte count len is minimum of buffer size and maxlen */
    int len = KSIZE < maxlen ? KSIZE : maxlen;
    memcpy(user_dest, kbuf, len);
    return len;
}

/* memcpy definition */
void *memcpy(void *dest, const void *src, size_t n);
```

```
#define MSIZE 528

void getstuff() {
    char mybuf[MSIZE];
    copy_from_kernel(mybuf, MSIZE);
    printf("%s\n", mybuf);
}
```

Malicious Usage

```
/* Kernel memory region holding user-accessible data */
#define KSIZE 1024
char kbuf[KSIZE];

/* Copy at most maxlen bytes from kernel region to user buffer */
int copy_from_kernel(void *user_dest, int maxlen) {
    /* Byte count len is minimum of buffer size and maxlen */
    int len = KSIZE < maxlen ? KSIZE : maxlen;
    memcpy(user_dest, kbuf, len);
    return len;
}

/* memcpy definition */
void *memcpy(void *dest, const void *src, size_t n);
```

```
#define MSIZE 528

void getstuff() {
    char mybuf[MSIZE];
    copy_from_kernel(mybuf, -MSIZE);
    ...
}
```

To Understand What Goes Wrong

```
/* Kernel memory region holding user-accessible data */
#define KSIZE 1024
char kbuf[KSIZE];

/* Copy at most maxlen bytes from kernel region to user buffer */
int copy_from_kernel(void *user_dest, int maxlen {
    /* Byte count len is minimum of buffer size and maxlen */
    int len = KSIZE < maxlen ? KSIZE : maxlen;
    memcpy(user_dest, kbuf, len);
    return len;
}

/* memcpy definition */
void *memcpy(void *dest, const void *src, size_t n);
```

```
#define MSIZE 528

void getstuff() {
    char mybuf[MSIZE];
    copy_from_kernel(mybuf, -MSIZE);
    ...
}
```

Computer Arithmetic

- Does not generate random values
 - Arithmetic operations have important mathematical properties
- Cannot assume all “usual” mathematical properties
 - Due to finiteness of representations
 - Integer operations satisfy “ring” properties
 - ▶ commutativity, associativity, distributivity
 - Floating point operations satisfy “ordering” properties
 - ▶ Monotonicity, values of signs
- Observation
 - Need to understand which abstractions apply in which contexts
 - Important issues for compiler writers and serious application programmers

Reality Check #2:

You've Got to Know Assembly

- Chances are, you'll never write programs in assembly
 - Compilers are much better & more patient than you are
- But: Understanding assembly is key to machine-level execution model
 - Behavior of programs in presence of bugs
 - ▶ High-level language models break down
 - Tuning program performance
 - ▶ Understand optimizations done / not done by the compiler
 - ▶ Understanding sources of program inefficiency
 - Implementing system software
 - ▶ Compiler has machine code as target
 - ▶ Operating systems must manage process state
 - Creating / fighting malware
 - ▶ x86 assembly is the language of choice!

Assembly Code Example

■ Time Stamp Counter

- Special 64-bit register in Intel-compatible machines
- Incremented every clock cycle
- Read with rdtsc instruction

■ Application

- Measure time (in clock cycles) required by procedure

```
double t;  
start_counter();  
P();  
t = get_counter();  
printf("P required %f clock cycles\n", t);
```

Code to Read Counter

- Write small amount of assembly code using GCC's asm facility
- Inserts assembly code into machine code generated by compiler

```
static unsigned cyc_hi = 0;
static unsigned cyc_lo = 0;

/* Set *hi and *lo to the high and low order bits
   of the cycle counter.
*/
void access_counter(unsigned *hi, unsigned *lo)
{
    asm("rdtsc; movl %%edx,%0; movl %%eax,%1"
        : "=r" (*hi), "=r" (*lo)
        :
        : "%edx", "%eax");
}
```

Reality Check #3: Random Access Memory

■ MatrixAdd.java

```
import java.util.*;

class MatrixAdd
{
    static int N=4096;
    static int R=16;

    static private long MatAdd(int N, int A[][],
                               int B[][], int sum[][])
    {
        long start = System.nanoTime();
        int i,j;

        for (i=0; i<N; i++) {
            for (j=0; j<N; j++) {
                sum[j][i] = A[j][i] + B[j][i];
            }
        }

        long stop = System.nanoTime();

        return stop - start;
    }
    ...
}
```

```
...

public static void main(String args[])
{
    System.out.format("Adding two %,d x %,d matrices.%n",
                      N, N);

    int a[][] = new int[N][N];
    int b[][] = new int[N][N];
    int c[][] = new int[N][N];
    int i,j;

    Random rand = new Random();

    for (j=0; j<N; j++) {
        for (i=0; i<N; i++) {
            a[j][i] = rand.nextInt(65536);
            b[j][i] = rand.nextInt(65536);
        }
    }

    long total = 0;

    for (i=1; i<=R; i++) {
        long t = MatAdd(N, a, b, c);
        total += t;
        System.out.format(" %2d. run: %,d%n", i, t);
    }

    System.out.format("Average runtime: %,d.%n",
                      total / R);
}
```

■ Compile and run

```
$ javac MatrixAdd.java
$ java MatrixAdd
```

Reality Check #3: Random Access Memory

■ Result

```
$ javac MatrixAdd.java
$ java MatrixAdd
Adding two 4,096 x 4,096 matrices.
 1. run: 1,167,219,175
 2. run: 1,170,854,849
 3. run: 1,157,460,403
 4. run: 1,146,491,235
 5. run: 1,131,016,361
 6. run: 1,127,804,280
 7. run: 1,176,178,946
 8. run: 1,174,429,930
 9. run: 1,123,144,337
10. run: 1,148,392,584
11. run: 1,158,618,876
12. run: 1,143,570,880
13. run: 1,136,663,775
14. run: 1,144,929,208
15. run: 1,130,314,085
16. run: 1,145,745,940
Average runtime: 1,148,927,179.
```

- Then your friend edits the file for 3 seconds, recompiles and runs it.

Reality Check #3: Random Access Memory

■ Result of modified source

```
$ javac MatrixAdd.java
$ java MatrixAdd
Adding two 4,096 x 4,096 matrices.
1. run: 22,371,258
2. run: 15,699,727
3. run: 17,192,265
4. run: 16,826,659
5. run: 17,135,681
6. run: 16,873,783
7. run: 16,837,033
8. run: 17,470,509
9. run: 17,435,977
10. run: 17,528,881
11. run: 17,451,922
12. run: 17,427,301
13. run: 17,342,101
14. run: 17,388,893
15. run: 17,271,563
16. run: 17,318,061
Average runtime: 17,473,225.
```

■ Factor $1,148,927,179 / 17,473,225 = 65$ faster ?!

Reality Check #3: Random Access Memory

■ Source comparison

Your code:

```
import java.util.*;

class MatrixAdd
{
    static int N=4096;
    static int R=16;

    static private long MatAdd(int N, int A[][],
                               int B[][], int sum[][])
    {
        long start = System.nanoTime();
        int i,j;

        for (i=0; i<N; i++) {
            for (j=0; j<N; j++) {
                sum[j][i] = A[j][i] + B[j][i];
            }
        }

        long stop  = System.nanoTime();

        return stop - start;
    }

    ...
}
```

Your friend's code:

```
import java.util.*;

class MatrixAdd
{
    static int N=4096;
    static int R=16;

    static private long MatAdd(int N, int A[][],
                               int B[][], int sum[][])
    {
        long start = System.nanoTime();
        int i,j;

        for (j=0; j<N; j++) {
            for (i=0; i<N; i++) {
                sum[j][i] = A[j][i] + B[j][i];
            }
        }

        long stop  = System.nanoTime();

        return stop - start;
    }

    ...
}
```

Random Access Memory Is an Unphysical Abstraction

- Memory is not unbounded
 - It must be allocated and managed
 - Many applications are memory dominated
- Memory referencing bugs especially pernicious
 - Effects are distant in both time and space
- Memory performance is not uniform
 - Cache and virtual memory effects can greatly affect program performance
 - Adapting programs to characteristics of a memory system can lead to major speed improvements

Memory System Performance Example

```
void copyij(int src[2048][2048],
            int dst[2048][2048])
{
    int i,j;
    for (i = 0; i < 2048; i++)
        for (j = 0; j < 2048; j++)
            dst[i][j] = src[i][j];
}
```




```
void copyji(int src[2048][2048],
            int dst[2048][2048])
{
    int i,j;
    for (j = 0; j < 2048; j++)
        for (i = 0; i < 2048; i++)
            dst[i][j] = src[i][j];
}
```

Effect of swapping two lines:
21 times slower (Pentium 4)
6 times slower (Core i7)

- Hierarchical memory organization
- Performance depends on access patterns
 - Including how step through multi-dimensional array

The Memory Mountain

Intel Core i7
2.67 GHz
32 KB L1 d-cache
256 KB L2 cache
8 MB L3 cache



Memory Referencing Bug Example

```
double fun(int i)
{
    volatile double d[1] = {3.14};
    volatile long int a[2];
    a[i] = 1073741824; /* Possibly out of bounds */
    return d[0];
}
```

■ Result is architecture specific

fun(0)	→	3.14
fun(1)	→	3.14
fun(2)	→	3.13999998664856
fun(3)	→	2.000000061035156
fun(4)	→	3.14, then segmentation fault

Memory Referencing Bug Example

```
double fun(int i)
{
    volatile double d[1] = {3.14};
    volatile long int a[2];
    a[i] = 1073741824; /* Possibly out of bounds */
    return d[0];
}
```

fun(0) → 3.14
fun(1) → 3.14
fun(2) → 3.13999998664856
fun(3) → 2.000000061035156
fun(4) → 3.14, then segmentation fault

Explanation:

Saved State	4	} Location accessed by fun(i)
d[0] (high 32bits)	3	
d[0] (low 32bit)	2	
a[1]	1	
a[0]	0	

Memory Referencing Errors


- C and C++ do not provide any memory protection
 - Out of bounds array references
 - Invalid pointer values
 - Abuses of malloc/free
- Can lead to nasty bugs
 - Whether or not bug has any effect depends on system and compiler
 - Action at a distance
 - ▶ Corrupted object logically unrelated to one being accessed
 - ▶ Effect of bug may be first observed long after it is generated
- How can I deal with this?
 - Program in a language with runtime bound checks (Java, Ruby, Pascal, ML,...)
 - Understand what possible interactions may occur
 - Use or develop tools to detect referencing errors (e.g. Valgrind)

Reality #4: There's more to performance than asymptotic complexity

- Constant factors matter too!
- And even exact op count does not predict performance
 - Easily see 10:1 performance range depending on how code written
 - Must optimize at multiple levels: algorithm, data representations, procedures, and loops
- Must understand system to optimize performance
 - How programs compiled and executed
 - How to measure program performance and identify bottlenecks
 - How to improve performance without destroying code modularity and generality

Example: Matrix Multiplication


Matrix-Matrix Multiplication (MMM) on 2 x Core 2 Duo 3 GHz (double precision)



- Standard desktop computer, vendor compiler, using optimization flags
- Both implementations have exactly the same operations count ($2n^3$)
- What is going on?

MMM Plot: Analysis

Matrix-Matrix Multiplication (MMM) on 2 x Core 2 Duo 3 GHz



- Reason for 20x: Blocking or tiling, loop unrolling, array scalarization, instruction scheduling, search to find best choice
- Effect: fewer register spills, L1/L2 cache misses, and TLB misses

Reality #5:

Computers do more than execute programs

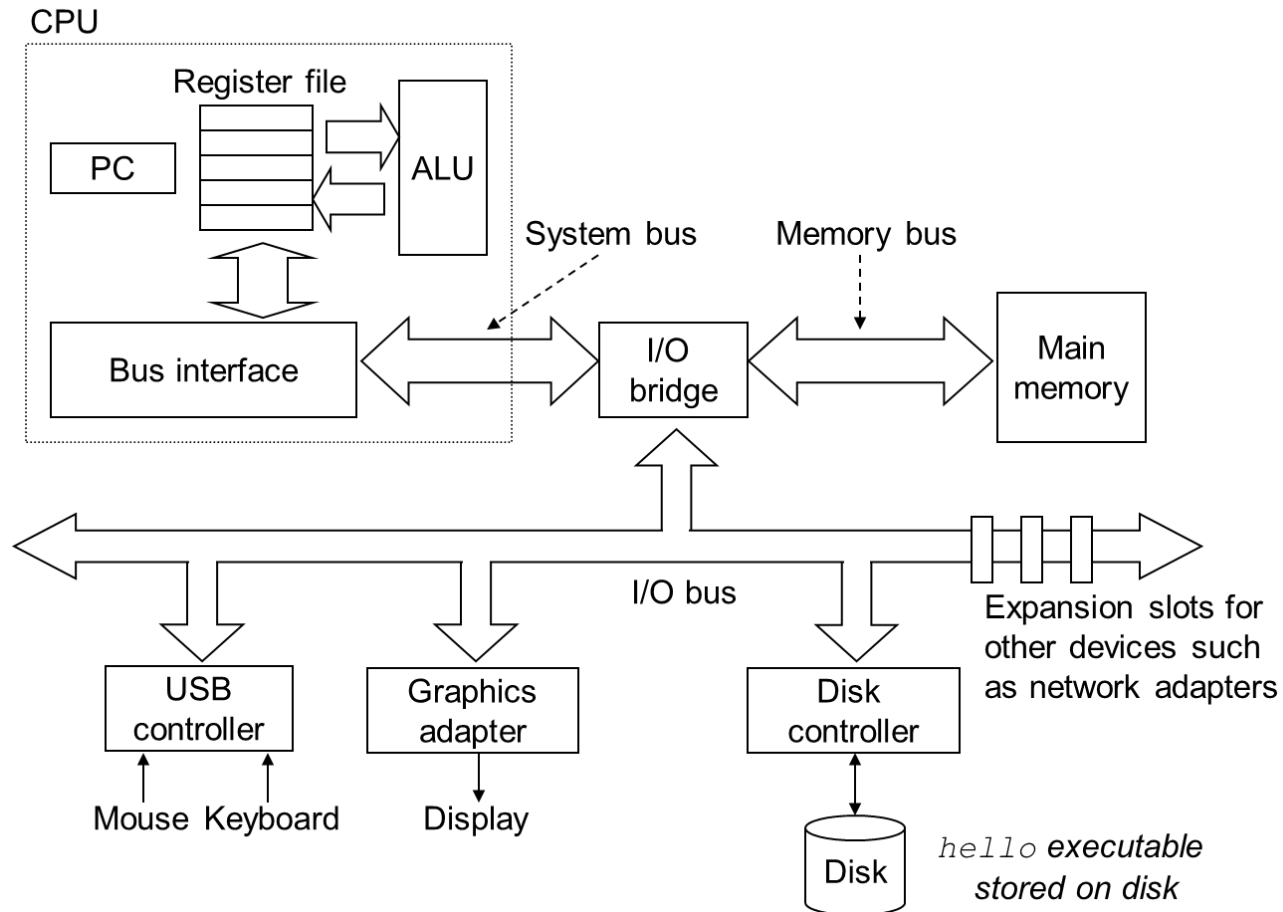
- They need to get data in and out
 - I/O system critical to program reliability and performance
- They communicate with each other over networks
 - Many system-level issues arise in presence of network
 - ▶ Concurrent operations by autonomous processes
 - ▶ Coping with unreliable media
 - ▶ Cross platform compatibility
 - ▶ Complex performance issues

Introduction to System Programming

- Abstractions are good but don't forget reality
- **Basic operation of a computer system**
- Summary


Hardware Organization

■ Typical hardware organization




Operating System Basics

- The operating system manages the hardware
 - protect H/W from misuse by buggy/malicious programs
 - provide simple and uniform mechanisms for manipulating hardware devices



- Fundamental abstractions provided by the OS

- processes
- virtual memory
- files



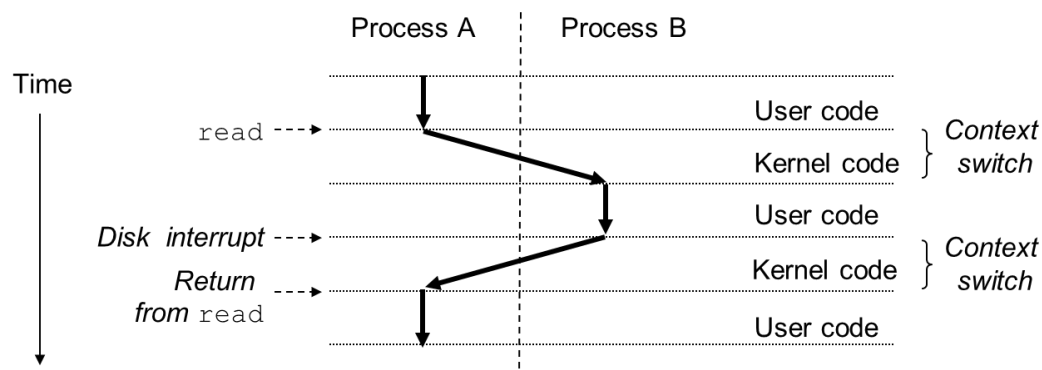
Processes and Threads

- A process is the operating systems abstraction for a running program
- multiple processes can run concurrently
 - multi-core processors: true parallelism
 - single-cores: apparent parallelism through context-switching
- OS provides the illusion that the process has exclusive access to the H/W
 - full memory address space
 - exclusive access to the I/O devices

Processes and Threads

■ Context-Switching

- OS keeps track of all state (context) that a process needs in order to run
 - ▶ current PC
 - ▶ register file
 - ▶ memory contents
 - ▶ open files
- switch to a new process
 - ▶ periodically
 - ▶ when a process has to wait for an event

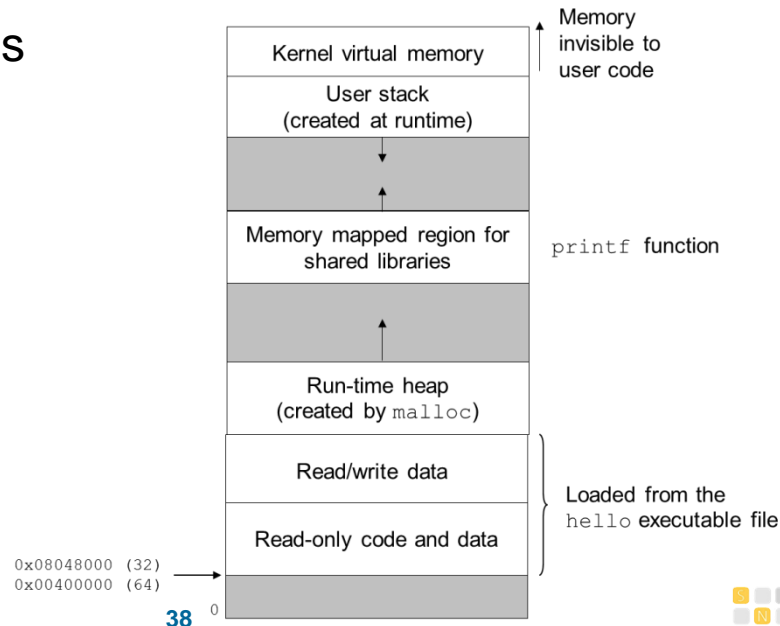


■ A process can consist of multiple *threads*

- not as heavy as full processes
- easier sharing of data
- typically more efficient scheduling

Virtual Memory

- An abstraction of the physical memory
- Provides each process with the illusion that it has exclusive use of the main memory
- Managed by the OS with the help of a hardware translation unit, the MMU (memory management unit)
- Typical layout on Linux systems



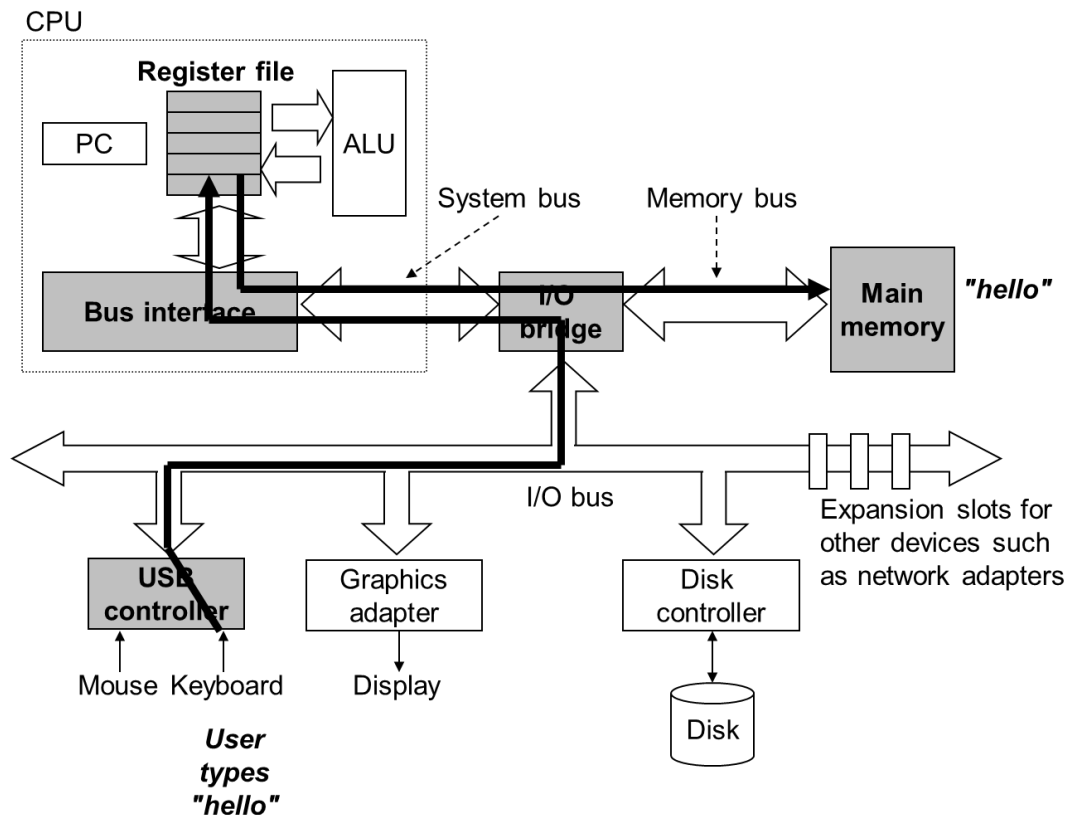
Files

- a file is a sequence of bytes
- in *nix systems, “everything” is modeled as a file
 - disk
 - keyboard
 - mouse
 - display
 - network
 - shared memory
- single interface to interact with files

Running a Program on a System

■ Running `hello`

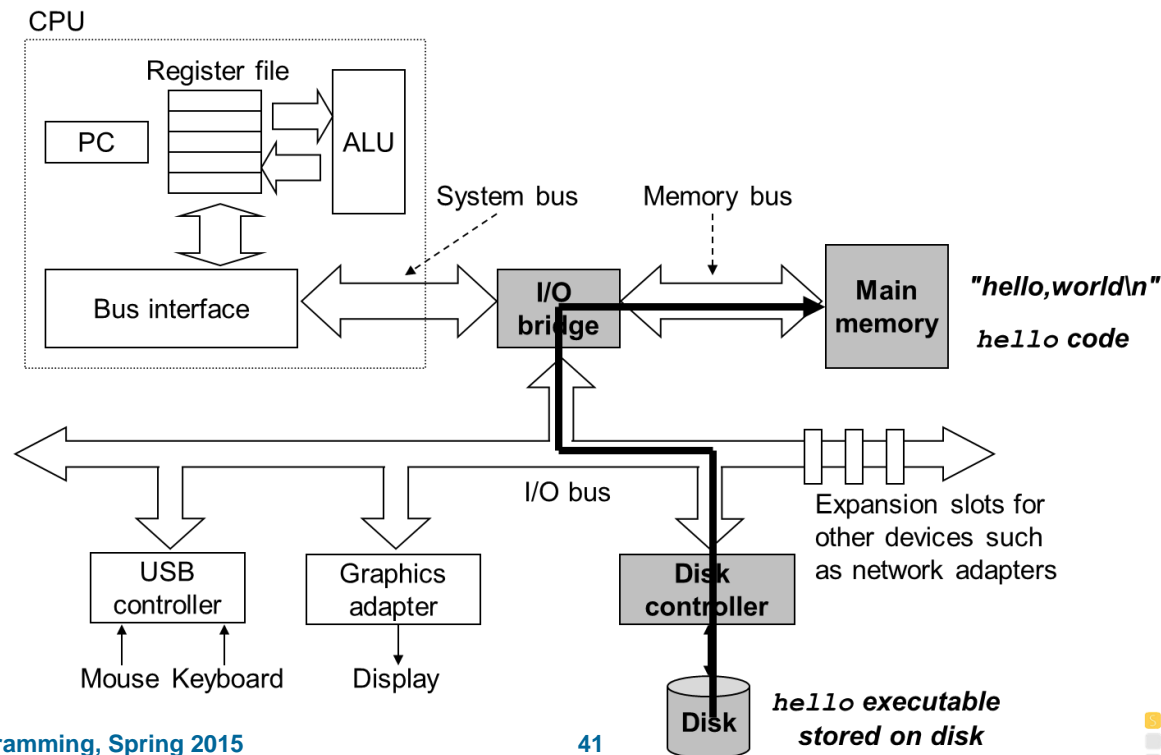
- shell reads input from keyboard, one character at a time into a register and from there to the main memory



Running a Program on a System

■ Running `hello`

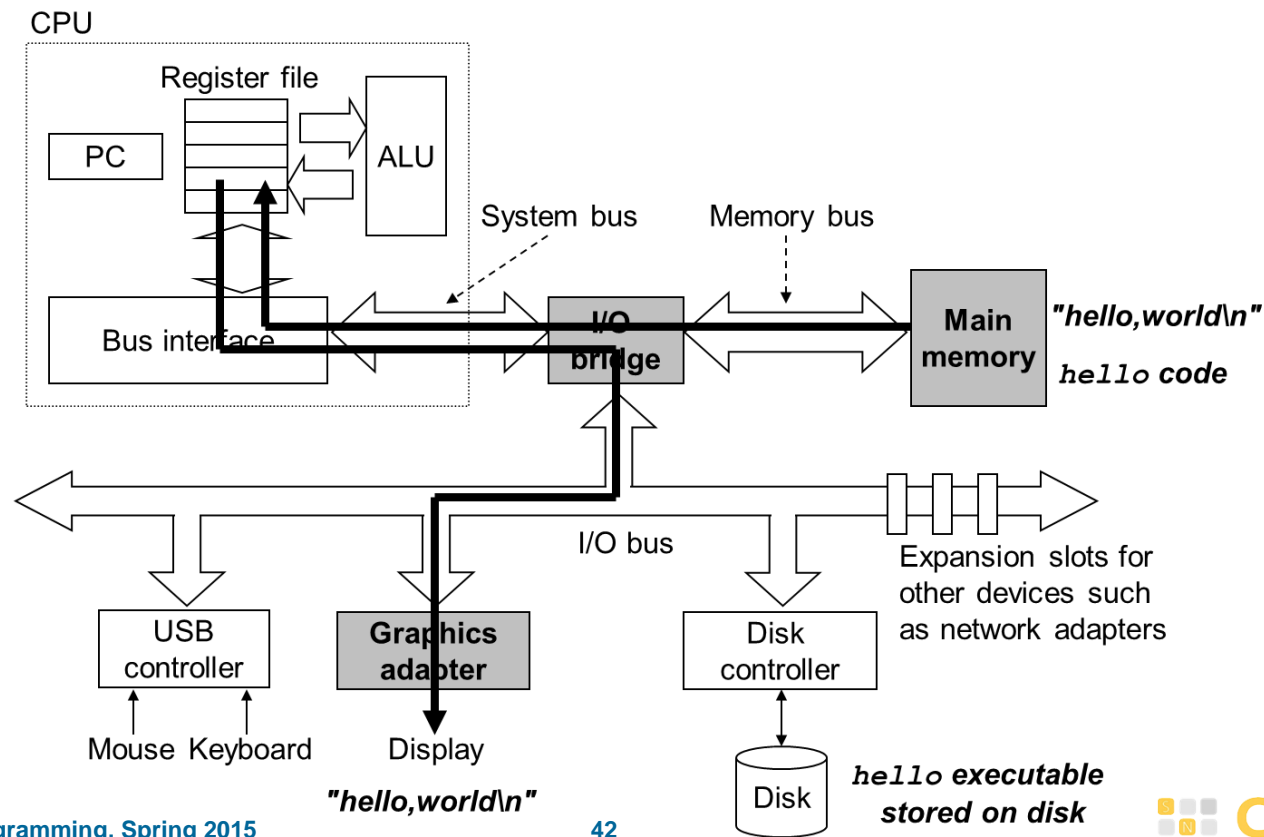
- `<enter>` signals the end of the command to the shell
- the shell searches for a file named 'hello' in the file system
- loads the program from disk into memory
- starts running the program



Running a Program on a System

■ Running `hello`

- the processor begins executing the machine-instructions
- copies the string 'hello, world\n' to registers, and from there to the display adapter via *system calls*



Introduction to System Programming

- Abstractions are good but don't forget reality
- Basic operation of a computer system
- **Summary**

System Programming

- System software (or systems software) is computer software or an operating system designed to operate and control the computer hardware and to provide a platform for running application software.
(source: Wikipedia)
- System Programming
 - concerns itself with writing system software
 - from relatively high-level (web servers) down to very low-level (boot loaders)

Wrapping it up

■ In this class you will learn

- learn how to interact with the operating system on a low level
- get an idea of what the operating system is doing
- how to write safer code
- how to write faster code

■ Prerequisites

- some programming experience
- computer architecture
- x86 assembly
- SP will require a substantial effort
 - ▶ read the book
 - ▶ do the homework
 - ▶ start with labs early!

Reading Assignment

- For Thursday, March 4
 - chapter 1 of the text book