

# Running Programs on a System

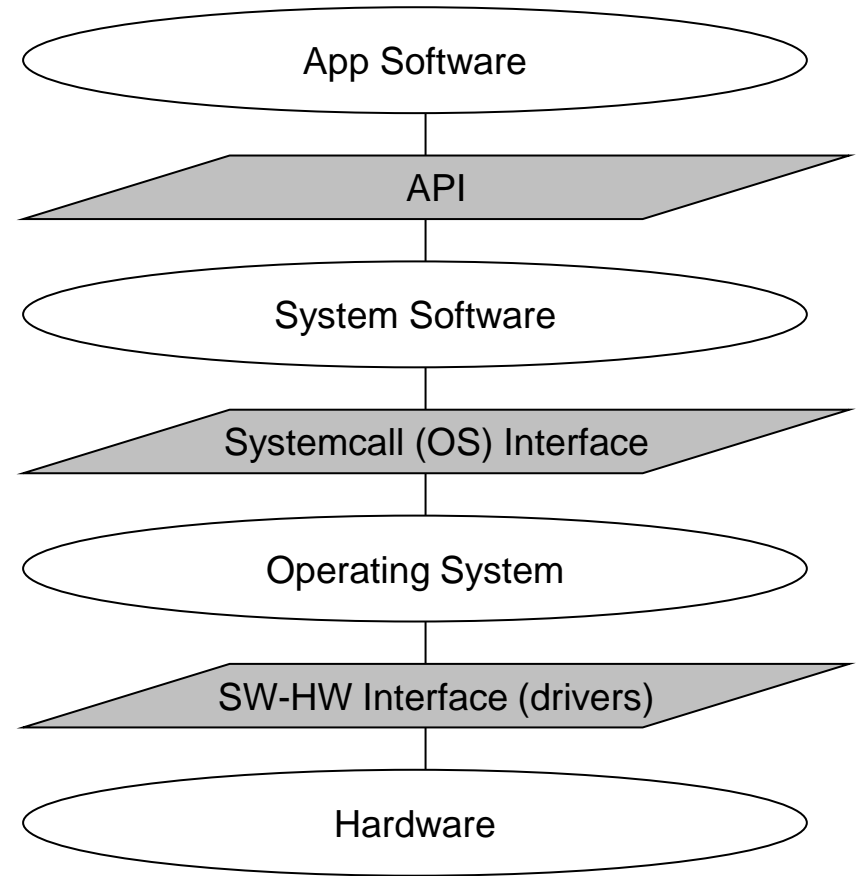
## Process Management



# Process Management

## ■ Linux (system)calls for process management

- `fork()`
- `exec()` and variants
- `wait()` and variants
- `exit()`




# Process Concept

- Definition: A process is an instance of a running program.
  - One of the most profound ideas in computer science
  - Not the same as “program” or “processor”
- Process provides each program with two key abstractions:
  - Logical control flow
    - ▶ Each program seems to have exclusive use of the CPU
  - Private virtual address space
    - ▶ Each program seems to have exclusive use of main memory
- How are these Illusions maintained?
  - Process execution in sequence, interleaved, or run on separate cores
  - Address spaces managed by virtual memory system

# Process Concept

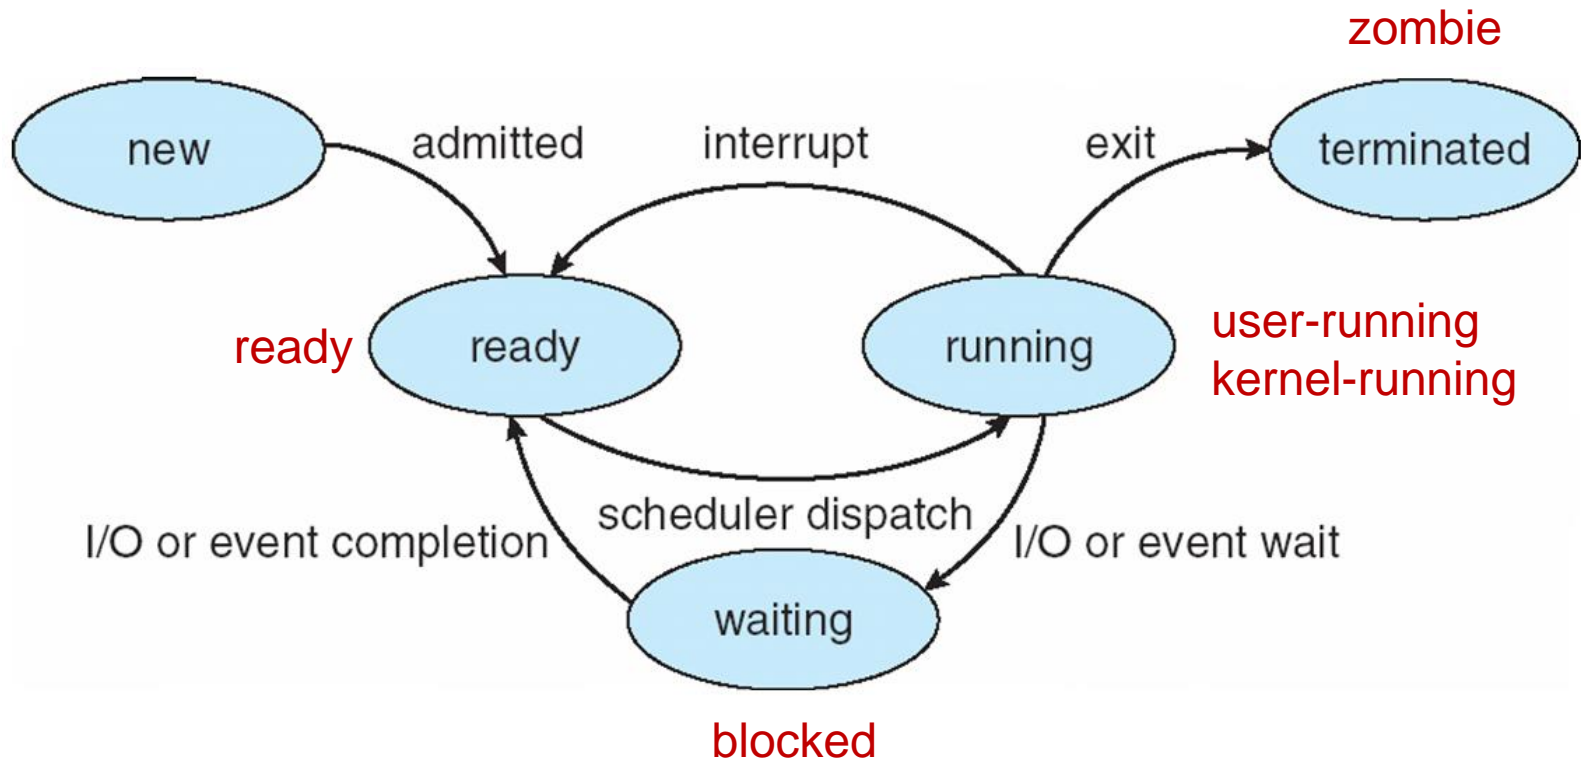
- A process includes:
  - program counter
  - stack
  - data section
  - open files



# Process State

- As a process executes, it changes state
  - new: The process is being created
  - running: Instructions are being executed
  - waiting: The process is waiting for some event to occur
  - ready: The process is waiting to be assigned to a processor
  - terminated: The process has finished execution

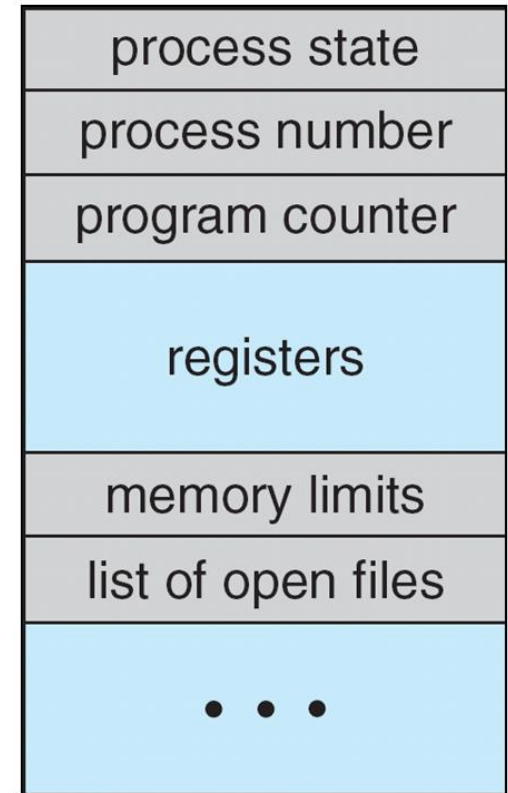
# Diagram of Process State



in red: Linux terminology

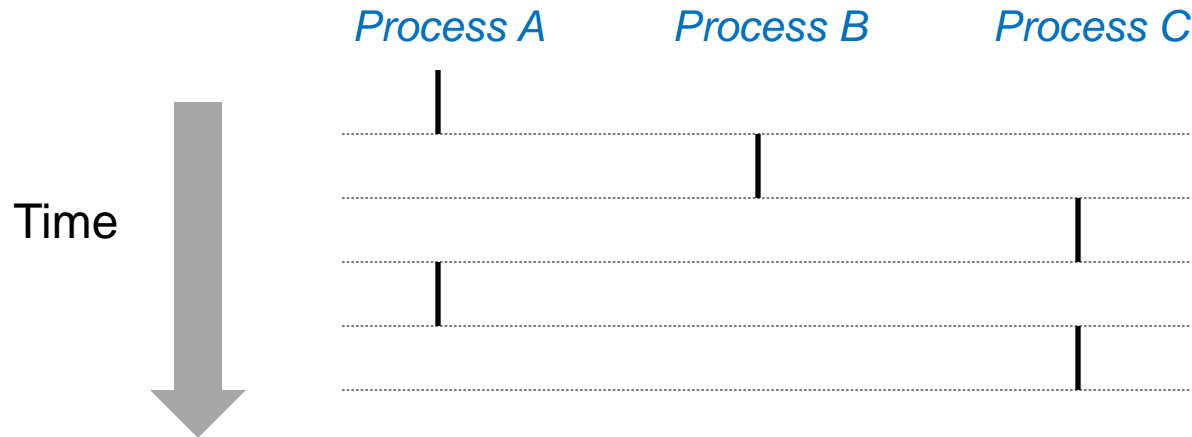
# Process Control Block (PCB)

- Information associated with each process
  - Process state
  - Program counter
  - CPU registers
  - CPU scheduling information
  - Memory-management information
  - Accounting information
  - I/O status information
- In Linux: task\_struct, thread\_struct
  - include/linux/sched.h



# Concurrent Processes

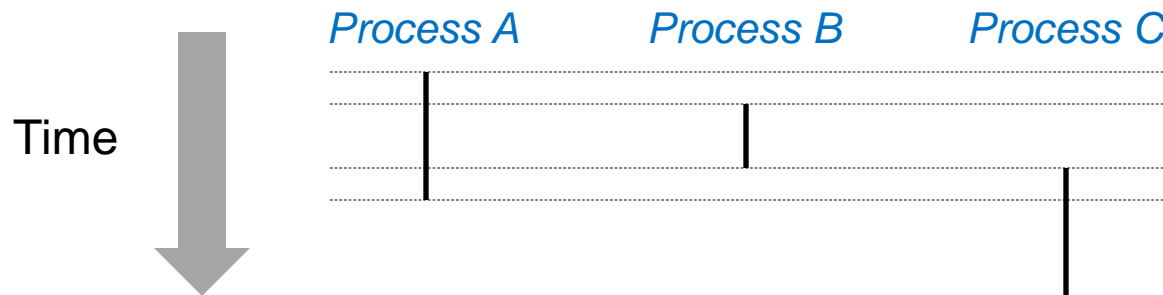
- Two processes run concurrently (are concurrent) if their flows overlap in time
- Otherwise, they are sequential
- Examples (running on single core):
  - Concurrent: A & B, A & C
  - Sequential: B & C





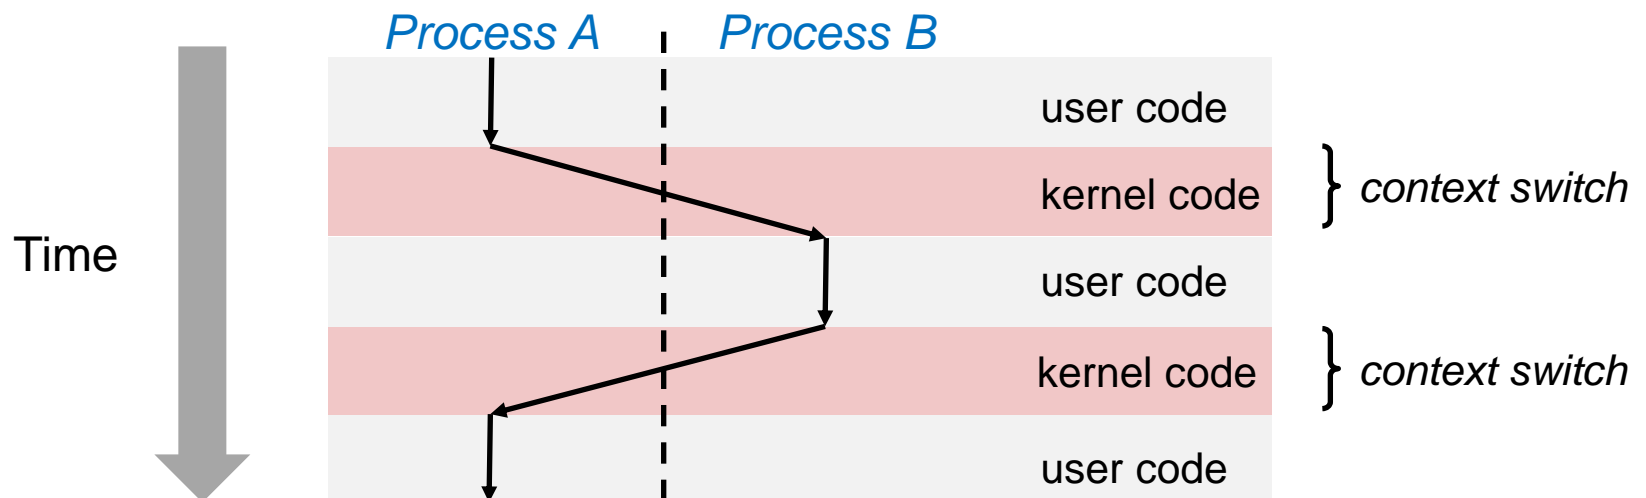
# User View of Concurrent Processes

- Control flows for concurrent processes are physically disjoint in time
- However, we can think of concurrent processes are running in parallel with each other




# Context Switching

- Processes are managed by a shared chunk of OS code called the kernel
  - the kernel is not a separate process, but rather runs as part of some user process
- Control flow passes from one process to another via a context switch
  - save state of old process, load saved state for new process (in PCB)
  - Context-switch time is overhead; the system does no useful work while switching
  - Time dependent on hardware support



# Context Switch Details

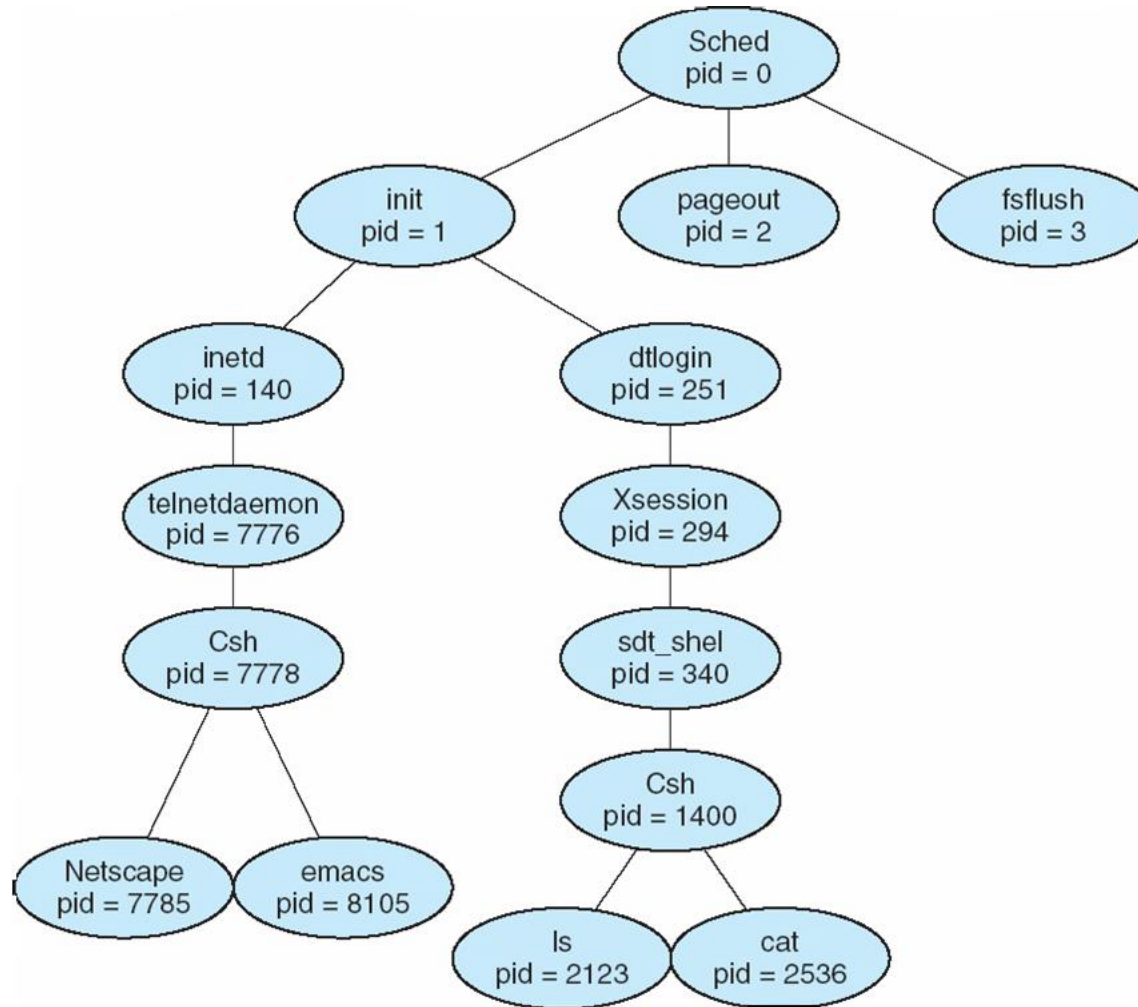


# Process Creation



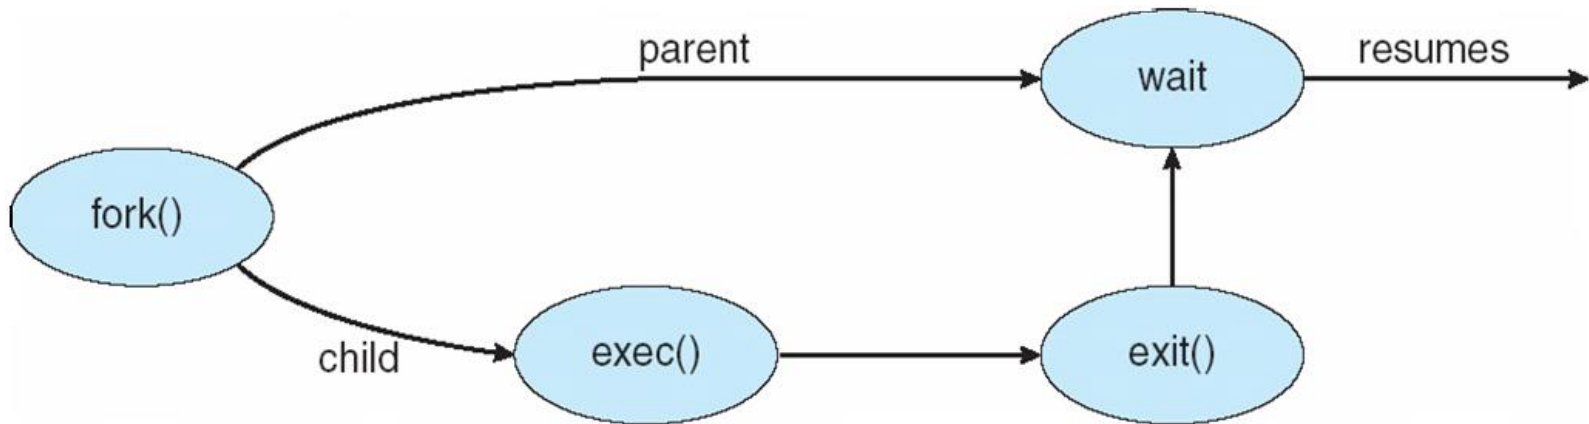
- A process is always created by a parent process
  - tree-like hierarchy of processes
- Process identified and managed via **a process identifier (PID)**
- Resource sharing options
  - Parent and children share all resources
  - Children share subset of parent's resources
  - Parent and child share no resources
- Execution
  - Parent and children execute concurrently
  - Parent waits until child/children terminate

# A tree of processes on a typical Solaris System



# Process Creation (Cont)

- Address space
  - Child duplicate of parent
  - Child has a program loaded into it



# C Program Forking Separate Process

```
int main()
{
    pid_t  pid;

    /* fork another process */
    pid = fork();
    if (pid < 0) { /* error occurred */
        fprintf(stderr, "Fork Failed");
        exit(-1);
    }
    else if (pid == 0) { /* child process */
        execlp("/bin/ls", "ls", NULL);
    }
    else { /* parent process */
        /* parent will wait for the child to complete */
        wait (NULL);
        printf ("Child Complete");
        exit(0);
    }
}
```

# Process Termination

- Process executes last statement and asks the operating system to delete it (exit)
  - Output data from child to parent (via wait)
  - Process' resources are deallocated by operating system
- Parent may terminate execution of children processes (abort)
  - Child has exceeded allocated resources
  - Task assigned to child is no longer required
  - If parent is exiting
    - ▶ Some operating system do not allow child to continue if its parent terminates
      - All children terminated - cascading termination



# fork(): Creating New Processes

## ■ `int fork(void)`

- creates a new process (child process) that is identical to the calling process (parent process)
- returns 0 to the child process
- returns child's pid to the parent process

```
pid_t pid = fork();
if (pid == 0) {
    printf("hello from child\n");
} else {
    printf("hello from parent\n");
}
```

- `fork()` is interesting (and often confusing) because it is called *once* but returns *twice*

# Understanding fork

## Process n



```
pid_t pid = fork();
if (pid == 0) {
    printf("hello from child\n");
} else {
    printf("hello from parent\n");
}
```



pid = m

```
pid_t pid = fork();
if (pid == 0) {
    printf("hello from child\n");
} else {
    printf("hello from parent\n");
}
```



```
pid_t pid = fork();
if (pid == 0) {
    printf("hello from child\n");
} else {
    printf("hello from parent\n");
}
```

hello from parent

## Child Process m



```
pid_t pid = fork();
if (pid == 0) {
    printf("hello from child\n");
} else {
    printf("hello from parent\n");
}
```



pid = 0

```
pid_t pid = fork();
if (pid == 0) {
    printf("hello from child\n");
} else {
    printf("hello from parent\n");
}
```



```
pid_t pid = fork();
if (pid == 0) {
    printf("hello from child\n");
} else {
    printf("hello from parent\n");
}
```

hello from child

*Which one is first?*

# fork Example #1


- Parent and child both run same code
  - Distinguish parent from child by return value from fork
- Start with same state, but each has private copy
  - Including shared output file descriptor
  - Relative ordering of their print statements undefined

```
void fork1()  
{  
    int x = 1;  
    pid_t pid = fork();  
    if (pid == 0) {  
        printf("Child has x = %d\n", ++x);  
    } else {  
        printf("Parent has x = %d\n", --x);  
    }  
    printf("Bye from process %d with x = %d\n", getpid(), x);  
}
```

# fork Example #2

- Both parent and child can continue forking


```
void fork2()  
{  
    printf("L0\n");  
    fork();  
    printf("L1\n");  
    fork();  
    printf("Bye\n");  
}
```



# fork Example #3

- Both parent and child can continue forking


```
void fork3()
{
    printf("L0\n");
    fork();
    printf("L1\n");
    fork();
    printf("L2\n");
    fork();
    printf("Bye\n");
}
```



# fork Example #4

- Both parent and child can continue forking


```
void fork4()  
{  
    printf("L0\n");  
    if (fork() != 0) {  
        printf("L1\n");  
        if (fork() != 0) {  
            printf("L2\n");  
            fork();  
        }  
    }  
    printf("Bye\n");  
}
```



# fork Example #5

- Both parent and child can continue forking

```
void fork5()  
{  
    printf("L0\n");  
    if (fork() == 0) {  
        printf("L1\n");  
        if (fork() == 0) {  
            printf("L2\n");  
            fork();  
        }  
    }  
    printf("Bye\n");  
}
```



# fork and Virtual Memory

- VM and memory mapping explain how fork provides private address space for each process.
- To create virtual address for new process
  - Create exact copies of current `mm_struct`, `vm_area_struct`, and page tables.
  - Flag each page in both processes as read-only
  - Flag each `vm_area_struct` in both processes as private COW
- On return, each process has exact copy of virtual memory
- Subsequent writes create new pages using COW mechanism.



# exit(): Ending a process

## ■ `void exit(int status)`

- exits a process
  - ▶ Normally return with status 0
- `atexit()` registers functions to be executed upon exit

```
void cleanup(void) {  
    printf("cleaning up\n");  
}  
  
void fork6() {  
    atexit(cleanup);  
    fork();  
    exit(0);  
}
```

# Zombies

## ■ Idea

- When process terminates, still consumes system resources
  - ▶ Various tables maintained by OS
- Called a “zombie”
  - ▶ Living corpse, half alive and half dead

## ■ Reaping

- Performed by parent on terminated child
- Parent is given exit status information
- Kernel discards process

## ■ What if parent doesn't reap?

- If any parent terminates without reaping a child, then child will be reaped by the init process
- So, only need explicit reaping in long-running processes
  - ▶ e.g., shells and servers

# Zombie Example

```
linux> ./forks 7 &
[1] 6639
Running Parent, PID = 6639
Terminating Child, PID = 6640
linux> ps
  PID TTY          TIME CMD
 6585 ttyp9        00:00:00 tcsh
 6639 ttyp9        00:00:03 forks
 6640 ttyp9        00:00:00 forks <defunct>
 6641 ttyp9        00:00:00 ps
linux> kill 6639
[1]      Terminated
linux> ps
  PID TTY          TIME CMD
 6585 ttyp9        00:00:00 tcsh
 6642 ttyp9        00:00:00 ps
```

```
void fork7()
{
    if (fork() == 0) {
        /* Child */
        printf("Terminating Child, PID = %d\n",
            getpid());
        exit(0);
    } else {
        printf("Running Parent, PID = %d\n",
            getpid());
        while (1)
            ; /* Infinite loop */
    }
}
```

- ps shows child process as “defunct”
- Killing parent allows child to be reaped by init

# Nonterminating Child Example

```
void fork8()
{
    if (fork() == 0) {
        /* Child */
        printf("Running Child, PID = %d\n",
               getpid());
        while (1)
            ; /* Infinite loop */
    } else {
        printf("Terminating Parent, PID = %d\n",
               getpid());
        exit(0);
    }
}
```

```
linux> ./forks 8
Terminating Parent, PID = 6675
Running Child, PID = 6676
linux> ps
  PID TTY          TIME CMD
 6585 tttyp9        00:00:00 tcsh
 6676 tttyp9        00:00:06 forks
 6677 tttyp9        00:00:00 ps
linux> kill 6676
linux> ps
  PID TTY          TIME CMD
 6585 tttyp9        00:00:00 tcsh
 6678 tttyp9        00:00:00 ps
```

- Child process still active even though parent has terminated
- Must kill explicitly, or else will keep running indefinitely


# wait(): Synchronizing with Children

■ `int wait(int *child_status)`

- suspends current process until one of its children terminates
- return value is the pid of the child process that terminated
- if `child_status != NULL`, then the object it points to will be set to a status indicating why the child process terminated

# wait: Synchronizing with Children

```
void fork9() {  
    int child_status;  
  
    if (fork() == 0) {  
        printf("HC: hello from child\n");  
    }  
    else {  
        printf("HP: hello from parent\n");  
        wait(&child_status);  
        printf("CT: child has terminated\n");  
    }  
    printf("Bye\n");  
    exit();  
}
```



# wait() Example

- If multiple children completed, will take in arbitrary order
- Can use macros WIFEXITED and WEXITSTATUS to get information about exit status

```
void fork10()
{
    pid_t pid[N];
    int i;
    int child_status;
    for (i = 0; i < N; i++)
        if ((pid[i] = fork()) == 0)
            exit(100+i); /* Child */
    for (i = 0; i < N; i++) {
        pid_t wpid = wait(&child_status);
        if (WIFEXITED(child_status))
            printf("Child %d terminated with exit status %d\n",
                wpid, WEXITSTATUS(child_status));
        else
            printf("Child %d terminated abnormally\n", wpid);
    }
}
```

# waitpid(): Waiting for a Specific Process

■ `waitpid(pid, &status, options)`

- suspends current process until specific process terminates
- various options (see textbook)

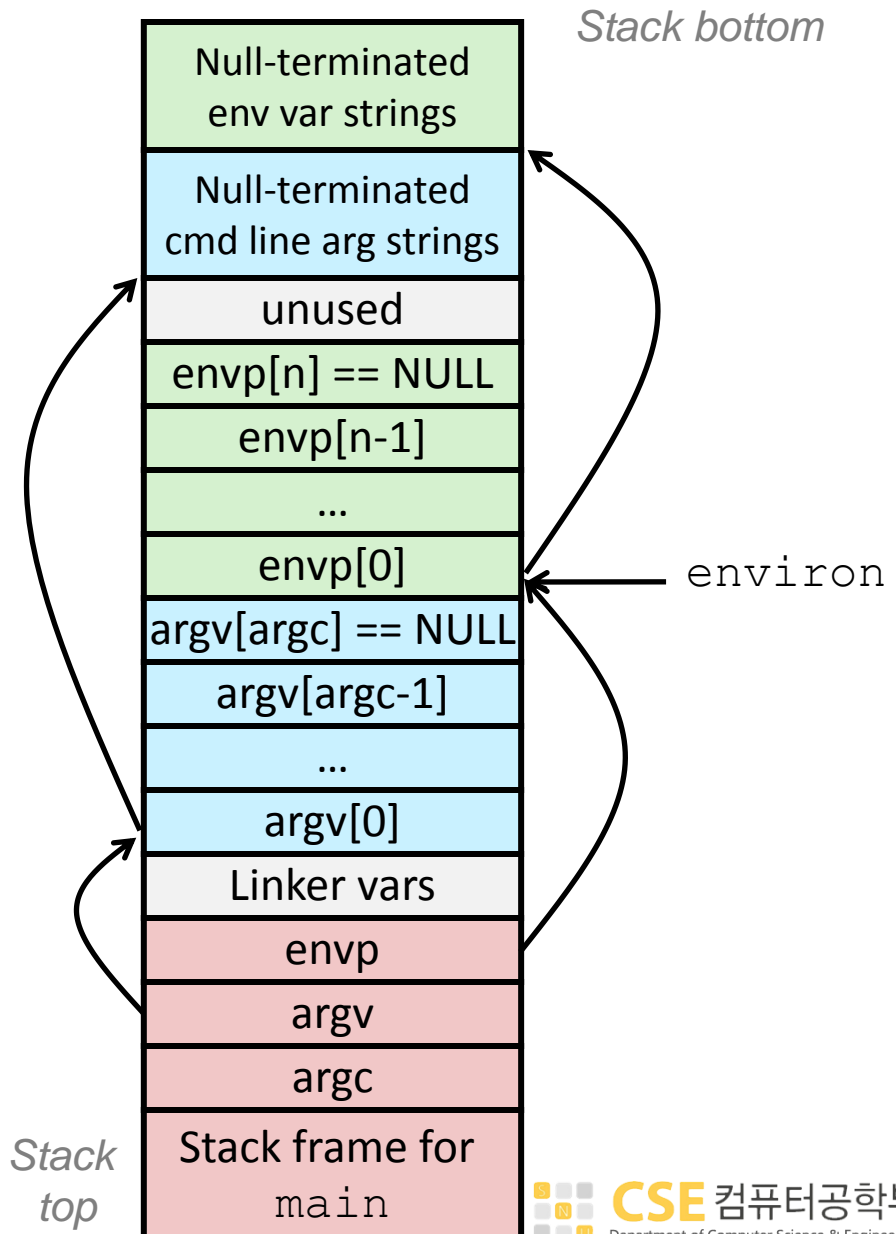
```
void fork11()
{
    pid_t pid[N];
    int i;
    int child_status;
    for (i = 0; i < N; i++)
        if ((pid[i] = fork()) == 0)
            exit(100+i); /* Child */
    for (i = N-1; i >= 0; i--) {
        pid_t wpid = waitpid(pid[i], &child_status, 0);
        if (WIFEXITED(child_status))
            printf("Child %d terminated with exit status %d\n",
                wpid, WEXITSTATUS(child_status));
        else
            printf("Child %d terminated abnormally\n", wpid);
    }
}
```



# execve(): Loading and Running Programs

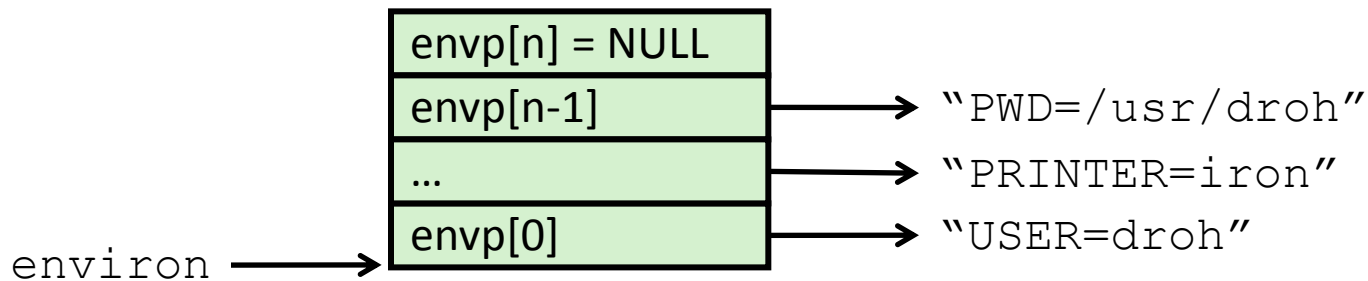
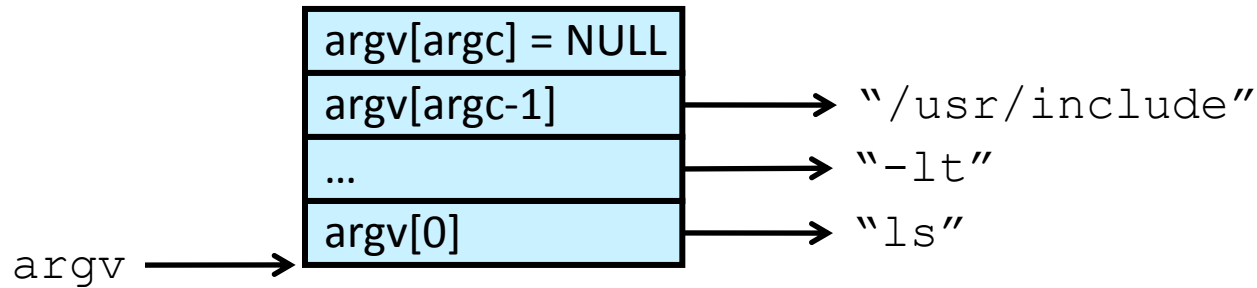
```
■ int execve(  
    char *filename,  
    char *argv[],  
    char *envp[]  
)
```

- Loads and runs in current process:
  - ▶ executable `filename`
  - ▶ argument list `argv`
  - ▶ environment variable list `envp`
- Does not return (unless error)
- Overwrites code, data, and stack
  - ▶ keeps pid, open files and signal context
- Environment variables:
  - ▶ “name=value” strings
  - ▶ `getenv` and `putenv`

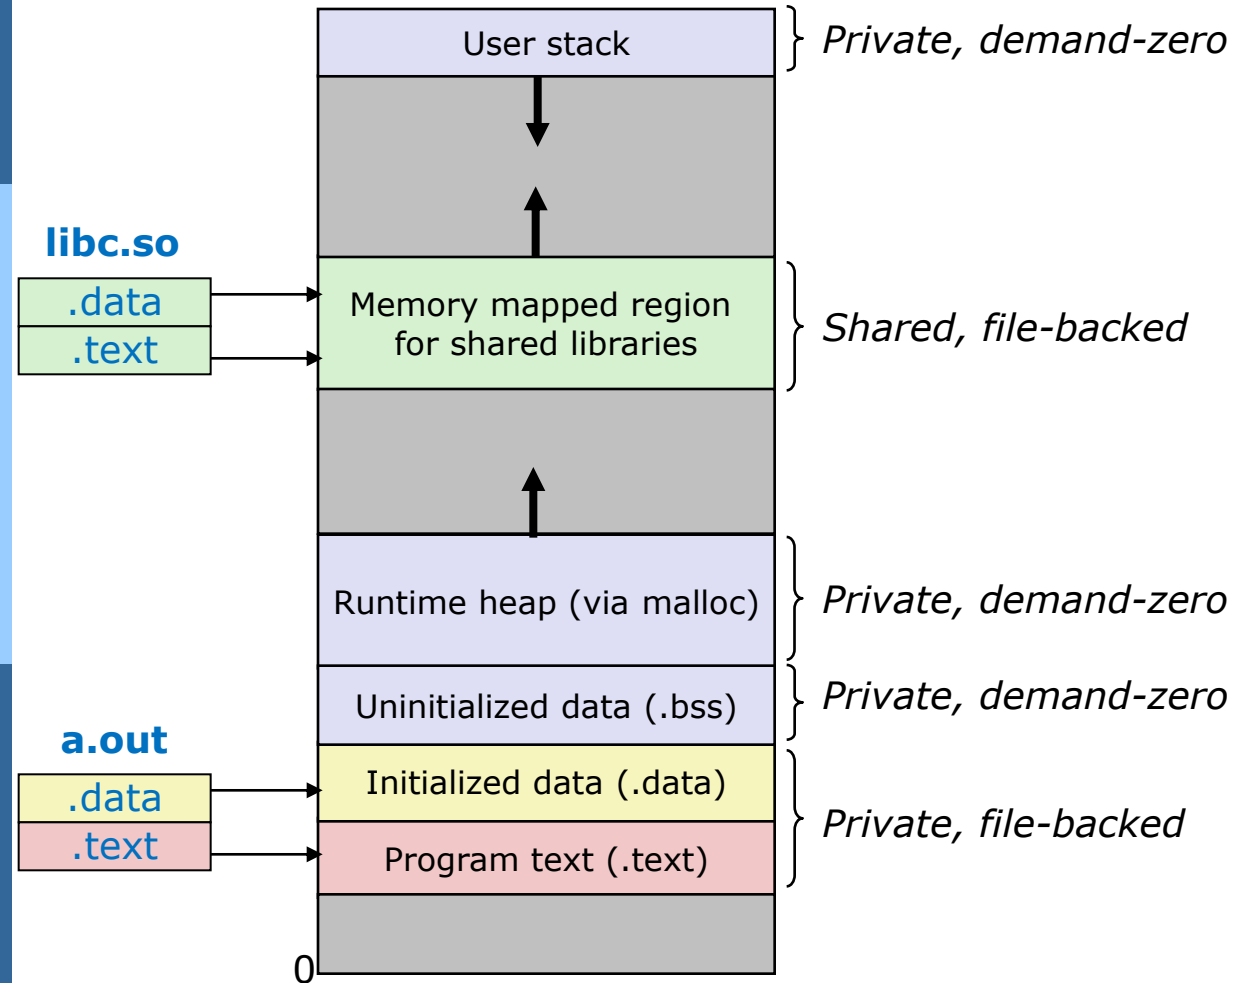


# execve Example

```
if ((pid = Fork()) == 0) { /* Child runs user job */  
    if (execve(argv[0], argv, environ) < 0) {  
        printf("%s: Command not found.\n", argv[0]);  
        exit(0);  
    }  
}
```



# execve and Virtual Memory



- To load and run a new program in the current process using `execve`:
- Free `vm_area_struct`'s and page tables for old areas
- Create `vm_area_struct`'s and page tables for new areas
  - Programs and initialized data backed by object files.
  - `.bss` and stack backed by anonymous files .
- Set PC to entry point in `.text`
  - Linux will fault in code and data pages as needed

# Signals

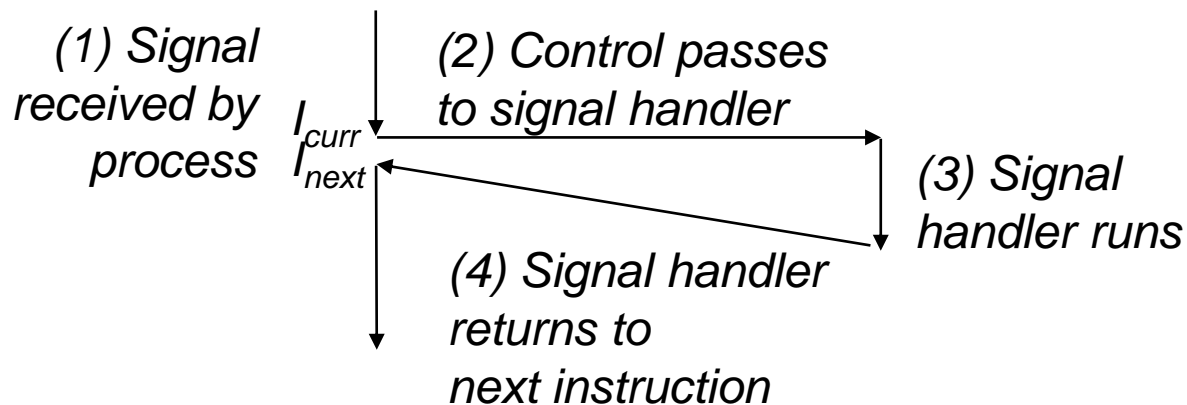
- higher-level form of exceptional control flow
- allows processes (and the kernel) to interrupt other processes
- different signals correspond to different events

#	Name	Default Action	Event
1	SIGHUP	terminate	terminal line hangup
2	SIGINT	terminate	interrupt from keyboard
3	SIGQUIT	terminate	quit from keyboard
4	SIGILL	terminate	illegal instruction
9	SIGKILL	terminate	kill program
11	SIGSEGV	terminate + coredump	segmentation fault (illegal memory ref)
14	SIGALRM	terminate	alarm() timer signal
15	SIGTERM	terminate	software termination signal
17	SIGCHLD	ignore	a child process has stopped or terminated
18	SIGCONT	ignore	continue process if stopped

→ see textbook fig. 8.25 for more details

# Sending and Receiving Signals

- Signals are sent by a program (or the kernel) and delivered by the kernel
  - `/bin/kill`: program to send arbitrary signals
  - keyboard: `ctrl-c`, `ctrl-z`
  - `int kill(pid_t, int sig)`
- The receiver can either ignore (block) the signal, terminate or catch the signal
- A signal that has been sent but not yet received is a *pending signal*



# Signal Handling Issues

- Subtle issues when catching multiple signals
  - pending signals are blocked
  - pending signals are not queued
  - system calls may be interrupted
- Example: simple shell
  - parent process creates some children
  - children run independently and then terminate
  - parent must reap children to avoid leaving zombies in the system
  - parent must keep running while the children are running
    - ▶ use SIGCHLD handler to reap the children

# Signal Handling Issues: Shell, Try 1

```
#include "csapp.h"

void handler1(int sig) {
    pid_t pid;

    if ((pid = waitpid(-1, NULL, 0)) < 0)
        unix_error("waitpid error");
    printf("Handler reaped child %d\n", (int)pid);
    Sleep(2);
    return;
}

int main() {
    int i, n;
    char buf[MAXBUF];

    if (signal(SIGCHLD, handler1) == SIG_ERR)
        unix_error("signal error");
```

```
/* Parent creates children */
for (i = 0; i < 3; i++) {
    if (Fork() == 0) {
        printf("Hello from child %d\n",
            (int)getpid());
        Sleep(1);
        exit(0);
    }
}

/* Parent waits for terminal input and then
processes it */
if ((n=read(STDIN_FILENO,buf,sizeof(buf))) < 0)
    unix_error("read");

printf("Parent processing input\n");
while (1)
    ;

exit(0);
}
```

→ signals are blocked and not queued

# Signal Handling Issues: Shell, Try 2

```
#include "csapp.h"

void handler2(int sig) {
    pid_t pid;

    while ((pid = waitpid(-1, NULL, 0)) > 0)
        printf("Handler reaped child %d\n",
            (int)pid);
    if (errno != ECHILD)
        unix_error("waitpid error");
    Sleep(2);
    return;
}

int main()
{
    int i, n;
    char buf[MAXBUF];

    if (signal(SIGCHLD, handler2) == SIG_ERR)
        unix_error("signal error");
```

```
/* Parent creates children */
for (i = 0; i < 3; i++) {
    if (Fork() == 0) {
        printf("Hello from child %d\n",
            (int)getpid());
        Sleep(1);
        exit(0);
    }
}

/* Parent waits for terminal input and then
processes it */
if ((n=read(STDIN_FILENO,buf,sizeof(buf))) < 0)
    unix_error("read");

printf("Parent processing input\n");
while (1)
    ;

    exit(0);
}
```

→ system calls can be interrupted



# Signal Handling Issues: Shell, Try 3

```
#include "csapp.h"

void handler2(int sig) {
    pid_t pid;

    while ((pid = waitpid(-1, NULL, 0)) > 0)
        printf("Handler reaped child %d\n",
            (int)pid);
    if (errno != ECHILD)
        unix_error("waitpid error");
    Sleep(2);
    return;
}

int main()
{
    int i, n;
    char buf[MAXBUF];

    if (signal(SIGCHLD, handler2) == SIG_ERR)
        unix_error("signal error");
```

```
/* Parent creates children */
for (i = 0; i < 3; i++) {
    if (Fork() == 0) {
        printf("Hello from child %d\n",
            (int)getpid());
        Sleep(1);
        exit(0);
    }
}

/* Manually restart the read call if it is
interrupted */
while ((n = read(STDIN_FILENO, buf,
    sizeof(buf))) < 0)
    if (errno != EINTR)
        unix_error("read error");

printf("Parent processing input\n");
while (1)
    ;

exit(0);
}
```

# Signal Handling

- Portable signal handling
  - semantics differ from system to system
    - ▶ it's a big mess
  - portable Posix standard:


```
int sigaction(int signum, struct sigaction *act, struct sigaction *oldact);
```

- Read textbook chapter 8.5 carefully


# Interprocess Communication

- Processes within a system may be independent or cooperating
- Cooperating process can affect or be affected by other processes, including sharing data
- Reasons for cooperating processes:
  - Information sharing
  - Computation speedup
  - Modularity
  - Convenience
- Cooperating processes need interprocess communication (IPC)
- Two models of IPC
  - Shared memory
  - Message passing

# Communications Models



(a)



(b)

# Producer-Consumer Problem

- Paradigm for cooperating processes, producer process produces information that is consumed by a consumer process
  - unbounded-buffer places no practical limit on the size of the buffer
  - bounded-buffer assumes that there is a fixed buffer size

# Bounded-Buffer – Shared-Memory Solution

- Shared data

```
#define BUFFER_SIZE 10
typedef struct {
    . . .
} item;

item buffer[BUFFER_SIZE];
int in = 0;
int out = 0;
```

- empty:  $in == out$ , full:  $((in+1) \% BUFFER\_SIZE) == out$
- Solution is correct, but can only use  $BUFFER\_SIZE-1$  elements

# Bounded-Buffer – Producer

```
while (true) {  
    /* Produce an item */  
    while (((in = (in + 1) % BUFFER SIZE count) == out)  
        ;    /* do nothing -- no free buffers */  
  
    buffer[in] = item;  
    in = (in + 1) % BUFFER SIZE;  
}
```

# Bounded Buffer – Consumer

```
while (true) {  
    while (in == out)  
        ; // do nothing -- nothing to consume  
  
    // remove an item from the buffer  
    item = buffer[out];  
    out = (out + 1) % BUFFER SIZE;  
    return item;  
}
```



# Interprocess Communication – Message Passing

- Mechanism for processes to communicate and to synchronize their actions
- Message system – processes communicate with each other without resorting to shared variables
- IPC facility provides two operations:
  - send(message) – message size fixed or variable
  - receive(message)
- If P and Q wish to communicate, they need to:
  - establish a communication link between them
  - exchange messages via send/receive
- Implementation of communication link
  - physical (e.g., shared memory, hardware bus)
  - logical (e.g., logical properties)

# Implementation Questions

- How are links established?
- Can a link be associated with more than two processes?
- How many links can there be between every pair of communicating processes?
- What is the capacity of a link?
- Is the size of a message that the link can accommodate fixed or variable?
- Is a link unidirectional or bi-directional?

# Direct Communication

- Processes must name each other explicitly:
  - `send (P, message)` – send a message to process P
  - `receive(Q, message)` – receive a message from process Q
- Properties of communication link
  - Links are established automatically
  - A link is associated with exactly one pair of communicating processes
  - Between each pair there exists exactly one link
  - The link may be unidirectional, but is usually bi-directional

# Indirect Communication

- Messages are directed and received from mailboxes (also referred to as ports)
  - Each mailbox has a unique id
  - Processes can communicate only if they share a mailbox
- Properties of communication link
  - Link established only if processes share a common mailbox
  - A link may be associated with many processes
  - Each pair of processes may share several communication links
  - Link may be unidirectional or bi-directional

# Synchronization

- Message passing may be either blocking or non-blocking
- Blocking is considered synchronous
  - Blocking send has the sender block until the message is received
  - Blocking receive has the receiver block until a message is available
- Non-blocking is considered asynchronous
  - Non-blocking send has the sender send the message and continue
  - Non-blocking receive has the receiver receive a valid message or null

# Buffering

- Queue of messages attached to the link; implemented in one of three ways
  1. Zero capacity – 0 messages  
Sender must wait for receiver (rendezvous)
  2. Bounded capacity – finite length of n messages  
Sender must wait if link full
  3. Unbounded capacity – infinite length  
Sender never waits

# Examples of IPC Systems - POSIX

## ■ POSIX Shared Memory

- Process first creates shared memory segment  
segment id = `shmget(IPC_PRIVATE, size, S_IRUSR | S_IWUSR);`
- Process wanting access to that shared memory must attach to it  
`shared memory = (char *) shmat(id, NULL, 0);`
- Now the process could write to the shared memory  
`sprintf(shared memory, "Writing to shared memory");`
- When done a process can detach the shared memory from its address space  
`shmdt(shared memory);`

# Communications in Client-Server Systems

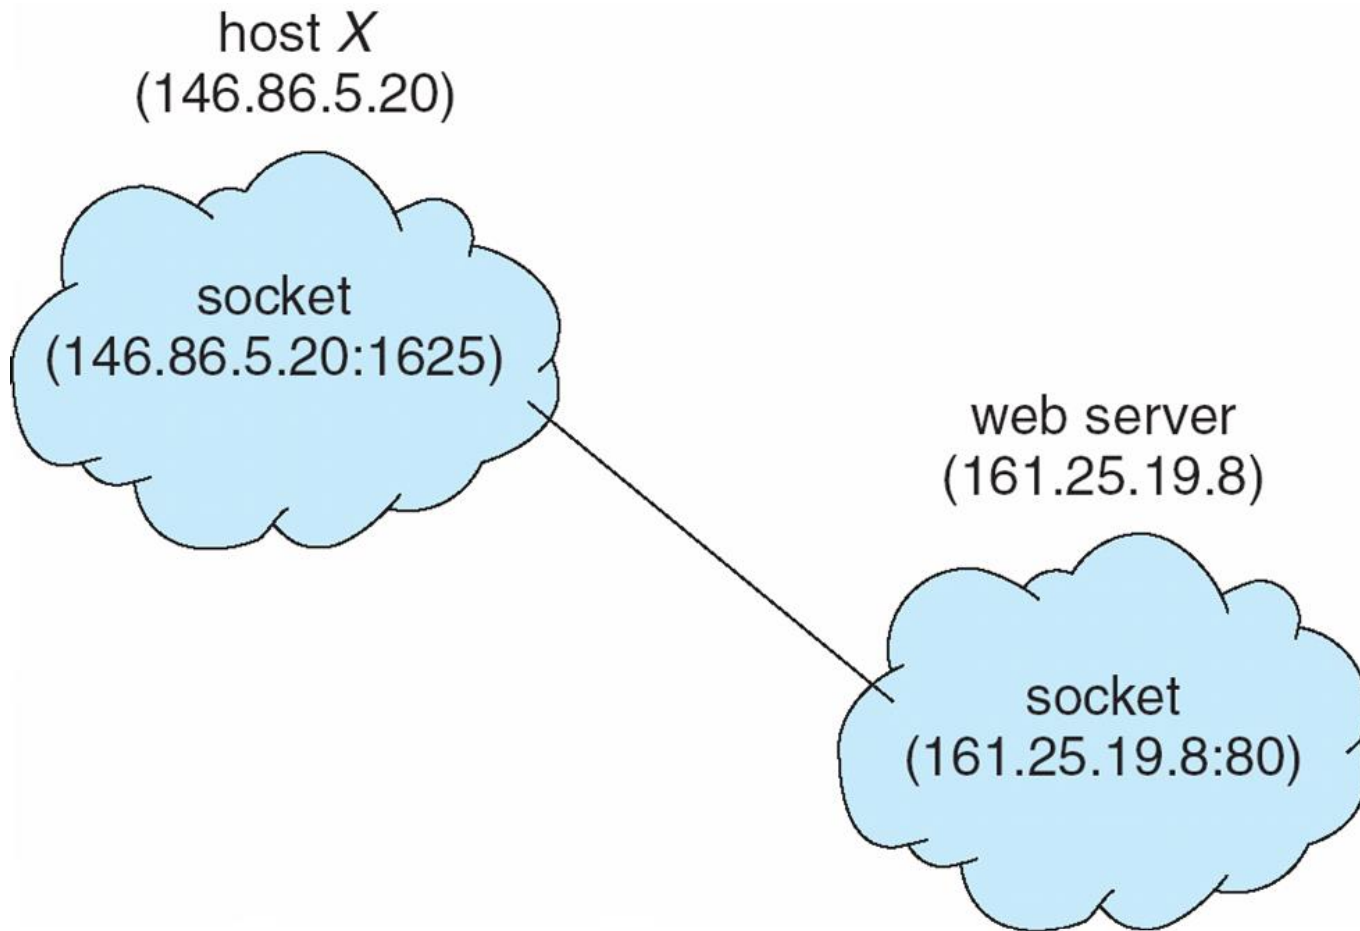
- Sockets
- Remote Procedure Calls  
(Remote Method Invocation (Java))
- Pipes



# Sockets

- A socket is defined as an endpoint for communication
- Concatenation of IP address and port
- The socket 161.25.19.8:1625 refers to port 1625 on host 161.25.19.8
- Communication consists between a pair of sockets

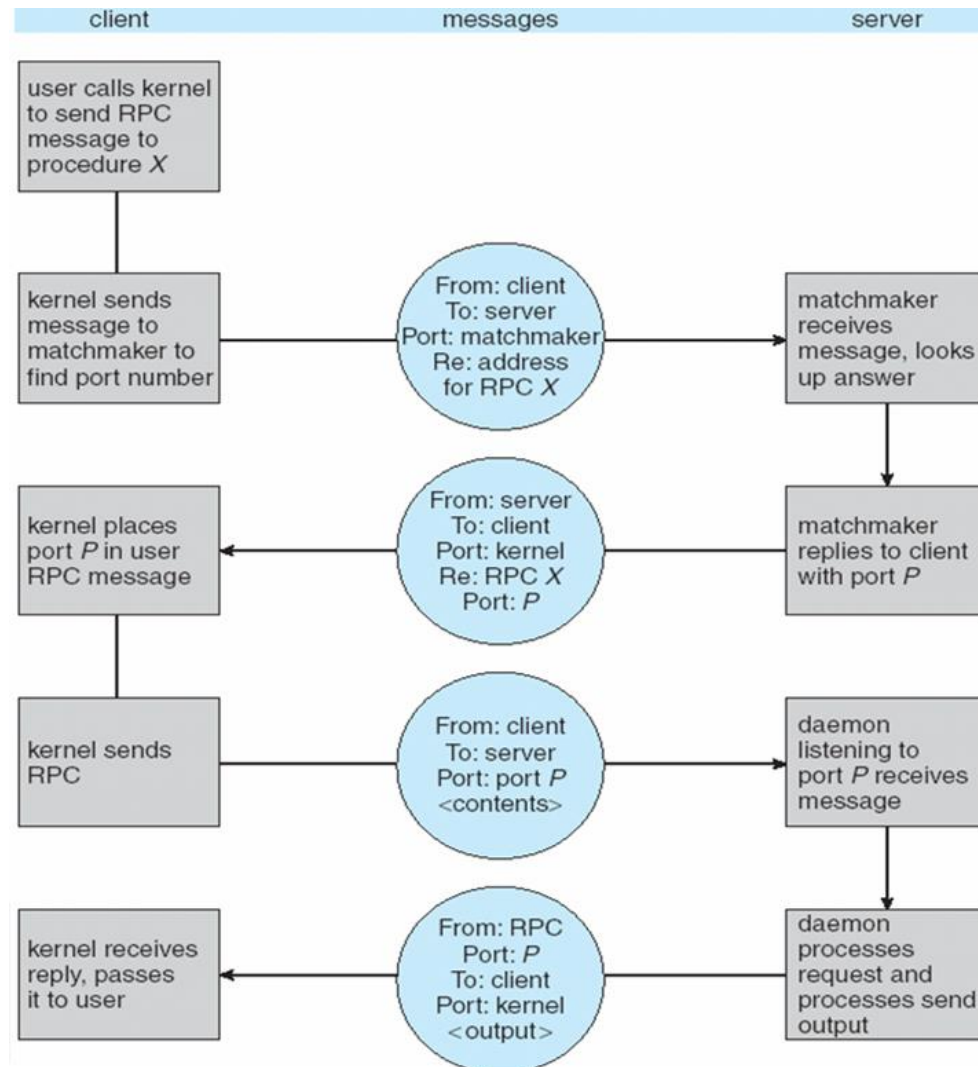
# Socket Communication



# Remote Procedure Calls


- Remote procedure call (RPC) abstracts procedure calls between processes on networked systems
- Stubs – client-side proxy for the actual procedure on the server
- The client-side stub locates the server and marshalls the parameters
- The server-side stub receives this message, unpacks the marshalled parameters, and performs the procedure on the server

# Execution of RPC

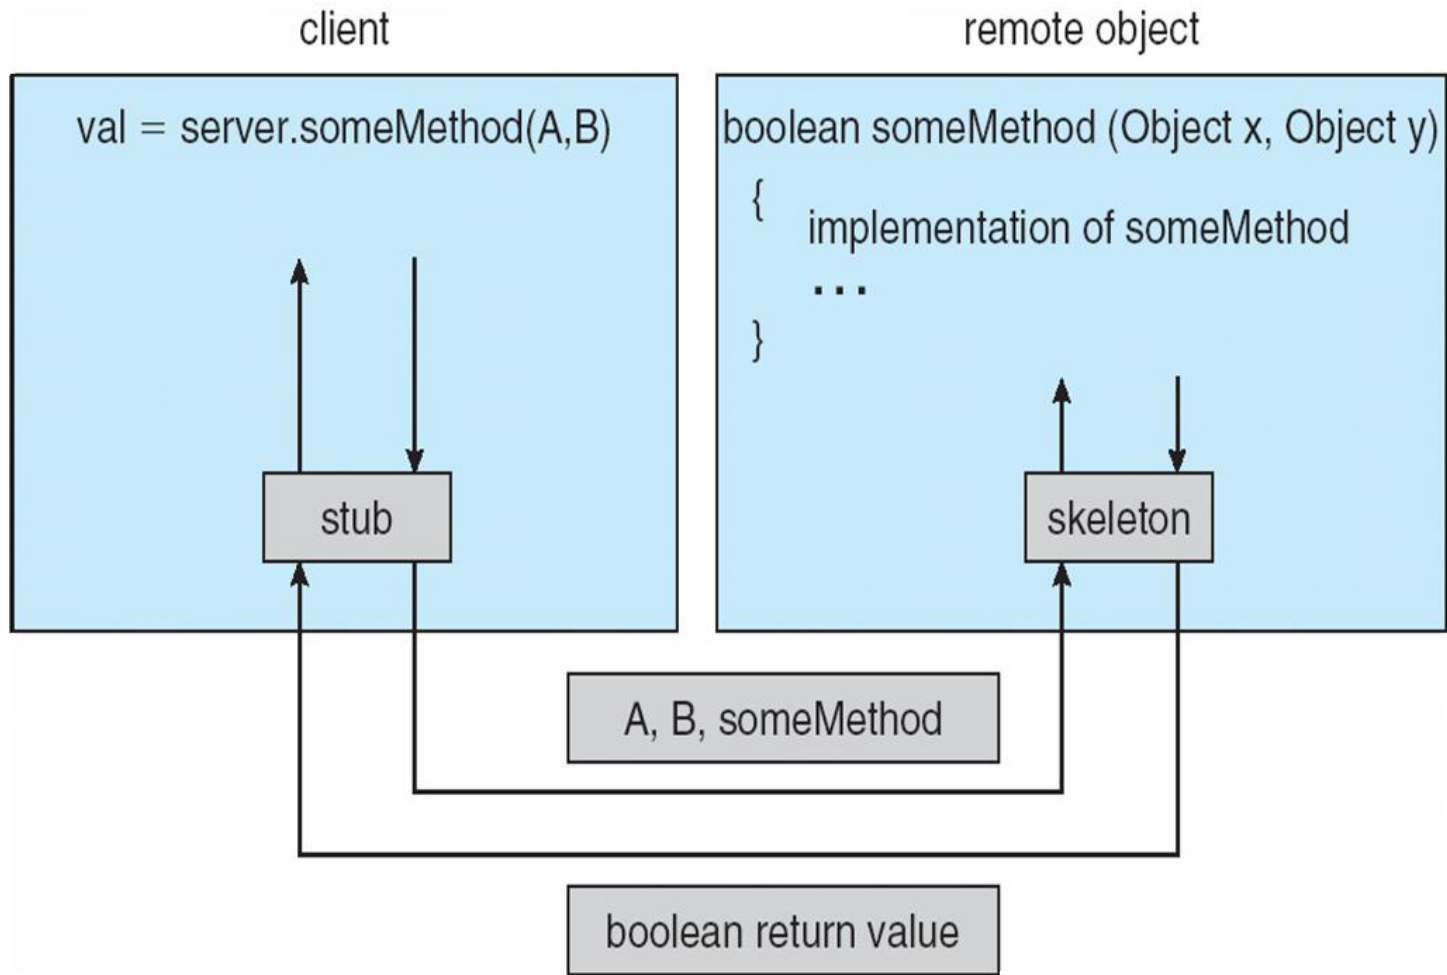


# Remote Method Invocation

- Remote Method Invocation (RMI) is a Java mechanism similar to RPCs
- RMI allows a Java program on one machine to invoke a method on a remote object



# Marshalling Parameters



# Pipes


- A pipe is a conduit allowing processes to communicate
- Unstructured contents ; structure imposed by communicating processes
- Issues to be considered
  - unidirectional/bidirectional?
  - for bidirectional pipes: half-duplex or full-duplex?
  - relationship between processes
  - local machine only or network transport possible?

# Ordinary Pipes

- In UNIX, pipe descriptors are file descriptors
  - use `read()/write()` to communicate, `close()` to close the pipe
  - `pipe(int fd[])` opens a pipe (`fd[0]`: read, `fd[1]`: write)
- In Windows, ordinary pipes are called anonymous pipes
  - same concept as in UNIX (file handle, use standard read/write calls)
- ordinary pipes are
  - unidirectional
  - cannot be accessed from outside the process that creates it
    - requires a parent-child relationship for communicating processes
    - only local communication possible



# UNIX Pipes



# Named Pipes

- Named pipes
  - typically “live” in the file system
  - can be bidirectional, no parent-child relationship required
  - persistent
  - UNIX:
    - ▶ local machine only
    - ▶ half-duplex bidirectional
  - Windows:
    - ▶ local/remote communication
    - ▶ full-duplex bidirectional

# Summary

## ■ Processes

- At any given time, system has multiple active processes
- Only one can execute at a time on a single core, though
- Each process appears to have total control of processor + private memory space
- Spawning processes: `fork()`
  - ▶ One call, two returns
- Process completion: `exit()`
  - ▶ One call, no return
- Reaping and waiting for Processes: `wait()`, `waitpid()`
- Loading and running Programs: `execve()` and variants
  - ▶ One call, no return (unless an error occurred)

# Summary (cont.)

## ■ Signals

- slow & limited way to send events between related processes

## ■ Interprocess communication

- faster, less limitations
- shared vs. message passing
- many different ways to do it