

Interaction and Communication between Programs

Network Programming

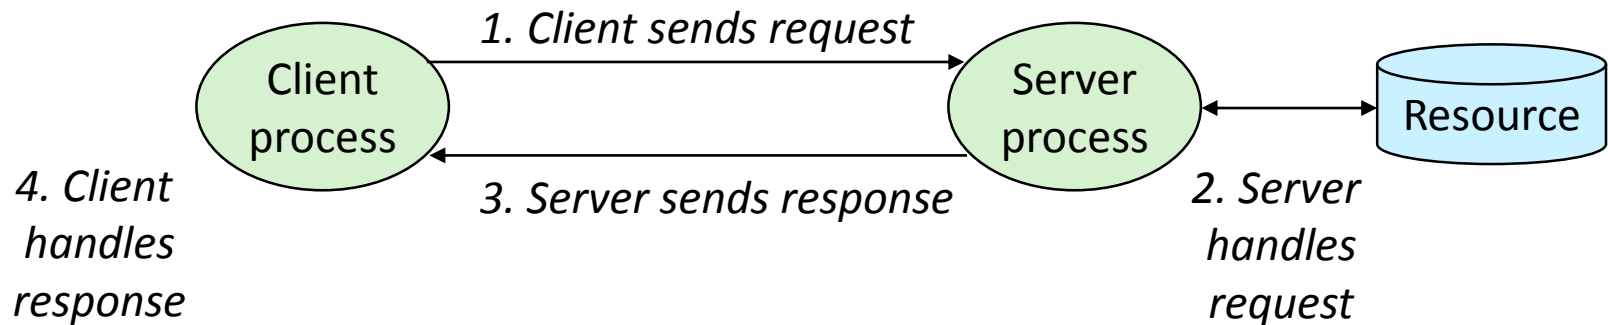


Network Programming

- **Client-Server Model and Computer Networks**
- Global IP Internet
 - naming
 - IPv4, IPv6
 - DNS
 - Internet connection
- Programmer's View
 - The sockets interface
 - Establishing an Internet connection
 - Copy data over an Internet connection

Acknowledgement: slides based on the cs:app2e material

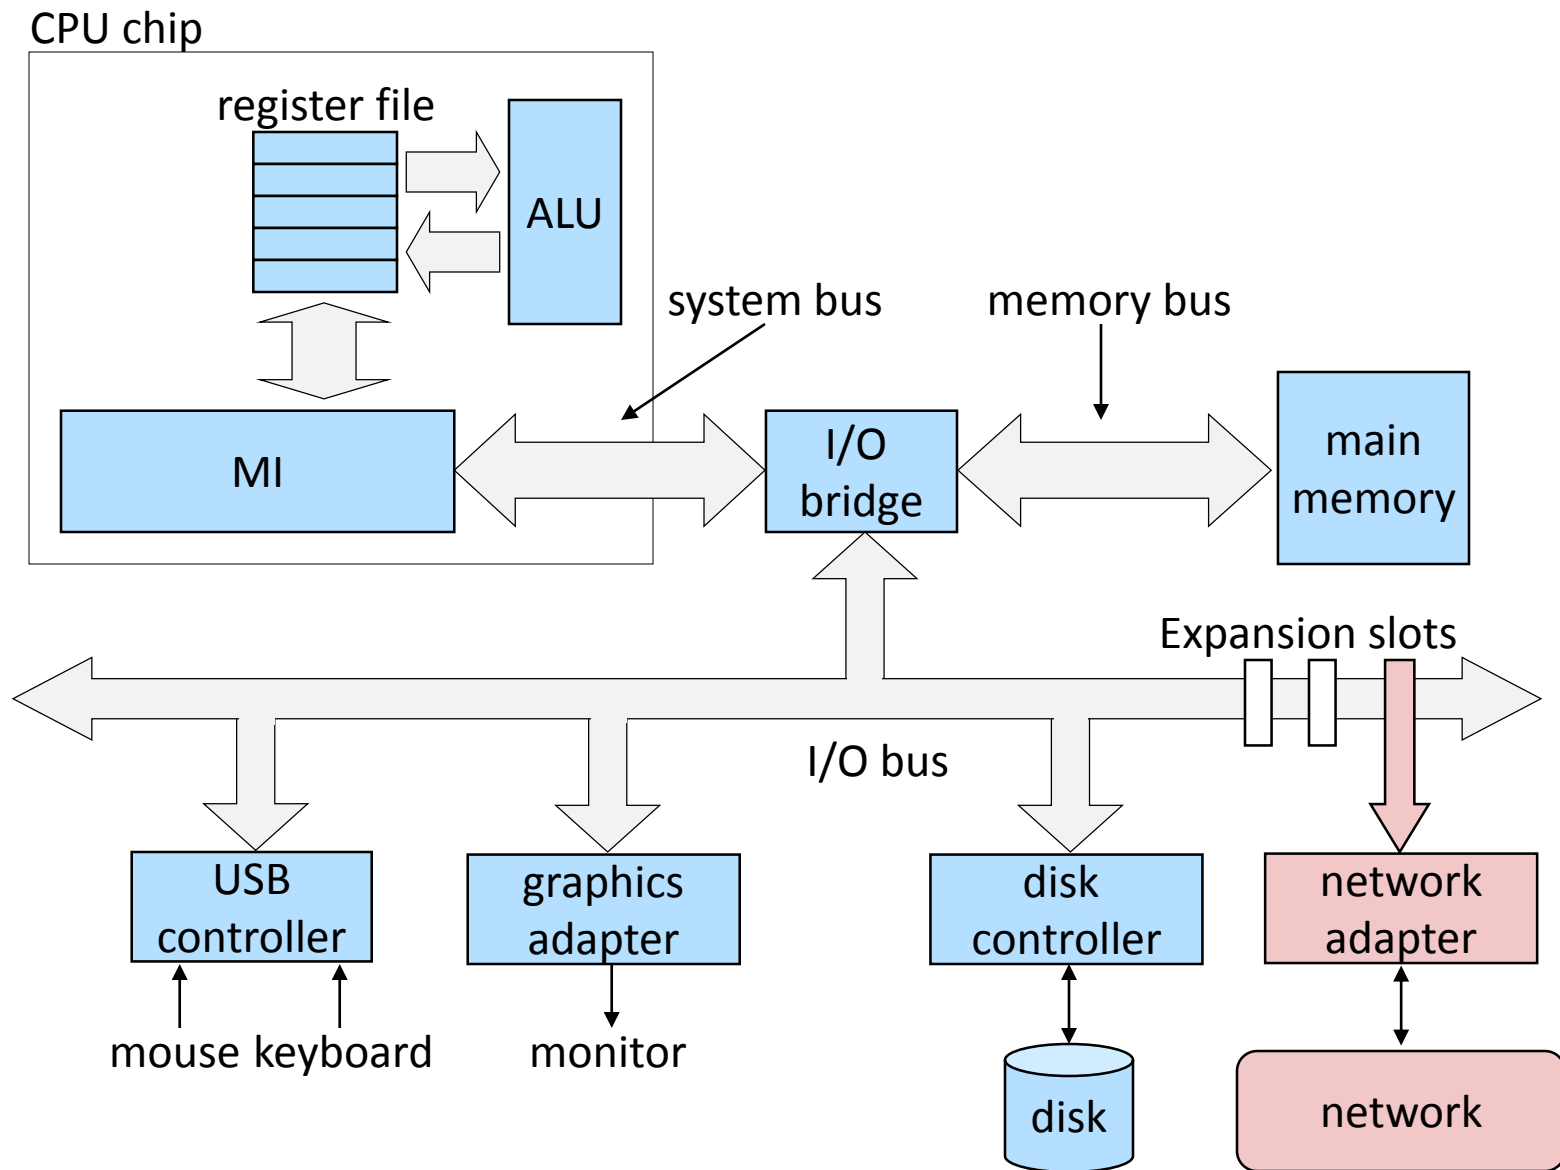
A Client-Server Transaction



Note: clients and servers are processes running on hosts (can be the same or different hosts)

- Most network applications are based on the client-server model:
 - A server process and one or more client processes
 - Server manages some resource
 - Server provides service by manipulating resource for clients
 - Server activated by request from client (vending machine analogy)


Hardware Organization of a Network Host



Computer Networks

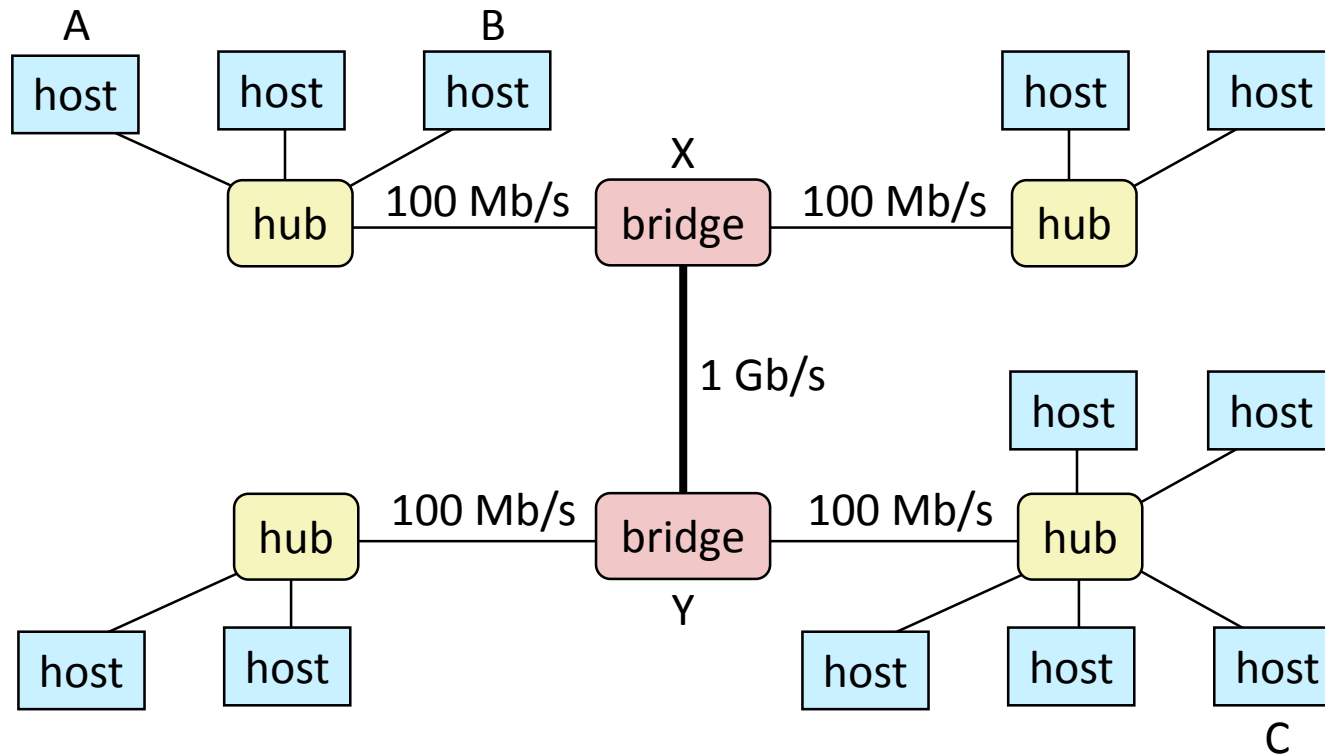
- A network is a hierarchical system of boxes and wires organized by geographical proximity
 - SAN (System Area Network) spans cluster or machine room
 - ▶ Switched Ethernet, Quadrics QSW, ...
 - LAN (Local Area Network) spans a building or campus
 - ▶ Ethernet is most prominent example
 - WAN (Wide Area Network) spans country or world
 - ▶ Typically high-speed point-to-point phone lines
- An internetwork (internet) is an interconnected set of networks
 - The Global IP Internet (uppercase “I”) is the most famous example of an internet (lowercase “i”)
- Let’s see how an internet is built from the ground up

Lowest Level: Ethernet Segment



- Ethernet segment consists of a collection of hosts connected by wires (twisted pairs) to a hub
- Spans room or floor in a building
- Operation
 - Each Ethernet adapter has a unique 48-bit address (MAC address)
 - ▶ E.g., 00:16:ea:e3:54:e6
 - Hosts send bits to any other host in chunks called frames
 - Hub slavishly copies each bit from each port to every other port
 - ▶ Every host sees every bit
 - ▶ Note: Hubs are on their way out. Bridges (switches, routers) became cheap enough to replace them (means no more broadcasting)


Next Level: Bridged Ethernet Segment



- Spans building or campus
- Bridges cleverly learn which hosts are reachable from which ports and then selectively copy frames from port to port


Conceptual View of LANs

- For simplicity, hubs, bridges, and wires are often shown as a collection of hosts attached to a single wire:




Next Level: internets

- Multiple incompatible LANs can be physically connected by specialized computers called routers
- The connected networks are called an internet



LAN 1 and LAN 2 might be completely different, totally incompatible (e.g., Ethernet and Wifi, 802.11, T1-links, DSL, ...)*

Logical Structure of an internet



- Ad hoc interconnection of networks
 - No particular topology
 - Vastly different router & link capacities
- Send packets from source to destination by hopping through networks
 - Router forms bridge from one network to another
 - Different packets may take different routes

The Notion of an internet Protocol

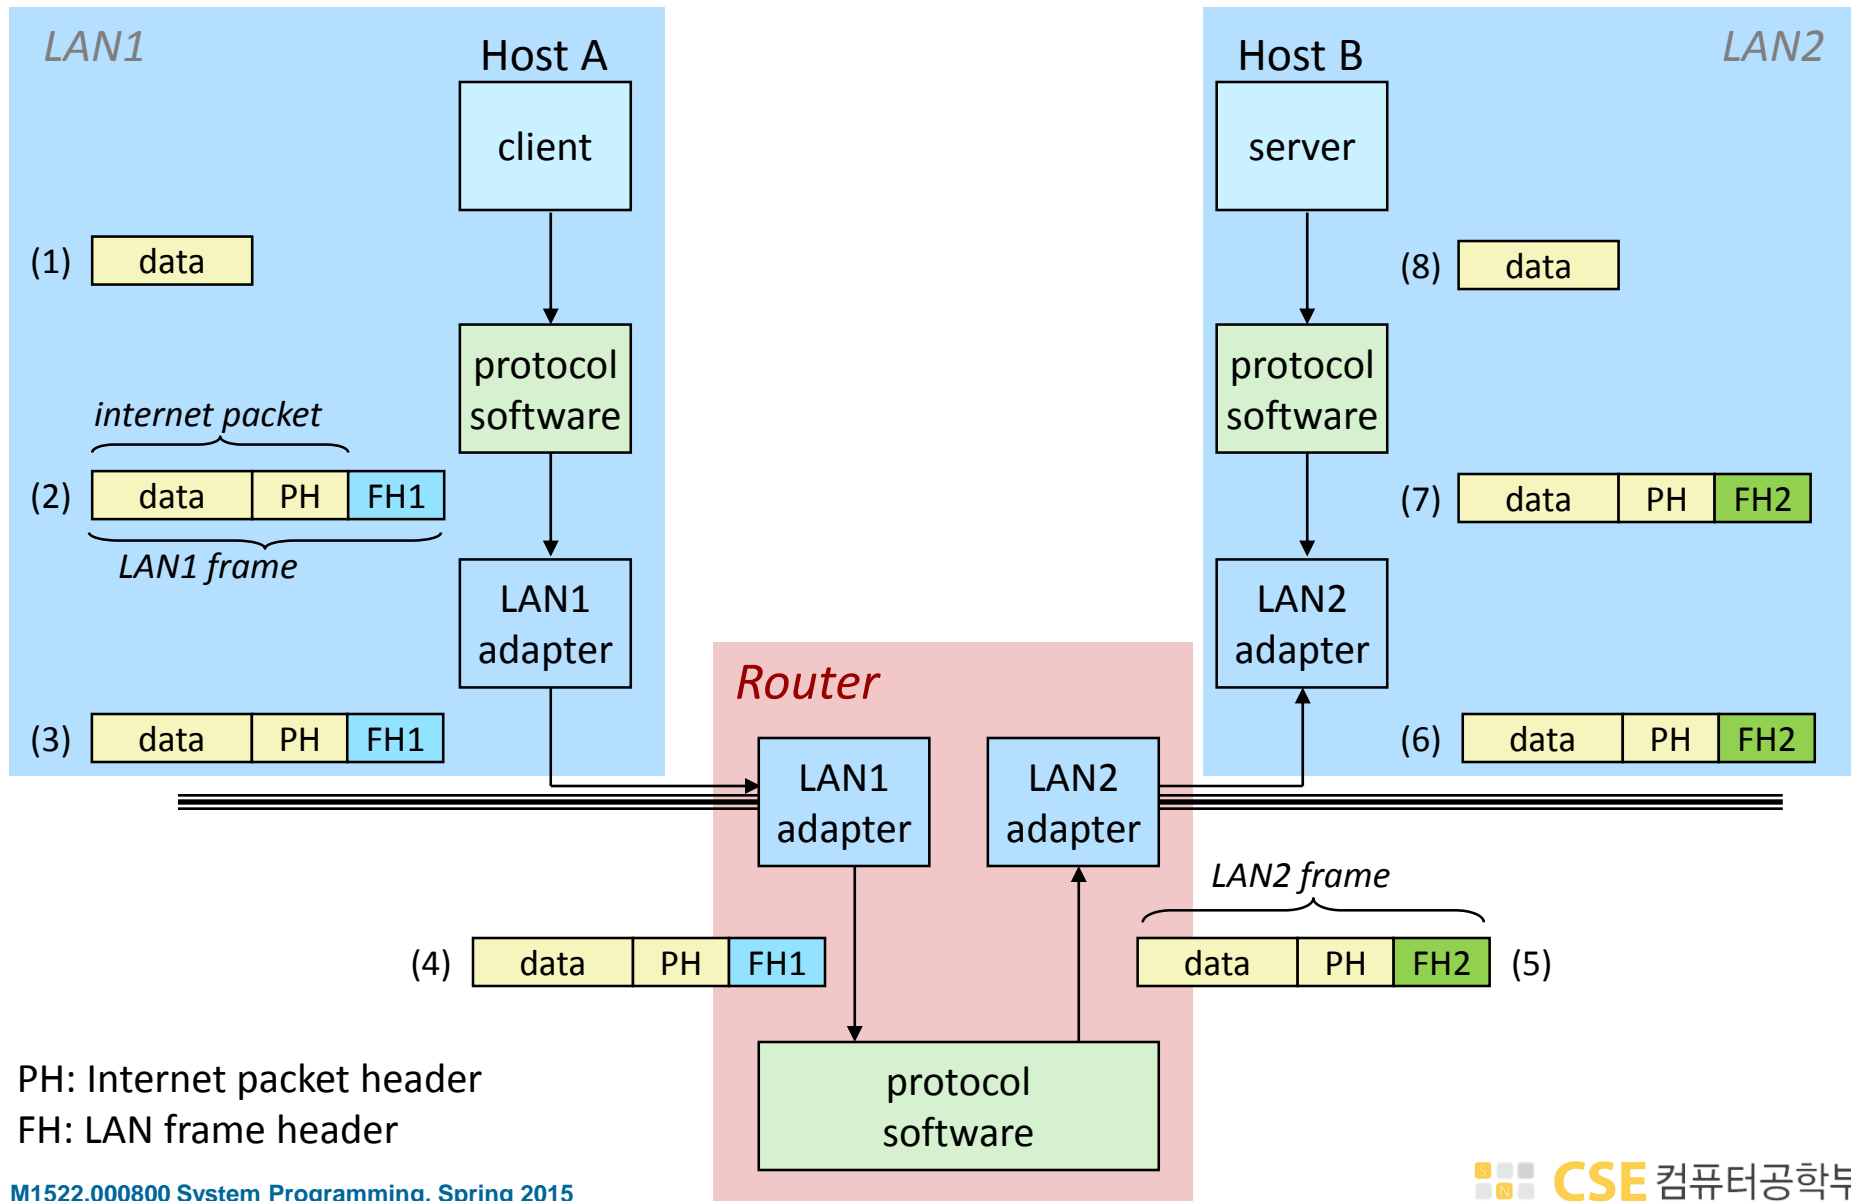
- How is it possible to send bits across incompatible LANs and WANs?
- Solution:
 - protocol software running on each host and router
 - smoothes out the differences between the different networks
- Implements an internet protocol (i.e., set of rules)
 - governs how hosts and routers should cooperate when they transfer data from network to network
 - TCP/IP is the protocol for the global IP Internet

What Does an internet Protocol Do?

- Provides a naming scheme
 - An internet protocol defines a uniform format for host addresses
 - Each host (and router) is assigned at least one of these internet addresses that uniquely identifies it

- Provides a delivery mechanism
 - An internet protocol defines a standard transfer unit (packet)
 - Packet consists of header and payload
 - ▶ Header: contains info such as packet size, source and destination addresses
 - ▶ Payload: contains data bits sent from source host

Transferring Data Over an internet



Other Issues

- We are glossing over a number of important questions:
 - What if different networks have different maximum frame sizes? (segmentation)
 - How do routers know where to forward frames?
 - How are routers informed when the network topology changes?
 - What if packets get lost?

- These (and other) questions are addressed by the area of systems known as computer networking

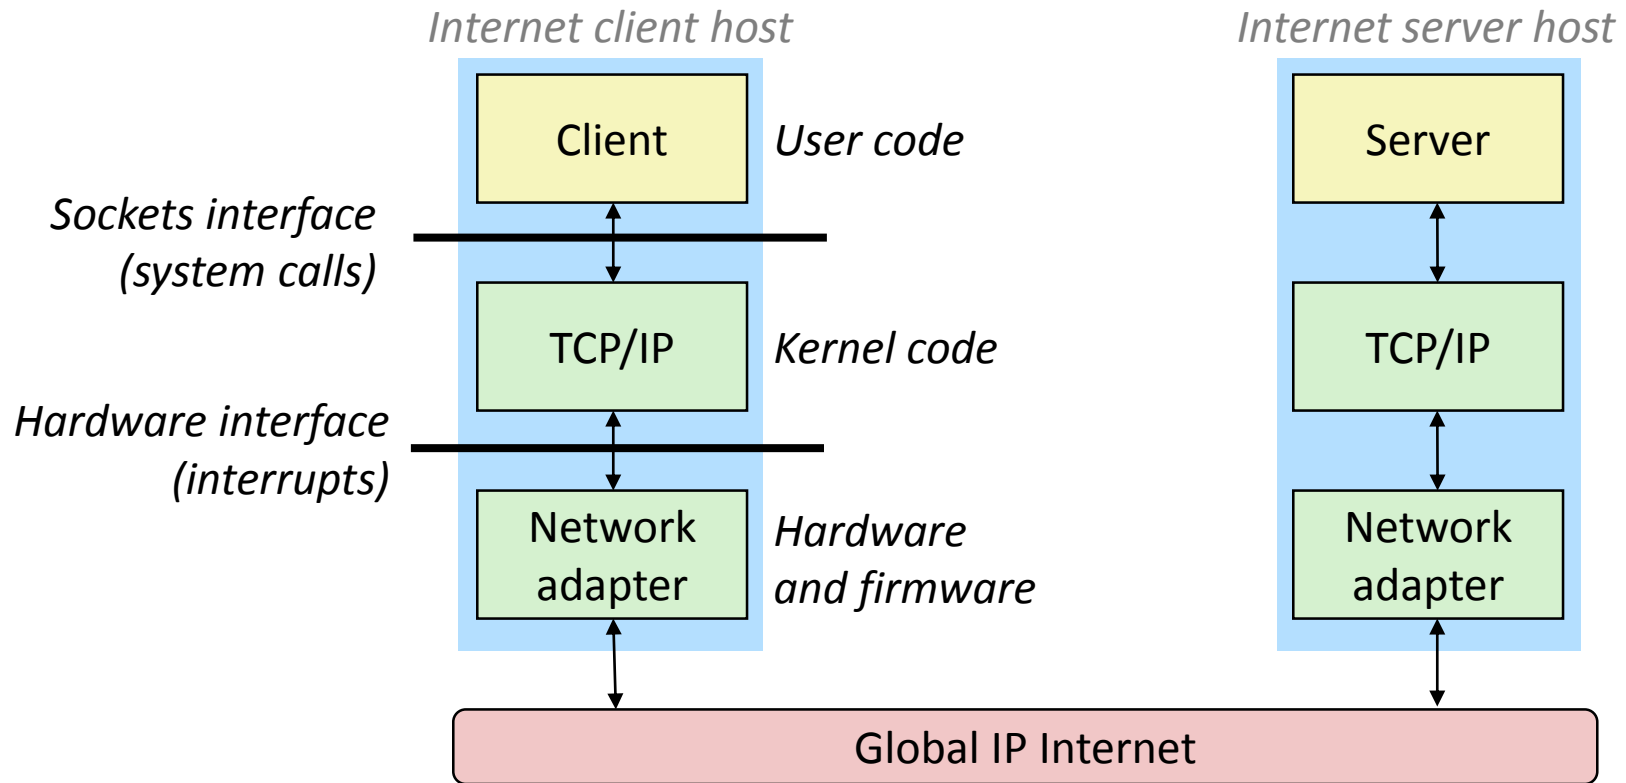
Network Programming

- Client-Server Model and Computer Networks
- **Global IP Internet**
 - naming
 - IPv4, IPv6
 - DNS
 - **Internet connection**
- Programmer's View
 - The sockets interface
 - Establishing an Internet connection
 - Copy data over an Internet connection

Global IP Internet

- Most famous example of an internet
- Based on the TCP/IP protocol family
 - IP (Internet protocol) :
 - ▶ Provides basic naming scheme and unreliable delivery capability of packets (datagrams) from host-to-host
 - UDP (Unreliable Datagram Protocol)
 - ▶ Uses IP to provide unreliable datagram delivery from process-to-process
 - TCP (Transmission Control Protocol)
 - ▶ Uses IP to provide reliable byte streams from process-to-process over connections
- Accessed via a mix of Unix file I/O and functions from the sockets interface

Hardware and Software Organization of an Internet Application



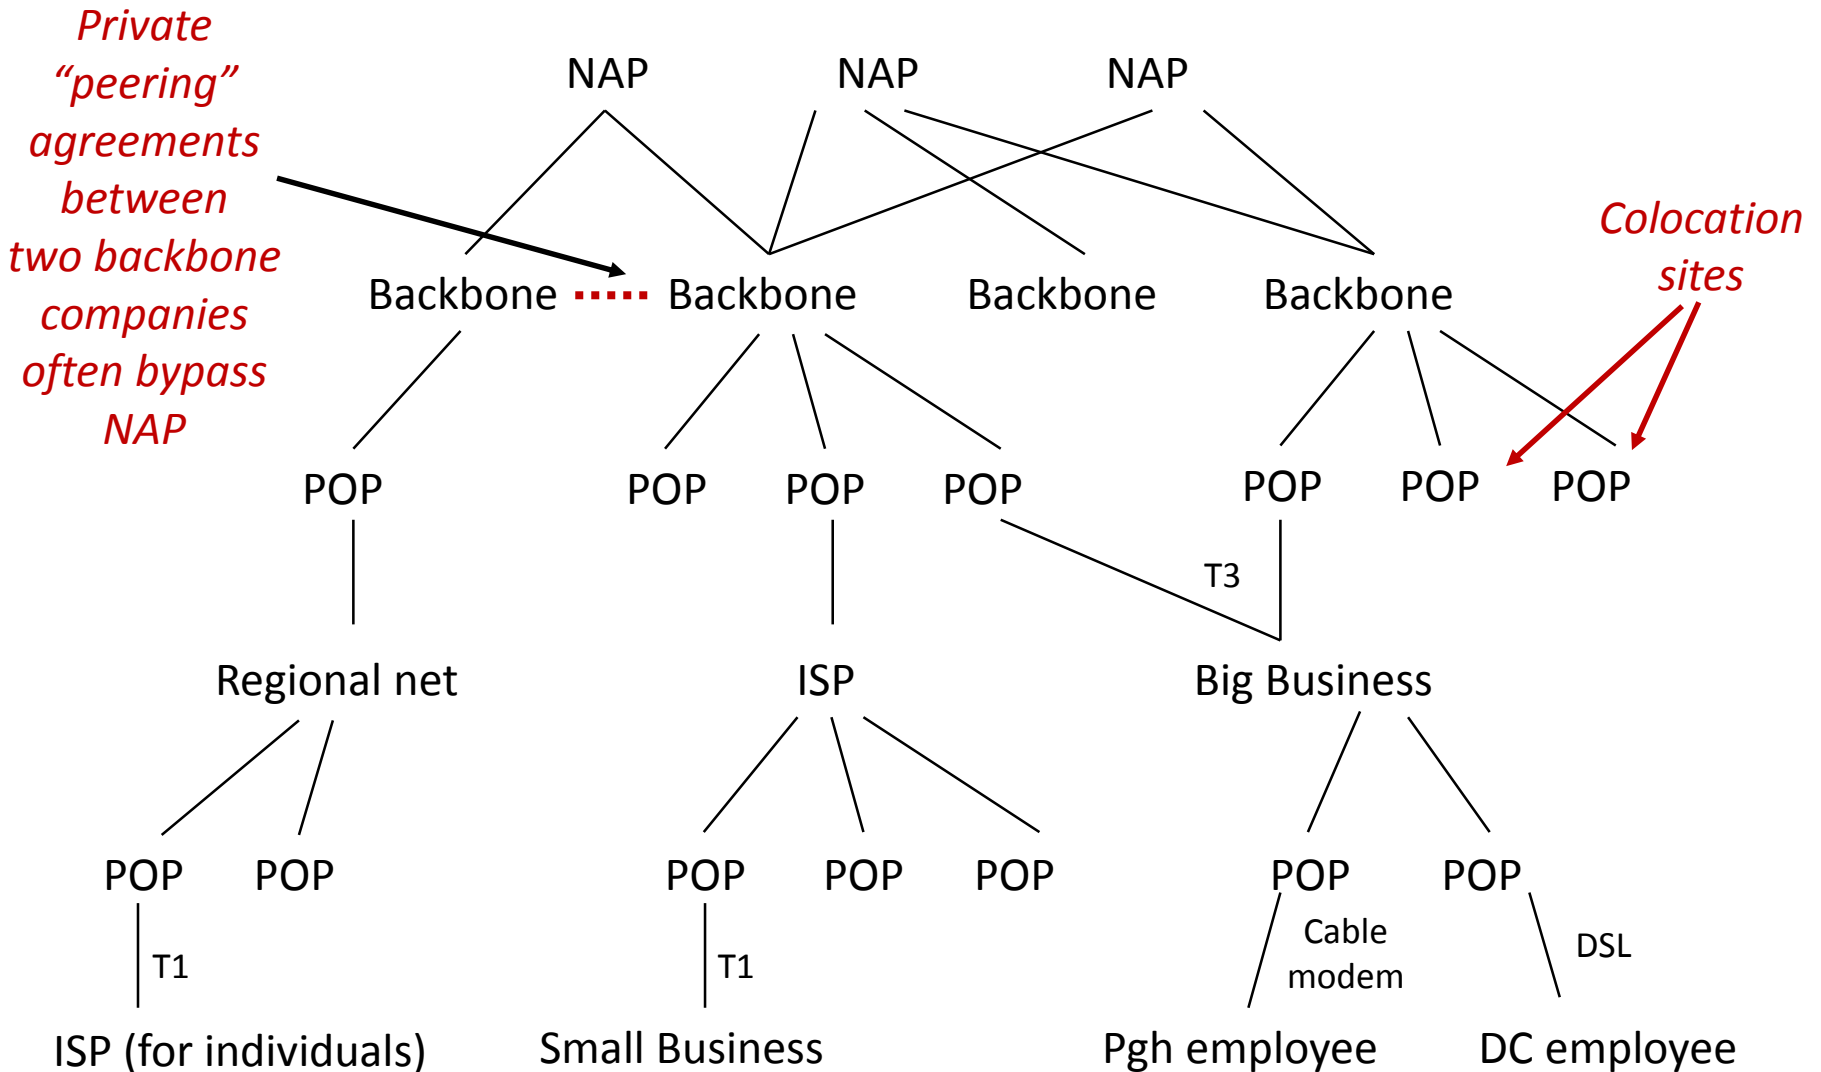
Basic Internet Components

- Internet backbone:
 - collection of routers (nationwide or worldwide) connected by high-speed point-to-point networks
- Network Access Point (NAP):
 - router that connects multiple backbones (often referred to as peers)
- Regional networks:
 - smaller backbones that cover smaller geographical areas (e.g., cities or states)
- Point of presence (POP):
 - machine that is connected to the Internet
- Internet Service Providers (ISPs):
 - provide dial-up or direct access to POPs

NAP-Based Internet Architecture

- NAPs link together commercial backbones provided by companies such as AT&T and Worldcom
- Currently in the US there are about 50 commercial backbones connected by ~12 NAPs (peering points)
- Similar architecture worldwide connects national networks to the Internet

Internet Connection Hierarchy



1¢ 14¢
FIXED-LINE MOBILE

Per-minute wholesale cost to send an international call to Switzerland.

27%

TRAFFIC DENT AS VOIP
Share of total international voice traffic transported as VoIP by carriers in 2009.

4/5

WIRELESS WORLD
Share of phones worldwide that are mobile.

12%

SKYPE VS. THE WORLD
Total international Skype traffic as a percentage of world-wide international voice traffic in 2009, approximately 33 billion minutes.

0.1 4.1
FIXED-LINE MOBILE


Billions of new subscribers added since 2000.

GLOBAL TRAFFIC MAP 2010

TeleGeography **ROGERS™**




MINUTES & VOLUME




REGIONAL TRAFFIC 2009 (in %)




CALL VOLUME & GROWTH



PRICING & REVENUES




RATES OF PRICE INDEX, VOLUME, AND GROWTH




WORLDWIDE TRAFFIC AND REVENUE BY DESTINATION REGION 2009




USERS & TECHNOLOGY



ROUTES & CONNECTIONS



DISSEMINATION BY TECHNOLOGY TYPE



INTERNATIONAL TRAFFIC FLOW 2009



Source: <https://www.telegeography.com/telecom-maps/global-traffic-map.1.html>

History and Forecast of Internet Traffic

Year	Global Internet Traffic	4TB HardDisks needed
1992	100 GB / day	1 every 40 days
1997	100 GB / hour	1 every 2 days
2002	100 GB / second	2'160 per day
2007	2000 GB / second	30 per minute
2013	28'875 GB / second	428 per minute
2018	50,000 GB / second	750 per minute

Source: Cisco VNI, 2014

Naming and Communicating on the Internet

■ Original Idea

- Every node on Internet would have unique IP address
 - ▶ Everyone would be able to talk directly to everyone
- No secrecy or authentication
 - ▶ Messages visible to routers and hosts on same LAN
 - ▶ Possible to forge source field in packet header

■ Shortcomings


- There aren't enough IP addresses available
- Don't want everyone to have access or knowledge of all other hosts
- Security issues mandate secrecy & authentication

Evolution of Internet: Naming

- Dynamic address assignment
 - Most hosts don't need to have known address
 - ▶ Only those functioning as servers
 - DHCP (Dynamic Host Configuration Protocol)
 - ▶ Local ISP assigns address for temporary use

- Example:
 - The CSAP server at SNU (wired connection)
 - ▶ IP address 147.46.174.108
 - ▶ assigned statically
 - My desktop at work
 - ▶ IP address 192.168.1.3
 - ▶ assigned dynamically by a router
 - ▶ only valid within (“behind”) the office router


Evolution of Internet: Firewalls



■ Firewalls

- Hides organizations nodes from rest of Internet
- Use local IP addresses within organization
- For external service, provides proxy service
 - ▶ Client request: src=10.2.2.2, dest=216.99.99.99
 - ▶ Firewall forwards: src=176.3.3.3, dest=216.99.99.99
 - ▶ Server responds: src=216.99.99.99, dest=176.3.3.3
 - ▶ Firewall forwards response: src=216.99.99.99, dest=10.2.2.2

Virtual Private Networks



- Supporting road warrior
 - Employee working remotely with assigned IP address 198.3.3.3
 - Wants to appear to rest of corporation as if working internally
 - ▶ From address 10.6.6.6
 - ▶ Gives access to internal services (e.g., ability to send mail)
- Virtual Private Network (VPN)
 - Overlays private network on top of regular Internet

A Programmer's View of the Internet

- Hosts are mapped to a set of 32-bit IP addresses
 - 147.46.174.108
- The set of IP addresses is mapped to a set of identifiers called Internet domain names
 - 147.46.174.108 is mapped to csap.snu.ac.kr
- A process on one Internet host can communicate with a process on another Internet host over a connection

IP Addresses (IPv4)

- 32-bit IP addresses are stored in an IP address struct
 - IP addresses are always stored in memory in network byte order (big-endian byte order)
 - True in general for any integer transferred in a packet header from one machine to another.
 - ▶ E.g., the port number used to identify an Internet connection.

```
/* Internet address structure */  
struct in_addr {  
    unsigned int s_addr; /* network byte order (big-endian) */  
};
```

Useful network byte-order conversion functions (“l” = 32 bits, “s” = 16 bits)

htonl: convert uint32_t from host to network byte order

htons: convert uint16_t from host to network byte order

ntohl: convert uint32_t from network to host byte order

ntohs: convert uint16_t from network to host byte order

Dotted Decimal Notation (IPv4)

- By convention, each byte in a 32-bit IP address is represented by its decimal value and separated by a period
 - ▶ IP address: 0x932EAE6C = 147.46.174.108
- Functions for converting between binary IP addresses and dotted decimal strings:
 - inet_aton: dotted decimal string → IP address in network byte order
 - inet_ntoa: IP address in network byte order → dotted decimal string
 - “n” denotes network representation
 - “a” denotes application representation

IPv4 Address Structure

- IPv4 Address space divided into classes:

	0	1	2	3	8	16	24	31
Class A	0	Net ID			Host ID			
Class B	1	0	Net ID				Host ID	
Class C	1	1	0	Net ID				Host ID
Class D	1	1	1	0	Multicast address			
Class E	1	1	1	1	Reserved for experiments			

- Network ID Written in form w.x.y.z/n
 - n = number of bits in host address
 - E.g., SNU written as 147.46.0.0/16
 - ▶ Class B address
- Unrouted (private) IP addresses:
 - 10.0.0.0/8 172.16.0.0/12 192.168.0.0/16

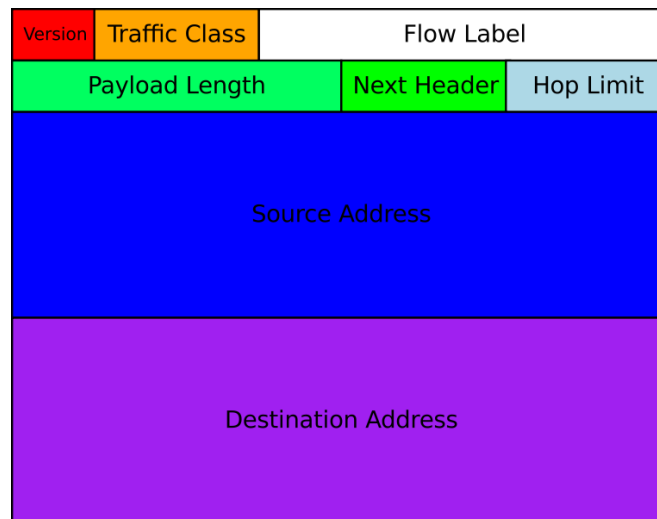
IPv6

■ Internet Protocol version 6

- successor of IPv4
- 128-bit addresses
 - ▶ $2^{128} = 340'282'366'920'938'463'463'374'607'431'768'211'456$ addresses
 - ▶ 8 billion people: $\sim 42'535'295'865'117'307'932'921'825'928$ IPs/person

■ (Improved) Support for

- multicasting
- stateless address autoconfiguration
- IPsec
- simplified routing
- mobility
- privacy



source: Wikipedia

IPv6 Dotted Decimal Notation

- By convention, each 16-bit word in a 128-bit IPv6 address is represented by its hexadecimal value and separated by a period
 - rules
 - ▶ omit leading zeroes
 - ▶ replace one or more groups of consecutive zeroes by a double quotation
 - IPv6 address: 0x20124190020300000000000020120611
 - text representation: 2012:4190:203::2012:0611
- loopback (localhost; 127.0.0.1 in IPv4)
 - IPv6: 0x00000000000000000000000000000001
 - textual representation: ::1

Internet Domain Names


- Domain name: strings (“labels”) separated by dots
 - max. length of a label: 63 characters
 - max. length of the full domain name: 253 characters
 - right-most label designates the top-level domain, each label to the left specifies a subdomain (up to max. 127 levels)
 - example: www.snu.ac.kr
- Fully-qualified vs. unqualified domain names
 - a fully qualified domain name includes a dot at the end
 - ▶ www.snu.ac.kr.
 - for unqualified domain names, in principle, the domain name resolver automatically append the system’s default domain name
 - ▶ “www” inside SNU resolves (should resolve) to www.snu.ac.kr.
 - browsers automatically append the missing dot to get a FQDN
 - ▶ try <http://www.snu.ac.kr/>

Top-Level Domains (TLD)

- Originally, only 7 generic top-level domains (gTLD) existed (predate ICANN)
 - .com, .org, .net, .int, .edu, .gov, .mil
 - current gTLDs:
 - ▶ .com, .info, .net, .org plus (the restricted) .biz, .name, .pro
 - ▶ .edu, .gov, .int, .mil are now considered sponsored TLDs (sTLD)

- Country code top-level domains (ccTLD; predate ICANN)
 - .kr, .ch, .tv, ...
 - recently: internationalized TLDs (IDN TLDs)
 - ▶ xn--3e0b707e → .한국

Generic Top-Level Domains (gTLD)



- application period closed (May 30, 2012): ~2'000 applications
- went live in November 2013, currently (May 15, 2015) >900 gTLDs have been assigned; over 5.5 mio domains using a gTLD have been registered
- so, can I register .egger and then use <http://bernhard.egger> for my homepage?
 - application fee to ICANN: \$185'000
 - annual fee to ICANN: \$25'000
 - plus operational cost
 - um, no, thanks.
- then, who did register new gTLDs?
 - Google (101), Amazon (76), Microsoft (11)
 - Donuts (domain name company): 307
- updated list of top-level domains:
<http://data.iana.org/TLD/tlds-alpha-by-domain.txt>

Domain Naming System (DNS)

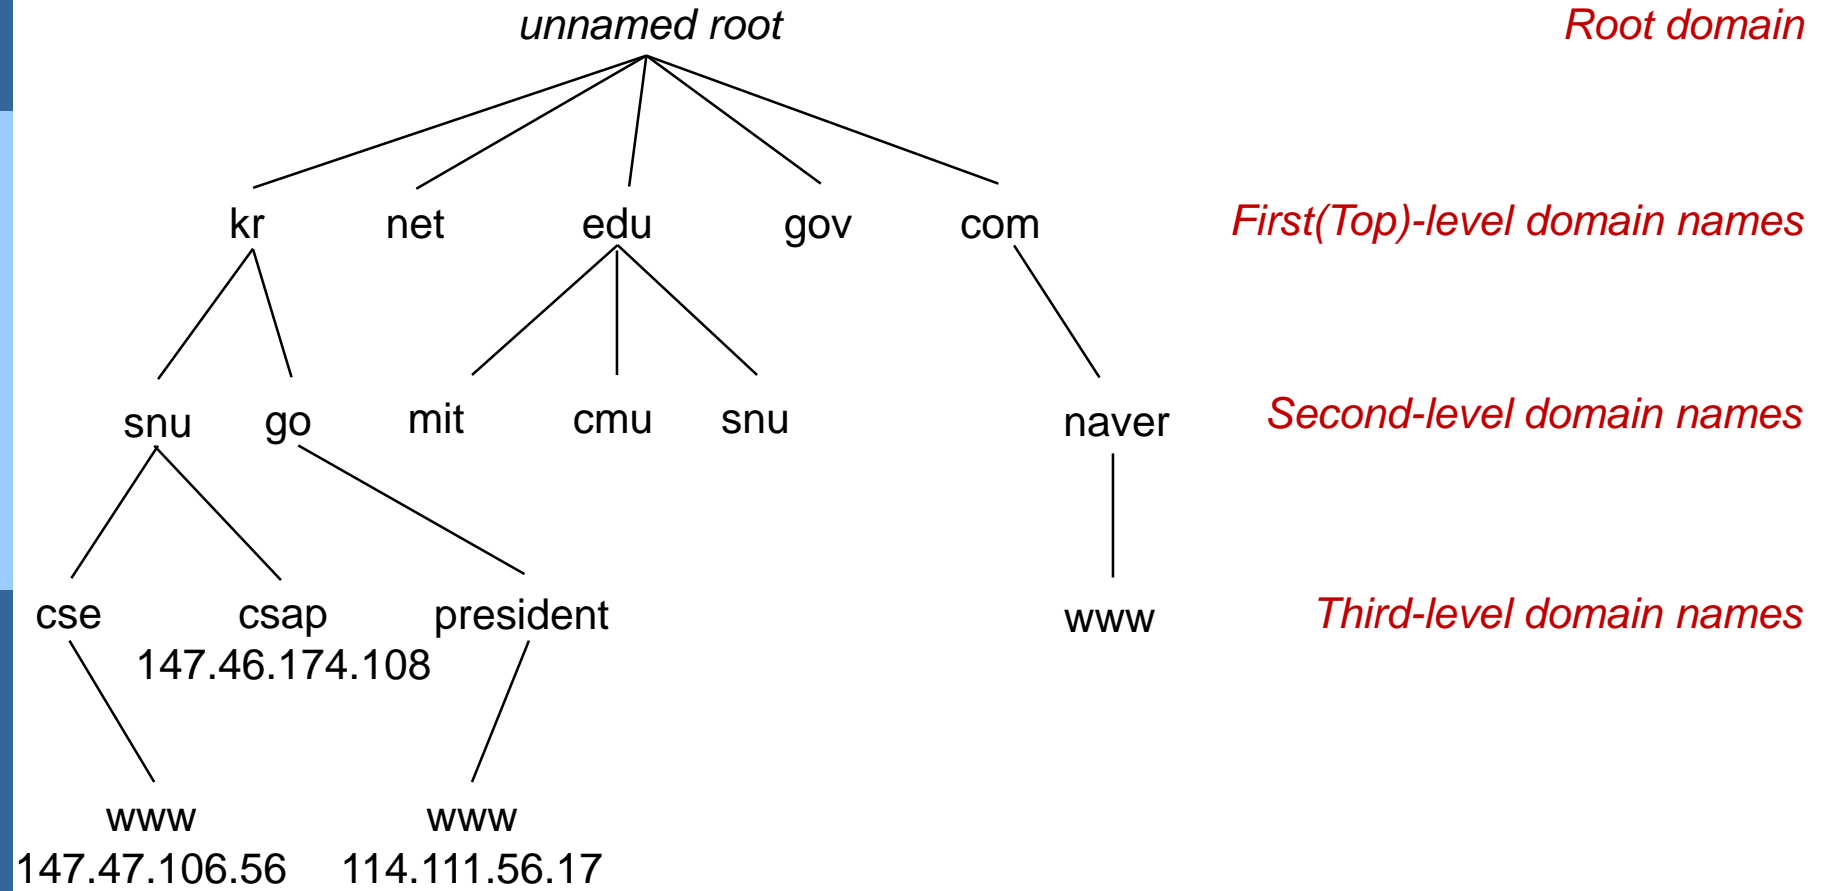
- The Internet maintains a mapping between IP addresses and domain names in a huge worldwide distributed database called DNS
 - Conceptually, programmers can view the DNS database as a collection of millions of host entry structures:

```
/* DNS host entry structure */
struct hostent {
    char    *h_name;           /* official domain name of host */
    char    **h_aliases;       /* null-terminated array of domain names */
    int      h_addrtype;        /* host address type (AF_INET) */
    int      h_length;          /* length of an address, in bytes */
    char    **h_addr_list;     /* null-terminated array of in_addr structs */
};
```

- Functions for retrieving host entries from DNS:
 - gethostbyname: query key is a DNS domain name.
 - gethostbyaddr: query key is an IP address.

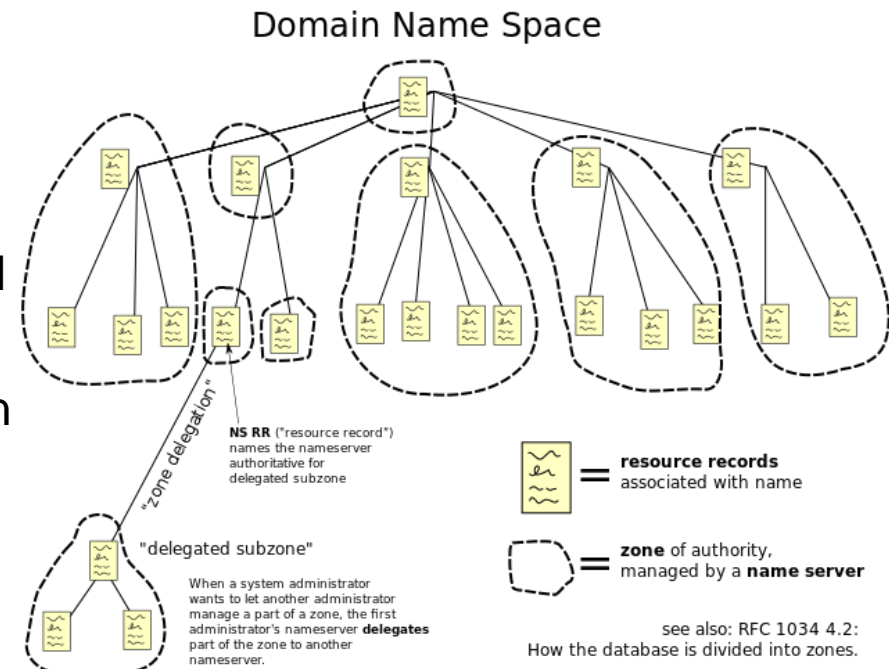
DNS Organization

- Hierarchical organization of DNS servers

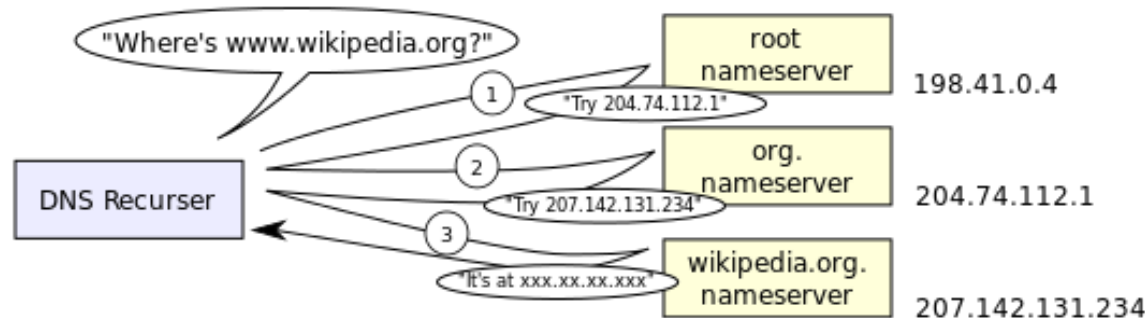


DNS Queries

- Each domain has at least two designated name servers
- These authoritative name servers contain the original DNS data for a domain
- Address resolution
 - recursively start at one of the 13 root name servers and find the authoritative name server for the given domain
 - to reduce traffic and increase lookup, DNS cache servers save results
 - ▶ hence the up to 24-hour period to propagate IP address changes



Source: http://en.wikipedia.org/wiki/Domain_Name_System



Source: http://en.wikipedia.org/wiki/Domain_Name_System

Properties of DNS Host Entries

- Each host entry is an equivalence class of domain names and IP addresses
- Each host has a locally defined domain name localhost which always maps to the loopback address 127.0.0.1
- Different kinds of mappings are possible:
 - Simple case: one-to-one mapping between domain name and IP address:
 - ▶ csap.snu.ac.kr maps to 147.46.174.108
 - Multiple domain names mapped to the same IP address:
 - ▶ eeecs.mit.edu and cs.mit.edu both map to 18.62.1.6
 - Multiple domain names mapped to multiple IP addresses:
 - ▶ google.com maps to multiple IP addresses
 - Some valid domain names don't map to any IP address:
 - ▶ for example: ics.cs.cmu.edu

A Program That Queries DNS

```
int main(int argc, char **argv) { /* argv[1] is a domain name */
    char **pp;                    /* or dotted decimal IP addr */
    struct in_addr addr;
    struct hostent *hostp;

    if (inet_aton(argv[1], &addr) != 0)
        hostp = Gethostbyaddr((const char *)&addr, sizeof(addr),
                               AF_INET);
    else
        hostp = Gethostbyname(argv[1]);
    printf("official hostname: %s\n", hostp->h_name);

    for (pp = hostp->h_aliases; *pp != NULL; pp++)
        printf("alias: %s\n", *pp);

    for (pp = hostp->h_addr_list; *pp != NULL; pp++) {
        addr.s_addr = ((struct in_addr *)*pp)->s_addr;
        printf("address: %s\n", inet_ntoa(addr));
    }
}
```


Using DNS Program

```
linux> ./dns csap.snu.ac.kr
official hostname: csap.snu.ac.kr
address 147.46.174.108

linux> ./dns 147.46.80.1
official hostname: ercc.snu.ac.kr
address: 147.46.80.1

linux> ./dns www.google.com
official hostname: www.l.google.com
alias: www.google.com
address: 74.125.71.105
address: 74.125.71.106
address: 74.125.71.147
address: 74.125.71.99
address: 74.125.71.103
address: 74.125.71.104
```

Querying DIG

- Domain Information Groper (dig) provides a scriptable command line interface to DNS

```
linux> dig +short csap.snu.ac.kr
147.46.174.108
linux> dig +short -x 147.46.80.1
ercc.snu.ac.kr.
linux> dig +short www.google.com
www.l.google.com.
74.125.31.106
74.125.31.147
74.125.31.99
74.125.31.103
74.125.31.104
74.125.31.105
linux> dig +short -x 74.125.71.105
hx-in-f105.1e100.net.
```

More Exotic Features of DIG

- Provides more information than you would ever want about DNS

```
linux> dig www.google.com a +trace

; <<>> DiG 9.8.1 <<>> www.google.com a +trace
;; global options: +cmd
.                351140      IN            NS           g.root-servers.net.
.                351140      IN            NS           i.root-servers.net.
.                351140      IN            NS           l.root-servers.net.
...
.                351140      IN            NS           b.root-servers.net.
.                351140      IN            NS           k.root-servers.net.
.                351140      IN            NS           d.root-servers.net.
;; Received 228 bytes from 147.46.37.10#53(147.46.37.10) in 18 ms

com.             172800      IN            NS           i.gtld-servers.net.
com.             172800      IN            NS           e.gtld-servers.net.
com.             172800      IN            NS           j.gtld-servers.net.
...
com.             172800      IN            NS           c.gtld-servers.net.
com.             172800      IN            NS           h.gtld-servers.net.
;; Received 492 bytes from 192.228.79.201#53(192.228.79.201) in 354 ms


google.com.      172800      IN            NS           ns2.google.com.
...
google.com.      172800      IN            NS           ns4.google.com.
;; Received 168 bytes from 192.55.83.30#53(192.55.83.30) in 206 ms

www.google.com.  604800      IN            CNAME        www.l.google.com.
www.l.google.com. 300         IN            A            74.125.71.103
...
www.l.google.com. 300         IN            A            74.125.71.104
;; Received 148 bytes from 216.239.34.10#53(216.239.34.10) in 98 ms
```

Internet Connections

- Clients and servers communicate by sending streams of bytes over connections:
 - Point-to-point, full-duplex (2-way communication), and reliable.
- A socket is an endpoint of a connection
 - Socket address is an IPaddress:port pair
- A port is a 16-bit integer that identifies a process:
 - Ephemeral port: Assigned automatically on client when client makes a connection request
 - Well-known port: Associated with some service provided by a server (e.g., port 80 is associated with Web servers)
- A connection is uniquely identified by the socket addresses of its endpoints (socket pair)
 - (cliaddr:cliport, servaddr:servport)

Putting it all Together: Anatomy of an Internet Connection



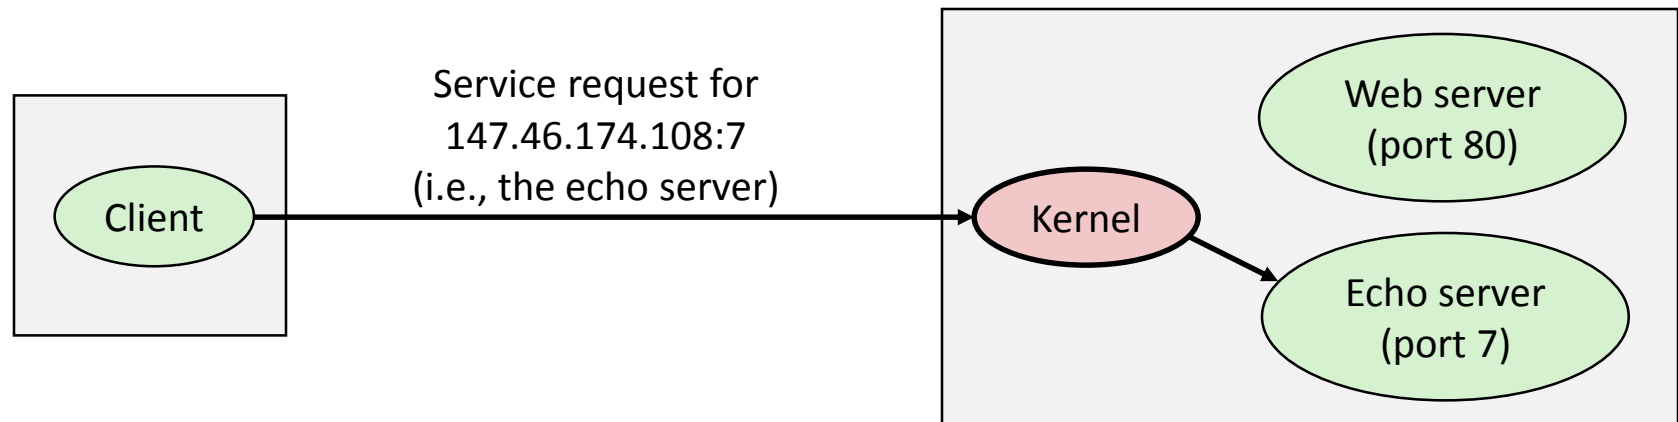
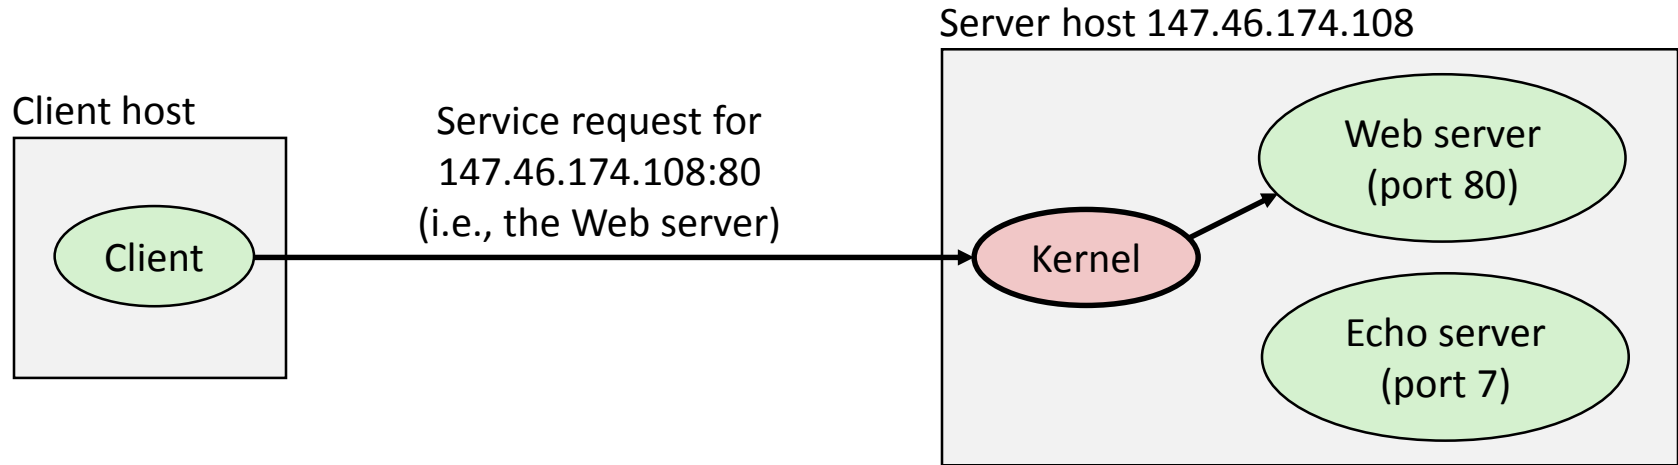
Network Programming

- Client-Server Model and Computer Networks
- Global IP Internet
 - naming
 - IPv4, IPv6
 - DNS
 - Internet connection
- **Programmer's View**
 - **The sockets interface**
 - **Establishing an Internet connection**
 - **Copy data over an Internet connection**

Clients

- Examples of client programs
 - Web browsers, ftp, telnet, ssh
- How does a client find the server?
 - The IP address in the server socket address identifies the host (more precisely, an adapter on the host)
 - The (well-known) port in the server socket address identifies the service, and thus implicitly identifies the server process that performs that service.
 - Examples of well know ports
 - ▶ Port 7: Echo server
 - ▶ Port 23: Telnet server
 - ▶ Port 25: Mail server
 - ▶ Port 80: Web server

Using Ports to Identify Services



Servers

- Servers are long-running processes (daemons)
 - Created at boot-time (typically) by the init process (process 1)
 - Run continuously until the machine is turned off
- Each server waits for requests to arrive on a well-known port associated with a particular service
 - Port 7: echo server
 - Port 23: telnet server
 - Port 25: mail server
 - Port 80: HTTP server
- A machine that runs a server process is also often referred to as a “server”

Server Examples

- Web server (port 80)
 - Resource: files/compute cycles (CGI programs)
 - Service: retrieves files and runs CGI programs on behalf of the client

- FTP server (20, 21)
 - Resource: files
 - Service: stores and retrieve files

- Telnet server (23)
 - Resource: terminal
 - Service: proxies a terminal on the server machine

- Mail server (25)
 - Resource: email “spool” file
 - Service: stores mail messages in spool file


See `/etc/services` for a comprehensive list of the port mappings on a Linux machine

Sockets Interface

- Created in the early 80's as part of the original Berkeley distribution of Unix that contained an early version of the Internet protocols
- Provides a user-level interface to the network
- Underlying basis for all Internet applications
- Based on client/server programming model

Sockets

- What is a socket?
 - To the kernel, a socket is an endpoint of communication
 - To an application, a socket is a file descriptor that lets the application read/write from/to the network
 - ▶ Remember: All Unix I/O devices, including networks, are modeled as files
- Clients and servers communicate with each other by reading from and writing to socket descriptors



- The main distinction between regular file I/O and socket I/O is how the application “opens” the socket descriptors

Example: Echo Client and Server

On Client

```
csap> ./echoserveri 15213
```

```
linux> echoclient csap.snu.ac.kr 15213
```

```
server connected to langnau.snu.ac.kr  
(147.47.174.115), port 64690
```

```
type: hello there
```

```
server received 12 bytes
```

```
echo: HELLO THERE  
type: ^D
```

```
Connection closed
```

On Server

Watching Echo Client / Server

Capturing from eno1 [Wireshark 1.12.4 (Git Rev Unknown from unknown)]

File Edit View Go Capture Analyze Statistics Telephony Tools Internals Help

Filter: tcp.port eq 1522 Expression... Clear Apply Save

No.	Time	Source	Destination	Protocol	Length	Info
7	2.867873000	192.168.1.3	147.46.174.108	TCP	74	43518→1522 [SYN] Seq=0 Win=29200 Len=0 MSS=1460 SA
8	2.868406000	147.46.174.108	192.168.1.3	TCP	74	1522→43518 [SYN, ACK] Seq=0 Ack=1 Win=14480 Len=0
9	2.868416000	192.168.1.3	147.46.174.108	TCP	66	43518→1522 [ACK] Seq=1 Ack=1 Win=29312 Len=0 TSva
24	12.577373000	192.168.1.3	147.46.174.108	TCP	93	43518→1522 [PSH, ACK] Seq=1 Ack=1 Win=29312 Len=27
25	12.577821000	147.46.174.108	192.168.1.3	TCP	66	1522→43518 [ACK] Seq=1 Ack=28 Win=14592 Len=0 TSva
26	12.577828000	147.46.174.108	192.168.1.3	TCP	93	1522→43518 [PSH, ACK] Seq=1 Ack=28 Win=14592 Len=27
27	12.577845000	192.168.1.3	147.46.174.108	TCP	66	43518→1522 [ACK] Seq=28 Ack=28 Win=29312 Len=0 TSv
46	19.777441000	192.168.1.3	147.46.174.108	TCP	66	43518→1522 [FIN, ACK] Seq=28 Ack=28 Win=29312 Len=0
47	19.777843000	147.46.174.108	192.168.1.3	TCP	66	1522→43518 [FIN, ACK] Seq=28 Ack=29 Win=14592 Len=0
48	19.777852000	192.168.1.3	147.46.174.108	TCP	66	43518→1522 [ACK] Seq=29 Ack=29 Win=29312 Len=0 TSv

▶ Frame 24: 93 bytes on wire (744 bits), 93 bytes captured (744 bits) on interface 0

▶ Ethernet II, Src: AsustekC_8a:e9:4a (78:24:af:8a:e9:4a), Dst: Netgear_fd:c5:fc (e4:f4:c6:fd:c5:fc)

▶ Internet Protocol Version 4, Src: 192.168.1.3 (192.168.1.3), Dst: 147.46.174.108 (147.46.174.108)

▶ Transmission Control Protocol, Src Port: 43518 (43518), Dst Port: 1522 (1522), Seq: 1, Ack: 1, Len: 27

▶ Data (27 bytes)

```

0000  e4 f4 c6 fd c5 fc 78 24 af 8a e9 4a 08 00 45 00  .....x$ ...J...E.
0010  00 4f ce 0c 40 00 40 06 69 56 c0 a8 01 03 93 2e  .0..@.@. iV.....
0020  ae 6c a9 fe 05 f2 44 54 c9 47 50 9c 0e 4b 80 18  .l....DT .GP..K..
0030  00 e5 03 88 00 00 01 01 08 0a 2e 4c 12 ad 3d 58  ..... ..L..=X
0040  57 f7 49 20 6c 6f 76 65 20 53 79 73 74 65 6d 20  W.I love System
0050  50 72 6f 67 72 61 6d 6d 69 6e 67 21 0a          Programm ing!.
```

eno1: <live capture in progress> Fil... Packets: 180 · Displayed: 10 (5.6%) Profile: Default

Ethical Issues

■ Packet Sniffer

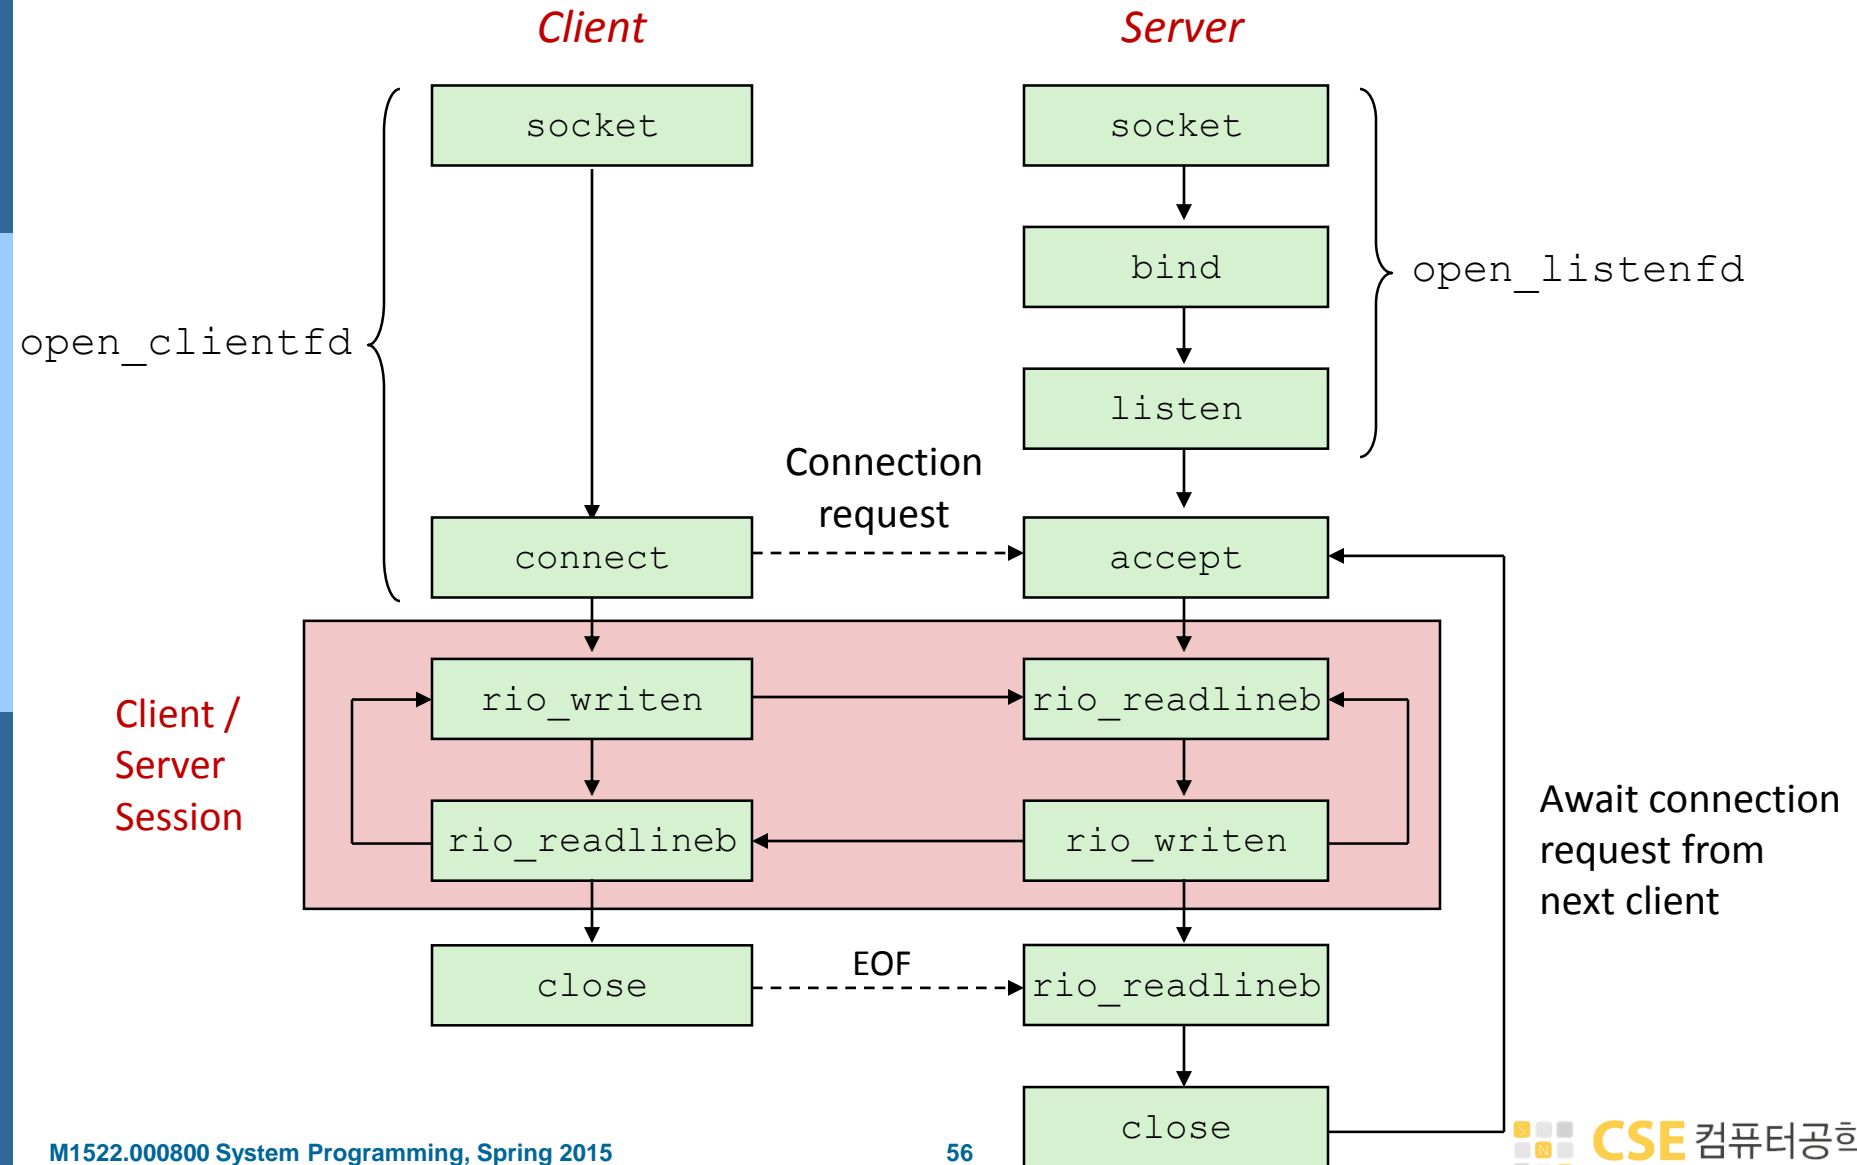
- Program that records network traffic visible at node
- Promiscuous mode: Record traffic that does not have this host as source or destination

■ Network traffic should be considered private; as a consequence, any packet interception of network information that is not intended for your use is not only an invasion of privacy but may also be a violation of university policy.

■ Pledge

- Instructors, TAs, and students of this class will not run packet sniffers in promiscuous mode inside SNU's network.

Overview of the Sockets Interface




Socket Address Structures

■ Generic socket address:

- For address arguments to connect, bind, and accept
- Necessary only because C did not have generic (void *) pointers when the sockets interface was designed

```
struct sockaddr {  
    unsigned short  sa_family;    /* protocol family */  
    char            sa_data[14]; /* address data. */  
};
```

sa_family



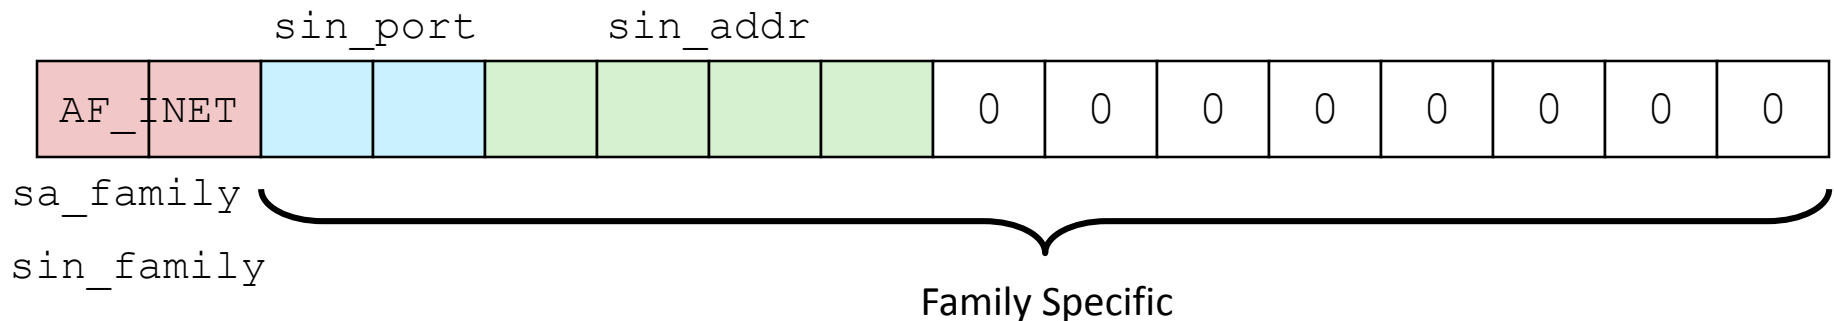
Family Specific

Socket Address Structures

■ Internet-specific socket address:

- Must cast (sockaddr_in *) to (sockaddr *) for connect, bind, and accept

```
struct sockaddr_in {  
    unsigned short    sin_family; /* address family (always AF_INET) */  
    unsigned short    sin_port;  /* port num in network byte order */  
    struct in_addr     sin_addr;  /* IP addr in network byte order */  
    unsigned char      sin_zero[8]; /* pad to sizeof(struct sockaddr) */  
};
```



Echo Client Main Routine

```
#include "csapp.h"

/* usage: ./echoclient host port */
int main(int argc, char **argv)
{
    int clientfd, port;
    char *host, buf[MAXLINE];
    rio_t rio;
    host = argv[1]; port = atoi(argv[2]);
    clientfd = Open_clientfd(host, port);
    Rio_readinitb(&rio, clientfd);
    printf("type:"); fflush(stdout);
    while (Fgets(buf, MAXLINE, stdin) != NULL) {
        Rio_writen(clientfd, buf, strlen(buf));
        Rio_readlineb(&rio, buf, MAXLINE);
        printf("echo:");
        Fputs(buf, stdout);
        printf("type:"); fflush(stdout);
    }
    Close(clientfd);
    exit(0);
}
```

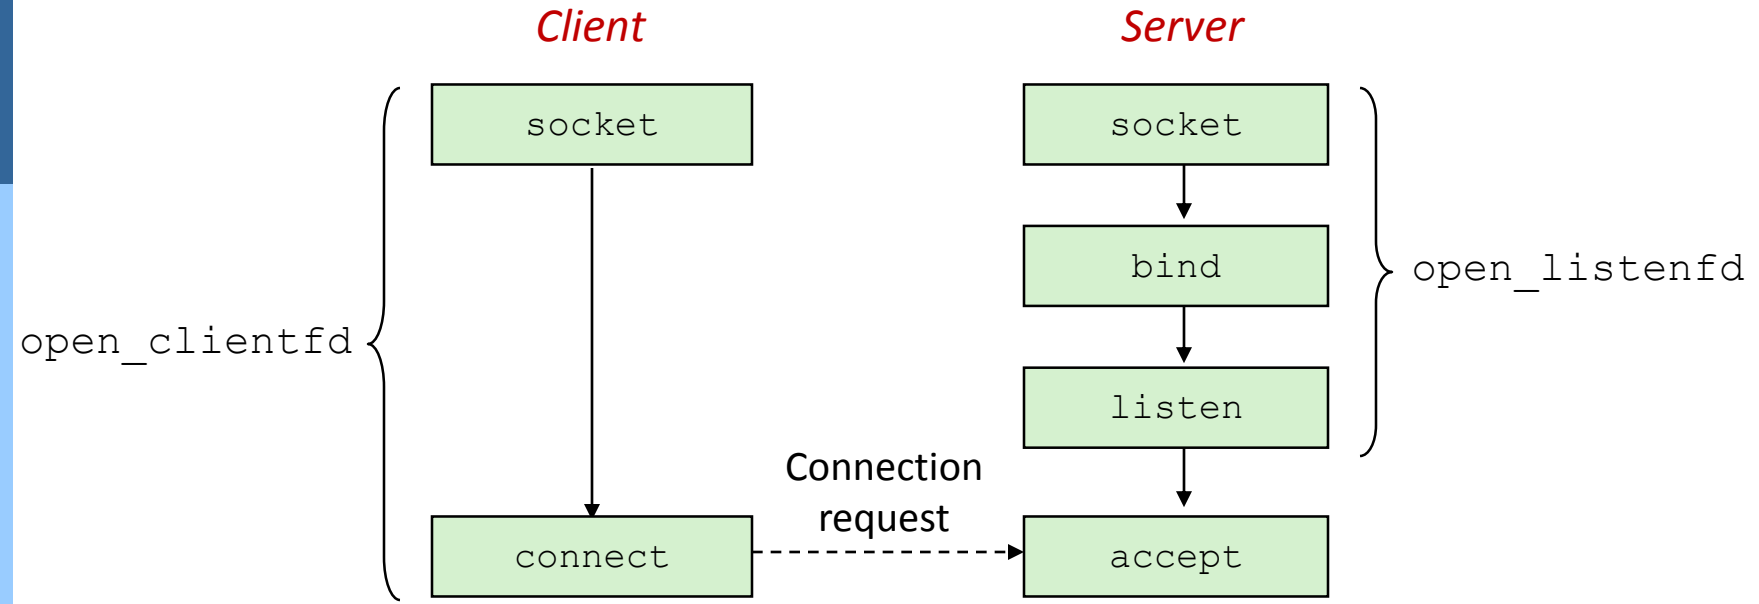
Send line to
server

Receive line
from server

Read input
line

Print server
response

Overview of the Sockets Interface



Echo Client: open_clientfd

```
int open_clientfd(char *hostname, int port) {
```

```
    int clientfd;
```

```
    struct hostent *hp;
```

```
    struct sockaddr_in serveraddr;
```

This function opens a connection from the client to the server at hostname:port

```
    if ((clientfd = socket(AF_INET, SOCK_STREAM, 0)) < 0)
        return -1; /* check errno for cause of error */
```

} Create socket

```
    /* Fill in the server's IP address and port */
```

```
    if ((hp = gethostbyname(hostname)) == NULL)
```

```
        return -2; /* check h_errno for cause of error */
```

```
    bzero((char *) &serveraddr, sizeof(serveraddr));
```

```
    serveraddr.sin_family = AF_INET;
```

```
    bcopy((char *) hp->h_addr_list[0],
```

```
          (char *) &serveraddr.sin_addr.s_addr, hp->h_length);
```

```
    serveraddr.sin_port = htons(port);
```

} Create address

```
    /* Establish a connection with the server */
```

```
    if (connect(clientfd, (SA *) &serveraddr,
                sizeof(serveraddr)) < 0)
```

```
        return -1;
```

```
    return clientfd;
```

```
}
```

} Establish connection

Echo Client: open_clientfd (socket)

- socket creates a socket descriptor on the client
 - Just allocates & initializes some internal data structures
 - AF_INET: indicates that the socket is associated with Internet protocols
 - SOCK_STREAM: selects a reliable byte stream connection
 - ▶ provided by TCP

```
int clientfd; /* socket descriptor */  
  
if ((clientfd = socket(AF_INET, SOCK_STREAM, 0)) < 0)  
    return -1; /* check errno for cause of error */  
  
... <more>
```

Echo Client: open_clientfd (gethostbyname)

- The client then builds the server's Internet address

```
int clientfd; /* socket descriptor */
struct hostent *hp; /* DNS host entry */
struct sockaddr_in serveraddr; /* server's IP address */

...

/* fill in the server's IP address and port */
if ((hp = gethostbyname(hostname)) == NULL)
    return -2; /* check h_errno for cause of error */
bzero((char *) &serveraddr, sizeof(serveraddr));
serveraddr.sin_family = AF_INET;
serveraddr.sin_port = htons(port);
bcopy((char *)hp->h_addr_list[0],
      (char *)&serveraddr.sin_addr.s_addr, hp->h_length);
```

Check
this out!

A Careful Look at bcopy Arguments

```
/* DNS host entry structure */
struct hostent {
    . . .
    int      h_length;      /* length of an address, in bytes */
    char     **h_addr_list; /* null-terminated array of in_addr structs */
};
```

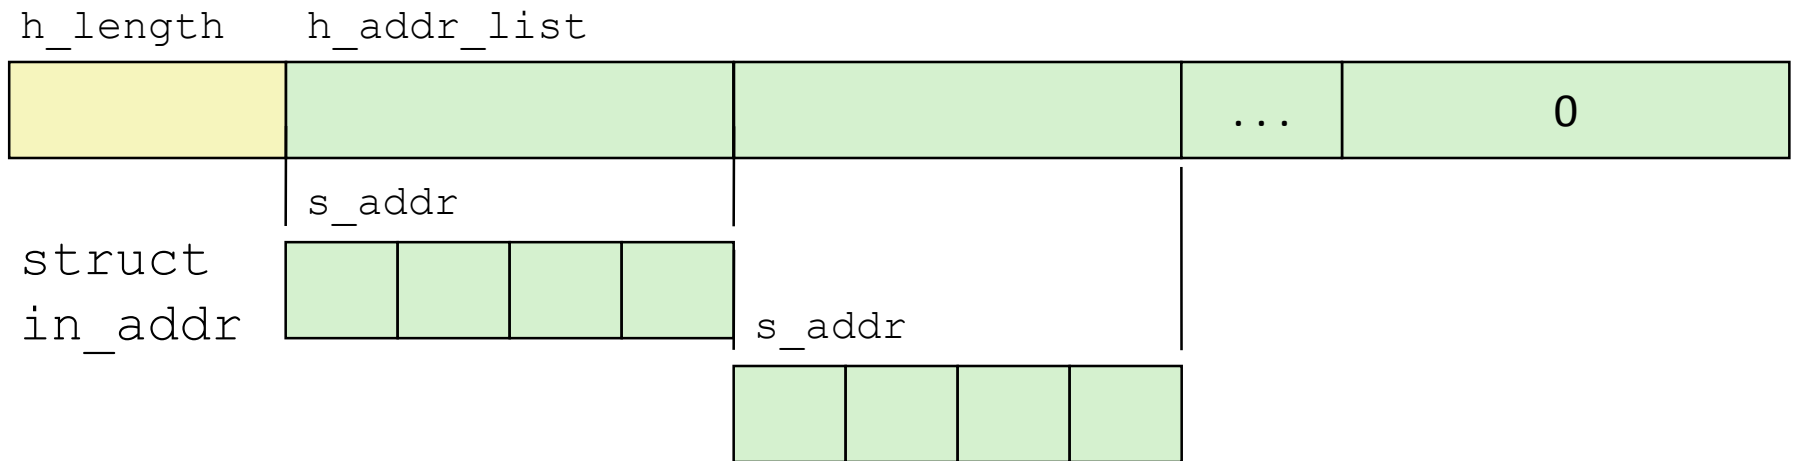
```
struct sockaddr_in {
    . . .
    struct in_addr  sin_addr; /* IP addr in network byte order */
    . . .
};
```

```
/* Internet address structure */
struct in_addr {
    unsigned int s_addr; /* network byte order (big-endian) */
};
```


```
struct hostent *hp; /* DNS host entry */
struct sockaddr_in serveraddr; /* server's IP address */
...
bcopy((char *)hp->h_addr_list[0], /* src, dest */
      (char *)&serveraddr.sin_addr.s_addr, hp->h_length);
```


Bcopy Argument Data Structures

struct hostent



struct sockaddr_in



Echo Client: open_clientfd (connect)

- Finally the client creates a connection with the server
 - Client process suspends (blocks) until the connection is created
 - After resuming, the client is ready to begin exchanging messages with the server via Unix I/O calls on descriptor clientfd

```
int clientfd;                /* socket descriptor */
struct sockaddr_in serveraddr; /* server address */
typedef struct sockaddr SA;   /* generic sockaddr */
...
/* Establish a connection with the server */
if (connect(clientfd, (SA *)&serveraddr, sizeof(serveraddr)) < 0)
    return -1;
return clientfd;
}
```

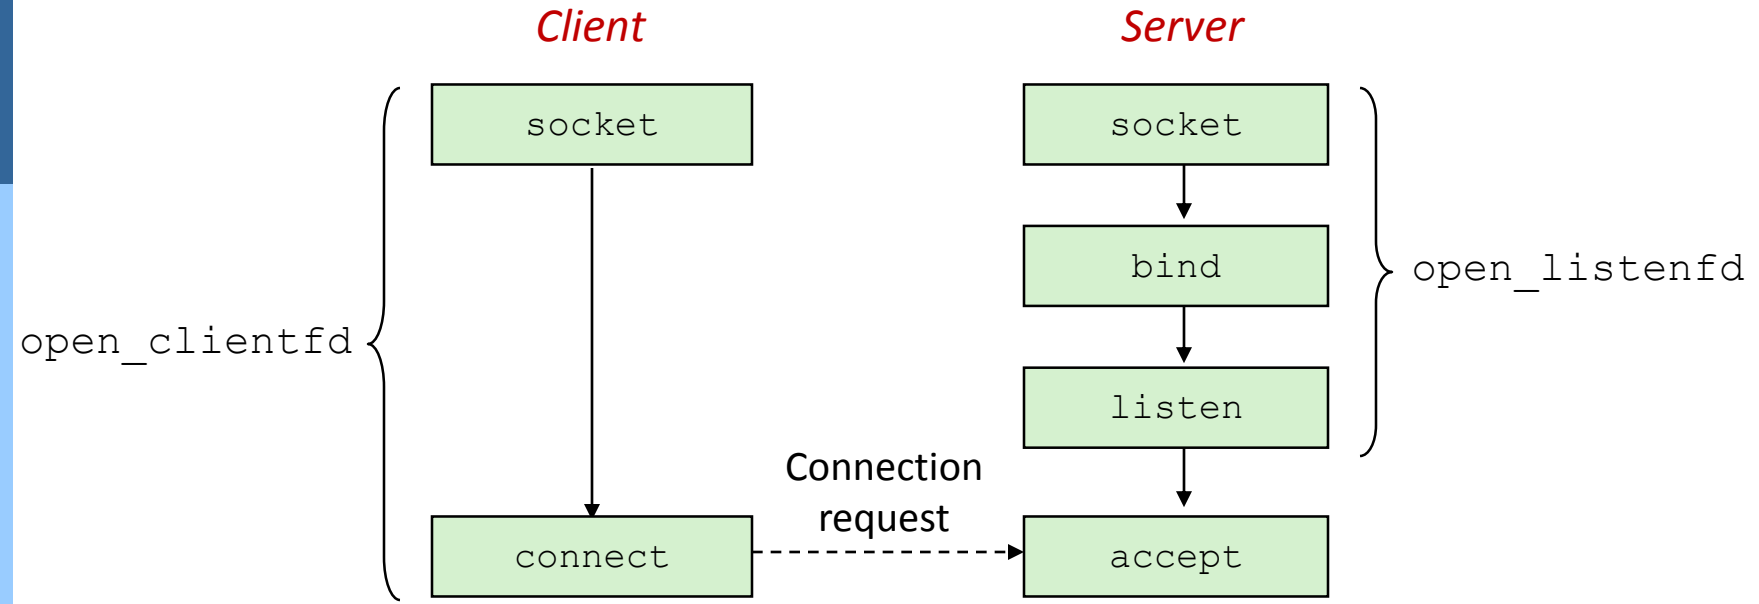
Echo Server: Main Routine

```
int main(int argc, char **argv) {
    int listenfd, connfd, port, clientlen;
    struct sockaddr_in clientaddr;
    struct hostent *hp;
    char *haddrp;
    unsigned short client_port;

    port = atoi(argv[1]); /* the server listens on a port passed
                           on the command line */
    listenfd = open_listenfd(port);

    while (1) {
        clientlen = sizeof(clientaddr);
        connfd = Accept(listenfd, (SA *)&clientaddr, &clientlen);
        hp = Gethostbyaddr((const char *)&clientaddr.sin_addr.s_addr,
                           sizeof(clientaddr.sin_addr.s_addr), AF_INET);
        haddrp = inet_ntoa(clientaddr.sin_addr);
        client_port = ntohs(clientaddr.sin_port);
        printf("server connected to %s (%s), port %u\n",
               hp->h_name, haddrp, client_port);
        echo(connfd);
        Close(connfd);
    }
}
```

Overview of the Sockets Interface



■ Office Telephone Analogy for Server

- Socket: Buy a phone
- Bind: Tell the local administrator what number you want to use
- Listen: Plug the phone in
- Accept: Answer the phone when it rings

Echo Server: open_listenfd

```
int open_listenfd(int port)
{
    int listenfd, optval=1;
    struct sockaddr_in serveraddr;

    /* Create a socket descriptor */
    if ((listenfd = socket(AF_INET, SOCK_STREAM, 0)) < 0)
        return -1;

    /* Eliminates "Address already in use" error from bind. */
    if (setsockopt(listenfd, SOL_SOCKET, SO_REUSEADDR,
                    (const void *)&optval , sizeof(int)) < 0)
        return -1;

    ... <more>
}
```

Echo Server: open_listenfd (cont.)

```
...

/* Listenfd will be an endpoint for all requests to port
   on any IP address for this host */
bzero((char *) &serveraddr, sizeof(serveraddr));
serveraddr.sin_family = AF_INET;
serveraddr.sin_addr.s_addr = htonl(INADDR_ANY);
serveraddr.sin_port = htons((unsigned short)port);
if (bind(listenfd, (SA *)&serveraddr, sizeof(serveraddr)) < 0)
    return -1;

/* Make it a listening socket ready to accept
   connection requests */
if (listen(listenfd, LISTENQ) < 0)
    return -1;

return listenfd;
}
```

Echo Server: open_listenfd (socket)

- socket creates a socket descriptor on the server
 - AF_INET: indicates that the socket is associated with Internet protocols
 - SOCK_STREAM: selects a reliable byte stream connection (TCP)

```
int listenfd; /* listening socket descriptor */

/* Create a socket descriptor */
if ((listenfd = socket(AF_INET, SOCK_STREAM, 0)) < 0)
    return -1;
```

Echo Server: open_listenfd (setsockopt)

- The socket can be given some attributes

```
...  
/* Eliminates "Address already in use" error from bind(). */  
if (setsockopt(listenfd, SOL_SOCKET, SO_REUSEADDR,  
              (const void *)&optval , sizeof(int)) < 0)  
    return -1;
```

- Handy trick that allows us to rerun the server immediately after we kill it
 - Otherwise we would have to wait about 15 seconds
 - Eliminates “Address already in use” error from bind()
- Strongly suggest you do this for all your servers to simplify debugging


Echo Server: open_listenfd (initialize socket address)

- Initialize socket with server port number
- Accept connection from any IP address

```
struct sockaddr_in serveraddr; /* server's socket addr */
...
/* listenfd will be an endpoint for all requests to port
   on any IP address for this host */
bzero((char *) &serveraddr, sizeof(serveraddr));
serveraddr.sin_family = AF_INET;
serveraddr.sin_port = htons((unsigned short)port);
serveraddr.sin_addr.s_addr = htonl(INADDR_ANY);
```

- IP addr and port stored in network (big-endian) byte order

sin_port sin_addr



sa_family

sin_family

Echo Server: open_listenfd (bind)

- bind associates the socket with the socket address we just created

```
int listenfd;                /* listening socket */
struct sockaddr_in serveraddr; /* server's socket addr */

...
/* listenfd will be an endpoint for all requests to port
   on any IP address for this host */
if (bind(listenfd, (SA *)&serveraddr, sizeof(serveraddr)) < 0)
    return -1;
```

Echo Server: open_listenfd (listen)

- listen indicates that this socket will accept connection (connect) requests from clients
- LISTENQ is constant indicating how many pending requests allowed

```
int listenfd; /* listening socket */  
  
...  
/* Make it a listening socket ready to accept connection requests */  
if (listen(listenfd, LISTENQ) < 0)  
    return -1;  
return listenfd;  
}
```

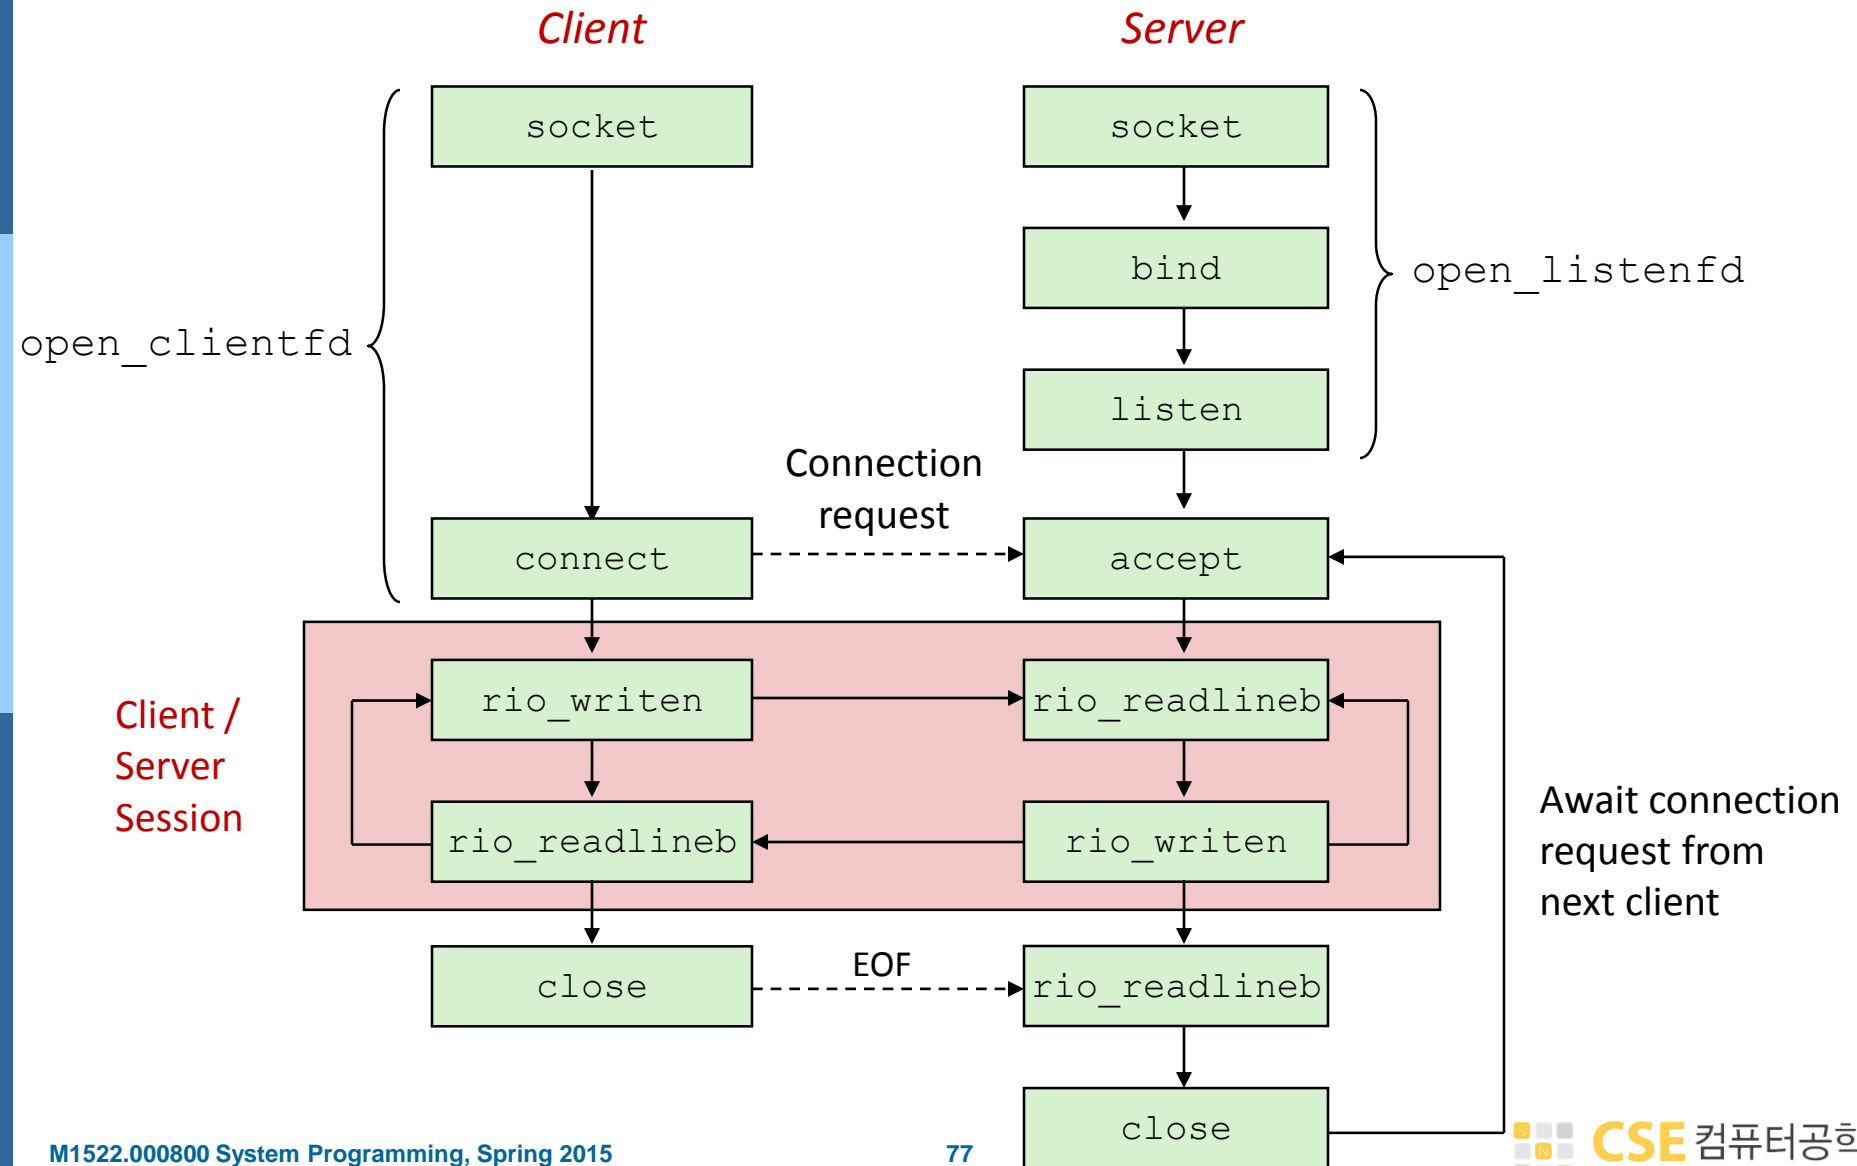
- We're finally ready to enter the main server loop that accepts and processes client connection requests.

Echo Server: Main Loop

- The server loops endlessly, waiting for connection requests, then reading input from the client, and echoing the input back to the client.

```
main() {  
  
    /* create and configure the listening socket */  
  
    while(1) {  
        /* Accept(): wait for a connection request */  
        /* echo(): read and echo input lines from client til EOF */  
        /* Close(): close the connection */  
    }  
}
```

Overview of the Sockets Interface



Echo Server: accept

- `accept()` blocks waiting for a connection request

```
int listenfd; /* listening descriptor */
int connfd;   /* connected descriptor */
struct sockaddr_in clientaddr;
int clientlen;


clientlen = sizeof(clientaddr);
connfd = Accept(listenfd, (SA *)&clientaddr, &clientlen);
```

- `accept` returns a connected descriptor (`connfd`) with the same properties as the listening descriptor (`listenfd`)
 - Returns when the connection between client and server is created and ready for I/O transfers
 - All I/O with the client will be done via the connected socket
- `accept` also fills in client's IP address


Echo Server: accept Illustrated



1. Server blocks in *accept*,
waiting for connection request
on listening descriptor
listenfd



2. Client makes connection request by
calling and blocking in *connect*



3. Server returns *connfd* from *accept*.
Client returns from *connect*.
Connection is now established between
clientfd and *connfd*

Connected vs. Listening Descriptors

- Listening descriptor
 - End point for client connection requests
 - Created once and exists for lifetime of the server

- Connected descriptor
 - End point of the connection between client and server
 - A new descriptor is created each time the server accepts a connection request from a client
 - Exists only as long as it takes to service client

- Why the distinction?
 - Allows for concurrent servers that can communicate over many client connections simultaneously
 - ▶ E.g., Each time we receive a new request, we fork a child to handle the request

Echo Server: Identifying the Client

- The server can determine the domain name, IP address, and port of the client

```
struct hostent *hp; /* pointer to DNS host entry */
char *haddrp;      /* pointer to dotted decimal string */
unsigned short client_port;
hp = Gethostbyaddr((const char *)&clientaddr.sin_addr.s_addr,
                  sizeof(clientaddr.sin_addr.s_addr), AF_INET);
haddrp = inet_ntoa(clientaddr.sin_addr);
client_port = ntohs(clientaddr.sin_port);
printf("server connected to %s (%s), port %u\n",
       hp->h_name, haddrp, client_port);
```

Echo Server: echo

- The server uses RIO to read and echo text lines until EOF (end-of-file) is encountered.
 - EOF notification caused by client calling `close(clientfd)`

```
void echo(int connfd)
{
    size_t n;
    char buf[MAXLINE];
    rio_t rio;

    Rio_readinitb(&rio, connfd);
    while((n = Rio_readlineb(&rio, buf, MAXLINE)) != 0) {
        upper_case(buf);
        Rio_writen(connfd, buf, n);
        printf("server received %d bytes\n", n);
    }
}
```

Testing Servers Using telnet

- The telnet program is invaluable for testing servers that transmit ASCII strings over Internet connections
 - Our simple echo server
 - Web servers
 - Mail servers
- Usage:
 - `unix> telnet <host> <portnumber>`
 - Creates a connection with a server running on <host> and listening on port <portnumber>

Testing the Echo Server With telnet

```
greatwhite> echoserver 15213
```

```
linux> telnet csap.snu.ac.kr 15213
Trying 147.46.174.108...
Connected to csap.snu.ac.kr.
Escape character is '^]'.
hi there
HI THERE
```

For More Information

- W. Richard Stevens, “Unix Network Programming: Networking APIs: Sockets and XTI”, Volume 1, Second Edition, Prentice Hall, 1998
 - THE network programming bible
- Unix Man Pages
 - Good for detailed information about specific functions
- Complete versions of the echo client and server are developed in the text
 - Updated versions linked to course website
 - Feel free to use this code in your assignments