

Interaction and Communication between Programs

Synchronization in Concurrent Programs



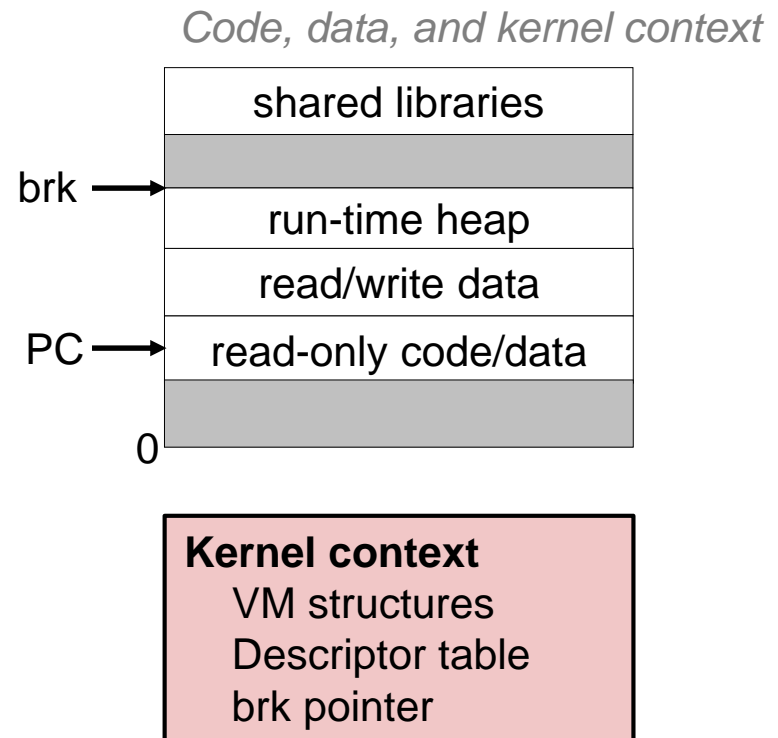
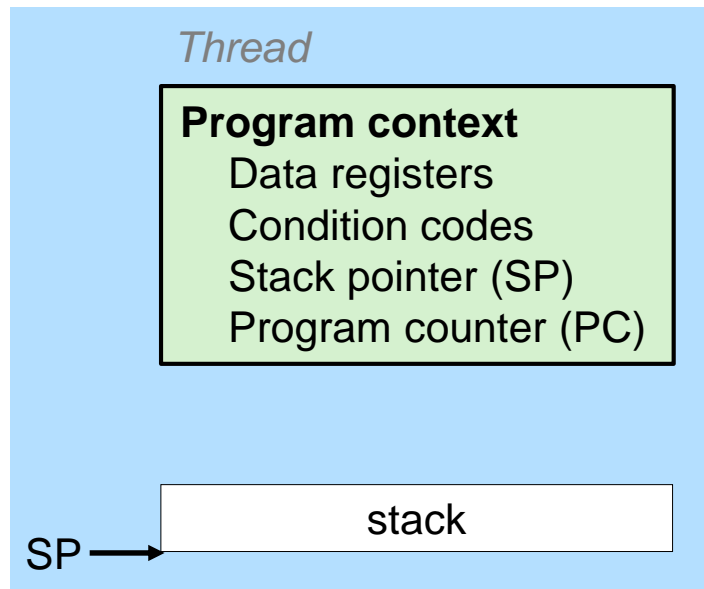
Synchronization in Concurrent Programs

- **Threads review**
- Sharing
- Mutual exclusion
- Semaphores
- Synchronization
 - Producer-consumer problem
 - Readers-writers problem
 - Thread safety
 - Races
 - Deadlocks

Acknowledgement: slides based on the cs:app2e material

Process: Single Thread View

- Process = main thread + code, data, and kernel context



Process with Two Threads


Thread 1

Program context:
Data registers
Condition codes
Stack pointer (SP)
Program counter (PC)




Thread 2

Program context:
Data registers
Condition codes
Stack pointer (SP)
Program counter (PC)



Code, data, and kernel context



Kernel context:
VM structures
Descriptor table
brk pointer

Threads vs. Processes

■ Threads and processes: similarities

- Each has its own logical control flow
- Each can run concurrently with others
- Each is context switched (scheduled) by the kernel


■ Threads and processes: differences

- Threads share code and data, processes (typically) do not
- Threads are less expensive than processes
 - ▶ Process control (creating and reaping) is more expensive as thread control
 - ▶ Context switches for processes more expensive than for threads


Threads vs. Processes (cont.)

- Processes form a tree hierarchy
- Threads form a pool of peers
 - Each thread can kill any other
 - Each thread can wait for any other thread to terminate
 - Main thread: first thread to run in a process

Process hierarchy



Thread pool



Posix Threads (Pthreads) Interface

- Pthreads: Standard interface for ~60 functions that manipulate threads from C programs
 - Creating and reaping threads
 - ▶ `pthread_create()`
 - ▶ `pthread_join()`
 - Determining your thread ID
 - ▶ `pthread_self()`
 - Terminating threads
 - ▶ `pthread_cancel()`
 - ▶ `pthread_exit()`
 - ▶ `exit()` [terminates all threads] , `RET` [terminates current thread]
 - Synchronizing access to shared variables
 - ▶ `pthread_mutex_init`
 - ▶ `pthread_mutex_[un]lock`
 - ▶ `pthread_cond_init`
 - ▶ `pthread_cond_[timed]wait`

The Pthreads "hello, world" Program

```
/*
 * hello.c - Pthreads "hello, world" program
 */
#include "csapp.h"

void *thread(void *vargp);

int main() {
    pthread_t tid;

    Pthread_create(&tid, NULL, thread, NULL);
    Pthread_join(tid, NULL);
    exit(0);
}
```

Thread attributes
(usually NULL)

Thread arguments
(void *p)

return value
(void **p)

```
/* thread routine */
void *thread(void *vargp) {
    printf("Hello, world!\n");
    return NULL;
}
```


Pros and Cons of Thread-Based Designs

- + Easy to share data structures between thread
 - e.g., logging information, file cache.
- + Threads are more efficient than processes
- Unintentional sharing can introduce subtle and hard-to-reproduce errors!
 - The ease with which data can be shared is both the greatest strength and the greatest weakness of threads.
 - Hard to know which data shared & which private
 - Hard to detect by testing
 - ▶ Probability of bad race outcome very low
 - ▶ But nonzero!

Synchronization in Concurrent Programs

- Threads review
- **Sharing**
- Mutual exclusion
- Semaphores
- Synchronization
 - Producer-consumer problem
 - Readers-writers problem
 - Thread safety
 - Races
 - Deadlocks

Shared Variables in Threaded C Programs

- Question: Which variables in a threaded C program are shared?
 - The answer is not as simple as “*global variables are shared*” and “*stack variables are private*”
- Requires answers to the following questions:
 - What is the memory model for threads?
 - How are instances of variables mapped to memory?
 - How many threads might reference each of these instances?
- **Definition:**
A variable **x** is shared if and only if multiple threads reference some instance of **x**.

Threads Memory Model

■ Conceptual model:

- Multiple threads run within the context of a single process
- Each thread has its own separate thread context
 - ▶ Thread ID, stack, stack pointer, PC, condition codes, and GP registers
- All threads share the remaining process context
 - ▶ Code, data, heap, and shared library segments of the process virtual address space
 - ▶ Open files and installed handlers

■ Operationally, this model is not strictly enforced:

- Register values are truly separate and protected, but...
- Any thread can read and write the stack of any other thread

→ This mismatch between the conceptual and operation model is a source of confusion and errors

Example Program to Illustrate Sharing

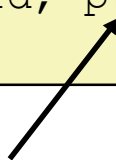
```
char **ptr;  /* global */

int main()
{
    int i;
    pthread_t tid;
    char *msgs[2] = {
        "Hello from foo",
        "Hello from bar"
    };
    ptr = msgs;

    for (i = 0; i < 2; i++)
        Pthread_create(&tid,
            NULL,
            thread,
            (void *)i);
    Pthread_exit(NULL);
}
```

```
/* thread routine */
void *thread(void *vargp)
{
    int myid = (int) vargp;
    static int cnt = 0;

    printf("[%d]: %s (svar=%d)\n",
        myid, ptr[myid], ++cnt);
}
```



Peer threads reference main thread's stack indirectly through global `ptr` variable

Mapping Variable Instances to Memory

■ Global variables

- Def: Variable declared outside of a function
- Virtual memory contains exactly one instance of any global variable

■ Local variables

- Def: Variable declared inside function without static attribute
- Each thread stack contains one instance of each local variable

■ Local static variables

- Def: Variable declared inside function with the static attribute
- Virtual memory contains exactly one instance of any local static variable.

Mapping Variable Instances to Memory

Global var. 1 instance (`ptr [data]`)

```
char **ptr; /* global */

int main()
{
    int i;
    pthread_t tid;
    char *msgs[2] = {
        "Hello from foo",
        "Hello from bar"
    };
    ptr = msgs;

    for (i = 0; i < 2; i++)
        Pthread_create(&tid,
            NULL,
            thread,
            (void *)i);
    Pthread_exit(NULL);
}
```

Local vars: 1 instance (`i.m, msgs.m`)

Local var: 2 instances (
 `myid.p0` [peer thread 0's stack],
 `myid.p1` [peer thread 1's stack]
)

```
/* thread routine */
void *thread(void *vargp)
{
    int myid = (int)vargp;
    static int cnt = 0;

    printf("[%d]: %s (svar=%d)\n",
        myid, ptr[myid], ++cnt);
}
```

Local static var. 1 instance (`cnt [data]`)

Shared Variable Analysis

■ Which variables are shared?

<i>Variable instance</i>	<i>Referenced by main thread?</i>	<i>Referenced by peer thread 0?</i>	<i>Referenced by peer thread 1?</i>
ptr	yes	yes	yes
cnt	no	yes	yes
i.m	yes	no	no
msgs.m	yes	yes	yes
myid.p0	no	yes	no
myid.p1	no	no	yes

■ Answer: A variable x is shared iff multiple threads reference at least one instance of x . Thus:

- ptr, cnt, and msgs are shared
- i and myid are not shared

Synchronization in Concurrent Programs

- Threads review
- Sharing
- **Mutual exclusion**
- Semaphores
- Synchronization
 - Producer-consumer problem
 - Readers-writers problem
 - Thread safety
 - Races
 - Deadlocks

badcnt.c: Improper Synchronization

```
volatile int cnt = 0; /* global */

int main(int argc, char **argv)
{
    int niters = atoi(argv[1]);
    pthread_t tid1, tid2;

    Pthread_create(&tid1, NULL,
                  thread, &niters);
    Pthread_create(&tid2, NULL,
                  thread, &niters);

    Pthread_join(tid1, NULL);
    Pthread_join(tid2, NULL);

    /* Check result */
    if (cnt != (2 * niters))
        printf("BOOM! cnt=%d\n", cnt);
    else
        printf("OK cnt=%d\n", cnt);
    exit(0);
}
```

```
/* Thread routine */
void *thread(void *vargp)
{
    int i, niters = *((int *)vargp);

    for (i = 0; i < niters; i++)
        cnt++;

    return NULL;
}
```

```
linux> ./badcnt 10000
OK cnt=20000
linux> ./badcnt 10000
BOOM! cnt=13051
linux>
```

cnt should equal 20,000.

What went wrong?

Assembly Code for Counter Loop

C code for counter loop in thread i

```
for (i=0; i < niters; i++)  
    cnt++;
```

Corresponding assembly code

<pre> movl (%rdi), %ecx movl \$0, %edx cmpl %ecx, %edx jge .L13</pre>	}	Head (H_i)
<pre>----- .L11: movl cnt(%rip), %eax incl %eax movl %eax, cnt(%rip)</pre>		
<pre>----- incl %edx cmpl %ecx, %edx jl .L11</pre>	}	Load cnt (L_i) Update cnt (U_i) Store cnt (S_i)
<pre>----- .L13:</pre>		
	}	Tail (T_i)

Concurrent Execution

- Key idea: In general, any sequentially consistent interleaving is possible, but some give an unexpected result
 - I_i denotes that thread i executes instruction I
 - $\%eax_i$ is the content of $\%eax$ in thread i 's context

i (thread)	$instr_i$	$\%eax_1$	$\%eax_2$	cnt
1	H_1	-	-	0
1	L_1	0	-	0
1	U_1	1	-	0
1	S_1	1	-	1
2	H_2	-	-	1
2	L_2	-	1	1
2	U_2	-	2	1
2	S_2	-	2	2
2	T_2	-	2	2
1	T_1	1	-	2



Thread 1
critical section



Thread 2
critical section

OK

Concurrent Execution (cont)

- Incorrect ordering: two threads increment the counter, but the result is 1 instead of 2

i (thread)	instr _i	%eax ₁	%eax ₂	cnt
1	H ₁	-	-	0
1	L ₁	0	-	0
1	U ₁	1	-	0
2	H ₂	-	-	0
2	L ₂	-	0	0
1	S ₁	1	-	1
1	T ₁	1	-	1
2	U ₂	-	1	1
2	S ₂	-	1	1
2	T ₂	-	1	1

Oops!

Concurrent Execution (cont)

- How about this ordering?


i (thread)	instr _i	%eax ₁	%eax ₂	cnt
1	H ₁			0
1	L ₁	0		
2	H ₂			
2	L ₂		0	
2	U ₂		1	
2	S ₂		1	1
1	U ₁	1		
1	S ₁	1		1
1	T ₁			
2	T ₂			1

Oops!

- We can analyze the behavior using a progress graph

Progress Graphs

Thread 2



A *progress graph* depicts the discrete *execution state space* of concurrent threads.


Each axis corresponds to the sequential order of instructions in a thread.

Each point corresponds to a possible *execution state* (Inst₁, Inst₂).

E.g., (L₁, S₂) denotes state where thread 1 has completed L₁ and thread 2 has completed S₂.

Trajectories in Progress Graphs

Thread 2




A *trajectory* is a sequence of legal state transitions that describes one possible concurrent execution of the threads.

Example:

$H_1, L_1, U_1, H_2, L_2, S_1, T_1, U_2, S_2, T_2$

Critical Sections and Unsafe Regions

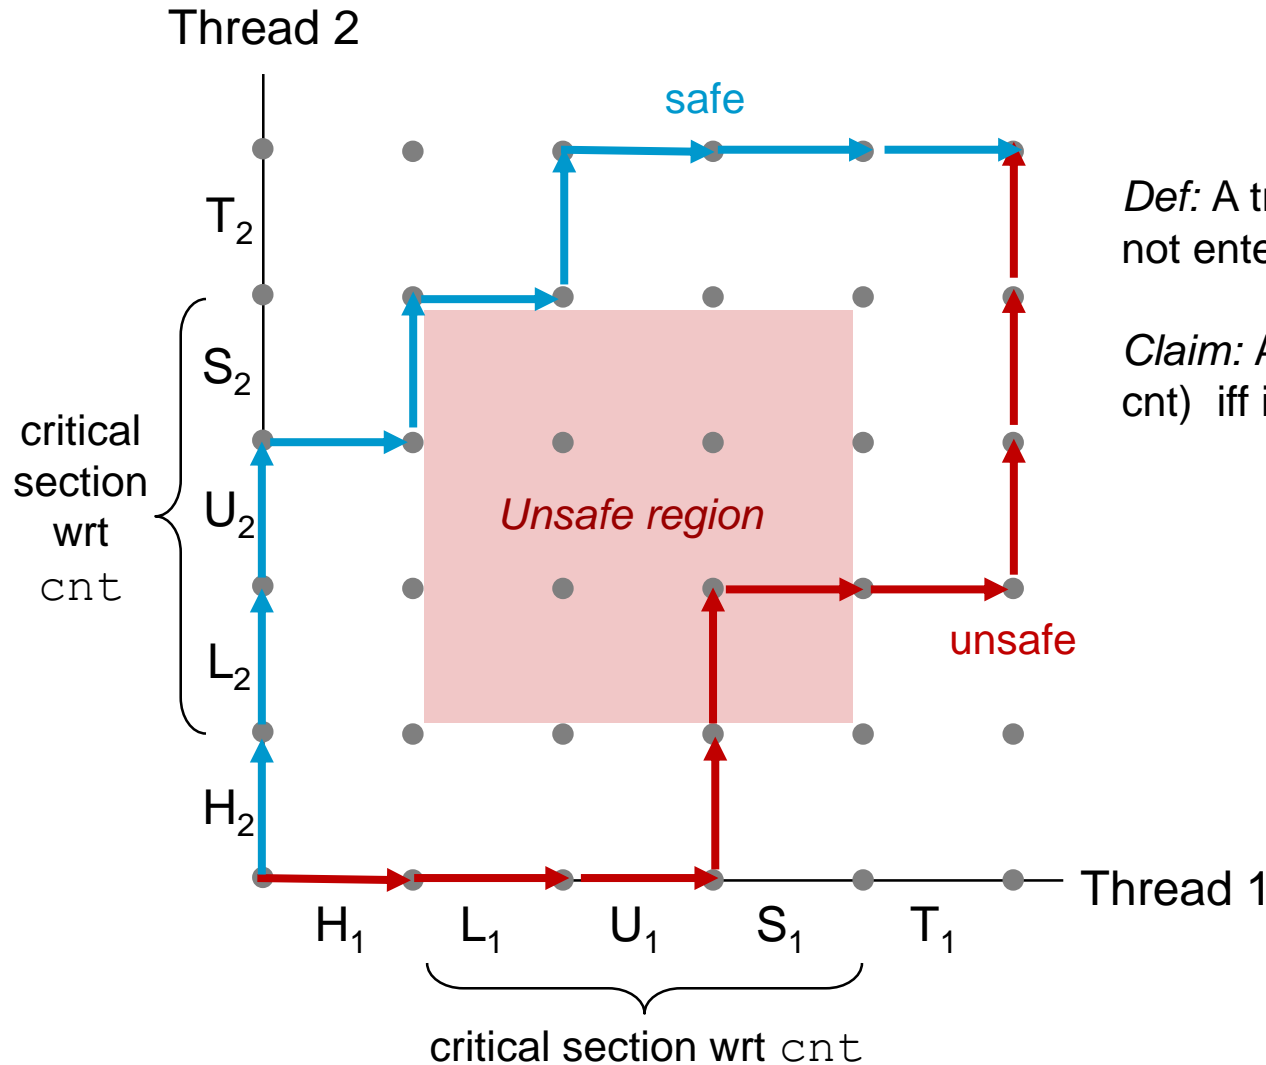


L, U, and S form a *critical section* with respect to the shared variable cnt

Instructions in critical sections (wrt to some shared variable) should not be interleaved

Sets of states where such interleaving occurs form *unsafe regions*

Critical Sections and Unsafe Regions



Def: A trajectory is *safe* iff it does not enter any unsafe region

Claim: A trajectory is correct (wrt cnt) iff it is safe

Enforcing Mutual Exclusion

- Question: How can we guarantee a safe trajectory?
- Answer: We must synchronize the execution of the threads so that they never have an unsafe trajectory.
 - i.e., need to guarantee mutually exclusive access to critical regions
- Classic solution:
 - Semaphores (Edsger Dijkstra)
- Other approaches (out of our scope)
 - Mutex and condition variables (Pthreads)
 - Monitors (Java)

Synchronization in Concurrent Programs

- Threads review
- Sharing
- Mutual exclusion
- **Semaphores**
- Synchronization
 - Producer-consumer problem
 - Readers-writers problem
 - Thread safety
 - Races
 - Deadlocks

Semaphores

- **Semaphore:** non-negative global integer synchronization variable
- Manipulated by P and V operations:
 - $P(s)$: [while (s == 0) wait(); s--;] “test” / “wait”
 - $V(s)$: [s++;] “increment” / “post”
- *OS kernel guarantees that operations between brackets [] are executed indivisibly*
 - ▶ Only one P or V operation at a time can modify s .
 - ▶ When while loop in P terminates, only that P can decrement s
- *Semaphore invariant.* ($s \geq 0$)

C Semaphore Operations

■ POSIX Pthreads semaphore functions

```
#include <semaphore.h>

int sem_init(sem_t *sem, 0, unsigned int val);} /* s = val */

int sem_wait(sem_t *s); /* P(s) */
int sem_post(sem_t *s); /* V(s) */
```

■ CA:APP wrapper functions

```
#include "csapp.h"

void P(sem_t *s); /* Wrapper function for sem_wait */
void V(sem_t *s); /* Wrapper function for sem_post */
```

badcnt.c: Improper Synchronization

```
volatile int cnt = 0; /* global */

int main(int argc, char **argv)
{
    int niters = atoi(argv[1]);
    pthread_t tid1, tid2;

    Pthread_create(&tid1, NULL,
                  thread, &niters);
    Pthread_create(&tid2, NULL,
                  thread, &niters);

    Pthread_join(tid1, NULL);
    Pthread_join(tid2, NULL);

    /* Check result */
    if (cnt != (2 * niters))
        printf("BOOM! cnt=%d\n", cnt);
    else
        printf("OK cnt=%d\n", cnt);
    exit(0);
}
```

```
/* Thread routine */
void *thread(void *vargp)
{
    int i, niters = *((int *)vargp);

    for (i = 0; i < niters; i++)
        cnt++;

    return NULL;
}
```

How can we fix this using
semaphores?

Using Semaphores for Mutual Exclusion

■ Basic idea:

- Associate a unique semaphore *mutex*, initially 1, with each shared variable (or related set of shared variables).
- Surround corresponding critical sections with $P(mutex)$ and $V(mutex)$ operations.

■ Terminology:

- *Binary semaphore*: semaphore whose value is always 0 or 1
- *Mutex*: binary semaphore used for mutual exclusion
 - ▶ P operation: “locking” the mutex
 - ▶ V operation: “unlocking” or “releasing” the mutex
 - ▶ “Holding” a mutex: locked and not yet unlocked.
- *Counting semaphore*: used as a counter for set of available resources.

goodcnt.c: Proper Synchronization

- Define and initialize a mutex for the shared variable cnt:

```
volatile int cnt = 0;      /* Counter */
sem_t mutex;              /* Semaphore that protects cnt */

Sem_init(&mutex, 0, 1);    /* mutex = 1 */
```

- Surround critical section with *P* and *V*:

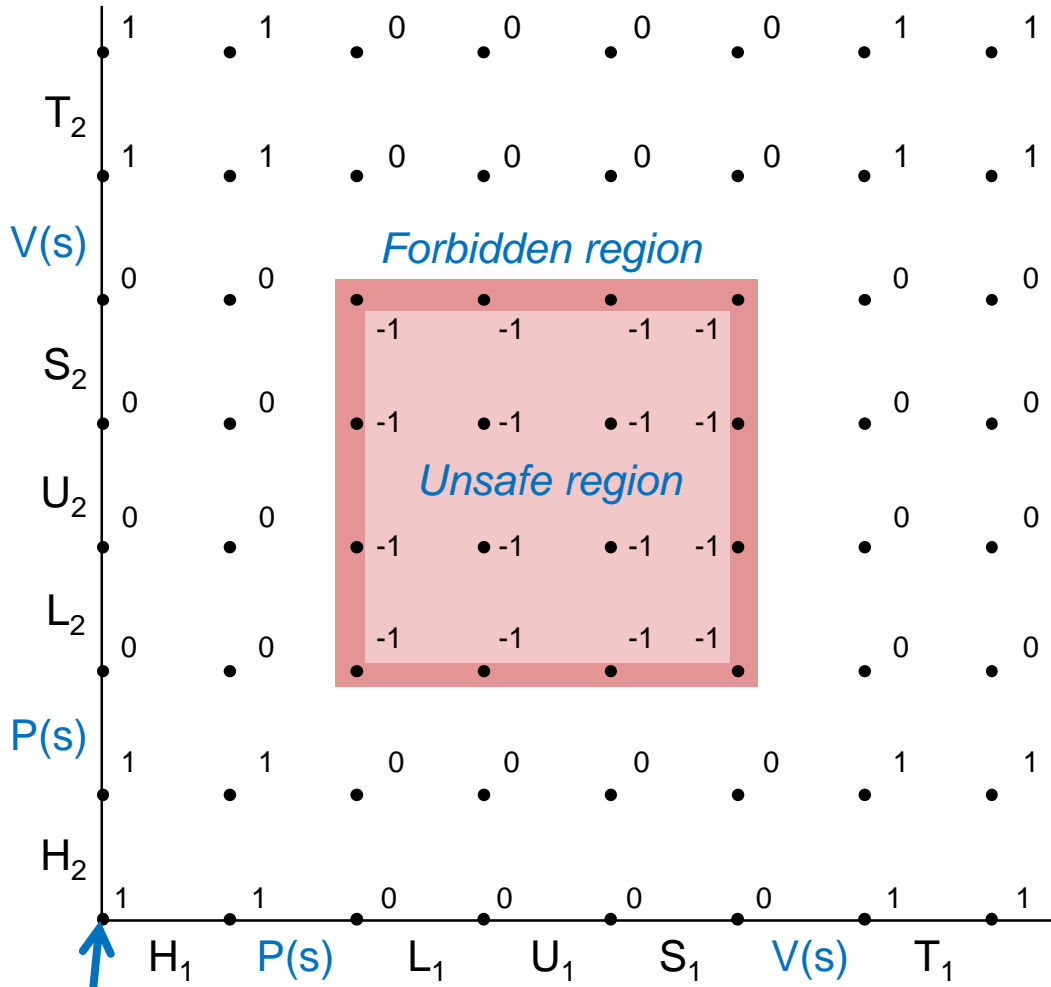
```
for (i = 0; i < niters; i++) {
    P(&mutex);
    cnt++;
    V(&mutex);
}
```

```
linux> ./goodcnt 10000
OK cnt=20000
linux> ./goodcnt 10000
OK cnt=20000
linux>
```

Observation: much slower than
badcnt.c.

Why Mutexes Work

Thread 2



Provide mutually exclusive access to shared variable by surrounding critical section with P and V operations on semaphore s (initially set to 1)

Semaphore invariant creates a *forbidden region* that encloses unsafe region that cannot be entered by any trajectory.

Thread 1

Initially $s = 1$


Synchronization in Concurrent Programs

- Threads review
- Sharing
- Mutual exclusion
- Semaphores
- **Synchronization**
 - **Producer-consumer problem**
 - Readers-writers problem
 - Thread safety
 - Races
 - Deadlocks

Using Semaphores to Schedule Access to Shared Resources

- Basic idea: Thread uses a semaphore operation to notify another thread that some condition has become true
 - Use counting semaphores to keep track of resource state.
 - Use binary semaphores to notify other threads.
- Two classic examples:
 - The Producer-Consumer Problem
 - The Readers-Writers Problem

Producer-Consumer Problem



■ Common synchronization pattern:

- Producer waits for empty slot, inserts item in buffer, and notifies consumer
- Consumer waits for item, removes it from buffer, and notifies producer

■ Examples

- Multimedia processing:
 - ▶ Producer creates MPEG video frames, consumer renders them
- Event-driven graphical user interfaces
 - ▶ Producer detects mouse clicks, mouse movements, and keyboard hits and inserts corresponding events in buffer
 - ▶ Consumer retrieves events from buffer and paints the display

Producer-Consumer on 1-element Buffer

```
#include "csapp.h"

#define NITERS 5

void *producer(void *arg);
void *consumer(void *arg);

struct {
    int buf; /* shared var */
    sem_t full; /* sems */
    sem_t empty;
} shared;

...
```

```
...

int main() {
    pthread_t tid_producer;
    pthread_t tid_consumer;

    /* Initialize the semaphores */
    Sem_init(&shared.empty, 0, 1);
    Sem_init(&shared.full, 0, 0);

    /* Create threads and wait */
    Pthread_create(&tid_producer, NULL,
                  producer, NULL);
    Pthread_create(&tid_consumer, NULL,
                  consumer, NULL);

    Pthread_join(tid_producer, NULL);
    Pthread_join(tid_consumer, NULL);

    exit(0);
}
```

Producer-Consumer on 1-element Buffer

- Initially: empty==1, full==0

Producer Thread

```
void *producer(void *arg) {
    int i, item;

    for (i=0; i<NITERS; i++) {
        /* Produce item */
        item = i;
        printf("produced %d\n",
            item);

        /* Write item to buf */
        P(&shared.empty);
        shared.buf = item;
        V(&shared.full);
    }
    return NULL;
}
```

Consumer Thread

```
void *consumer(void *arg) {
    int i, item;

    for (i=0; i<NITERS; i++) {
        /* Read item from buf */
        P(&shared.full);
        item = shared.buf;
        V(&shared.empty);

        /* Consume item */
        printf("consumed %d\n", item);
    }
    return NULL;
}
```

Producer-Consumer on an n-element Buffer

- Requires a mutex and two counting semaphores:
 - mutex: enforces mutually exclusive access to the the buffer
 - slots: counts the available slots in the buffer
 - items: counts the available items in the buffer
- Implemented using a shared buffer package called `sbuf`

sbuf Package - Declarations

■ Data structure and interface

```
#include "csapp.h"

typedef struct {
    int *buf;           /* Buffer array */
    int n;              /* Maximum number of slots */
    int front;          /* buf[(front+1)%n] is first item */
    int rear;           /* buf[rear%n] is last item */
    sem_t mutex;        /* Protects accesses to buf */
    sem_t slots;        /* Counts available slots */
    sem_t items;        /* Counts available items */
} sbuf_t;

void sbuf_init(sbuf_t *sp, int n);
void sbuf_deinit(sbuf_t *sp);
void sbuf_insert(sbuf_t *sp, int item);
int sbuf_remove(sbuf_t *sp);
```

sbuf.h

sbuf Package - Implementation

■ Initializing and deinitializing a shared buffer

```
/* Create an empty, bounded, shared FIFO buffer with n slots */
void sbuf_init(sbuf_t *sp, int n)
{
    sp->buf = Calloc(n, sizeof(int));
    sp->n = n; /* Buffer holds max of n items */
    sp->front = sp->rear = 0; /* Empty buffer iff front == rear */
    Sem_init(&sp->mux, 0, 1); /* Binary semaphore for locking */
    Sem_init(&sp->slots, 0, n); /* Initially, buf has n empty slots */
    Sem_init(&sp->items, 0, 0); /* Initially, buf has zero items */
}

/* Clean up buffer sp */
void sbuf_deinit(sbuf_t *sp)
{
    Free(sp->buf);
}
```

sbuf.c

sbuf Package - Implementation

■ Inserting an item into a shared buffer

```
/* Insert item onto the rear of shared buffer sp */
void sbuf_insert(sbuf_t *sp, int item)
{
    P(&sp->slots);                /* Wait for available slot */
    P(&sp->mutex);                 /* Lock the buffer */
    sp->buf[(++sp->rear)%(sp->n)] = item; /* Insert the item */
    V(&sp->mutex);                 /* Unlock the buffer */
    V(&sp->items);                 /* Announce available item */
}
```

sbuf.c

sbuf Package - Implementation

■ Removing an item from a shared buffer

```
/* Remove and return the first item from buffer sp */
int sbuf_remove(sbuf_t *sp)
{
    int item;

    P(&sp->items);                /* Wait for available item */
    P(&sp->mutex);                 /* Lock the buffer */
    item = sp->buf[(++sp->front)%(sp->n)]; /* Remove the item */
    V(&sp->mutex);                 /* Unlock the buffer */
    V(&sp->slots);                 /* Announce available slot */
    return item;
}
```

sbuf.c

Synchronization in Concurrent Programs

- Threads review
- Sharing
- Mutual exclusion
- Semaphores
- **Synchronization**
 - Producer-consumer problem
 - **Readers-writers problem**
 - Thread safety
 - Races
 - Deadlocks

Readers-Writers Problem

- Generalization of the mutual exclusion problem
- Problem statement:
 - *Reader* threads only read the object
 - *Writer* threads modify the object
 - Writers must have exclusive access to the object
 - Unlimited number of readers can access the object
- Occurs frequently in real systems, e.g.,
 - Online airline reservation system
 - Multithreaded caching Web proxy

Variants of Readers-Writers

- First readers-writers problem (favors readers)
 - No reader should be kept waiting unless a writer has already been granted permission to use the object.
 - A reader that arrives after a waiting writer gets priority over the writer.
- Second readers-writers problem (favors writers)
 - Once a writer is ready to write, it performs its write as soon as possible
 - A reader that arrives after a writer must wait, even if the writer is also waiting.
- Third readers-writers problem (equal priority to both)
 - Under the assumption of a FIFO wake-up sequence for semaphores
- *Starvation* (where a thread waits indefinitely) is possible in both cases.

Solution to First Readers-Writers Problem

■ Reader code

```
int readcnt;      /* Initially 0 */
sem_t mutex, w; /* Both initially 1 */

void reader(void)
{
    while (1) {
        P(&mutex);
        readcnt++;
        if (readcnt == 1) /* First in */
            P(&w);
        V(&mutex);

        /* Reading happens here */

        P(&mutex);
        readcnt--;
        if (readcnt == 0) /* Last out */
            V(&w);
        V(&mutex);
    }
}
```

■ Writer code

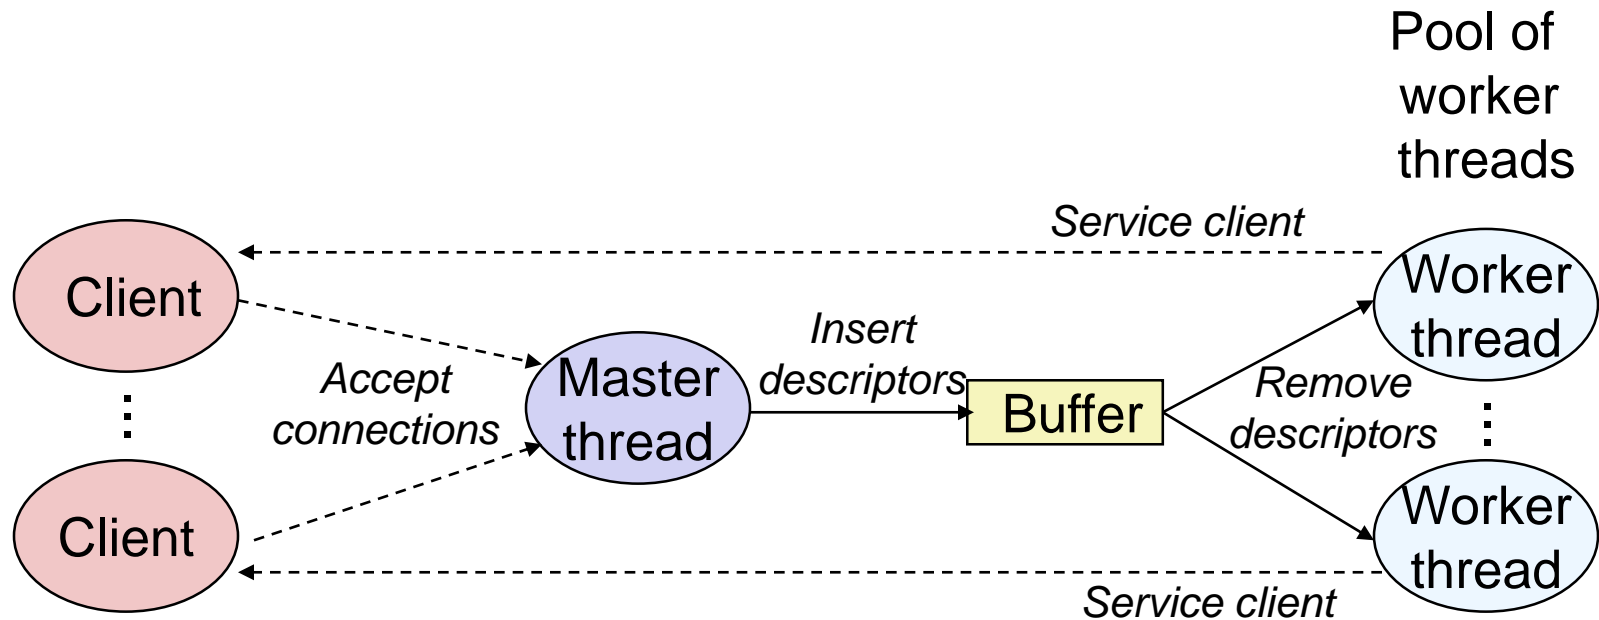
```
void writer(void)
{
    while (1) {
        P(&w);

        /* Writing here */

        V(&w);
    }
}
```

rw1.c

Case Study: Prethreaded Concurrent Server



Prethreaded Concurrent Server

```
sbuf_t sbuf; /* Shared buffer of connected descriptors */

int main(int argc, char **argv)
{
    int i, listenfd, connfd, port;
    socklen_t clientlen=sizeof(struct sockaddr_in);
    struct sockaddr_in clientaddr;
    pthread_t tid;

    port = atoi(argv[1]);
    sbuf_init(&sbuf, SBUFSIZE);
    listenfd = Open_listenfd(port);

    for (i = 0; i < NTHREADS; i++) /* Create worker threads */
        Pthread_create(&tid, NULL, thread, NULL);

    while (1) {
        connfd = Accept(listenfd, (SA *) &clientaddr, &clientlen);
        sbuf_insert(&sbuf, connfd); /* Insert connfd in buffer */
    }
}
```

echoserver_pre.c

Prethreaded Concurrent Server

Worker thread routine:

```
void *thread(void *vargp)
{
    Pthread_detach(pthread_self());
    while (1) {
        int connfd = sbuf_remove(&sbuf); /* Remove connfd from
                                          buffer */
        echo_cnt(connfd);                /* Service client */
        Close(connfd);
    }
}
```

echoserver_pre.c

Prethreaded Concurrent Server

echo_cnt initialization routine:

```
static int byte_cnt; /* Byte counter */
static sem_t mutex; /* and the mutex that protects it */

static void init_echo_cnt(void)
{
    Sem_init(&mutex, 0, 1);
    byte_cnt = 0;
}
```

echo_cnt.c

Prethreaded Concurrent Server

Worker thread service routine:

```
void echo_cnt(int connfd)
{
    int n;
    char buf[MAXLINE];
    rio_t rio;
    static pthread_once_t once = PTHREAD_ONCE_INIT;

    Pthread_once(&once, init_echo_cnt);
    Rio_readinitb(&rio, connfd);
    while((n = Rio_readlineb(&rio, buf, MAXLINE)) != 0) {
        P(&mutex);
        byte_cnt += n;
        printf("thread %d received %d (%d total) bytes"
               " on fd %d\n",
               (int) pthread_self(), n, byte_cnt, connfd);
        V(&mutex);
        Rio_writen(connfd, buf, n);
    }
}
```

echo_cnt.c

Synchronization in Concurrent Programs

- Threads review
- Sharing
- Mutual exclusion
- Semaphores
- **Synchronization**
 - Readers-writers problem
 - Producer-consumer problem
 - **Thread safety**
 - Races
 - Deadlocks

Crucial concept: Thread Safety

- Functions called from a thread must be **thread-safe**
- *Def.* A function is *thread-safe* iff it will always produce correct results when called repeatedly from multiple concurrent threads.
- Classes of thread-unsafe functions:
 - Class 1: Functions that do not protect shared variables.
 - Class 2: Functions that keep state across multiple invocations.
 - Class 3: Functions that return a pointer to a static variable.
 - Class 4: Functions that call thread-unsafe functions.

Thread-Unsafe Functions (Class 1)

- Failing to protect shared variables
 - Fix: Use semaphore operations to protect parts that manipulate shared variables
 - Example: `goodcnt.c`
 - Issue: Synchronization operations will slow down code

Thread-Unsafe Functions (Class 2)

- Relying on persistent state across multiple function invocations
 - Example: Random number generator that relies on static state

```
static unsigned int next = 1;

/* rand: return pseudo-random integer on 0..32767 */
int rand(void)
{
    next = next*1103515245 + 12345;
    return (unsigned int)(next/65536) % 32768;
}

/* srand: set seed for rand() */
void srand(unsigned int seed)
{
    next = seed;
}
```

Thread-Safe Random Number Generator

- Fix: Pass state as part of argument
 - and, thereby, eliminate static state

```
/* rand_r - return pseudo-random integer on 0..32767 */  
  
int rand_r(int *nextp)  
{  
    *nextp = *nextp*1103515245 + 12345;  
    return (unsigned int)(*nextp/65536) % 32768;  
}
```

- Consequence: programmer using rand_r must maintain seed

Thread-Unsafe Functions (Class 3)

- Returning a pointer to a static variable
- Fix 1: Rewrite function so caller passes address of variable to store result
 - Requires changes in caller and callee
- Fix 2: Lock-and-copy
 - Requires simple changes in caller (and none in callee)
 - However, caller must free memory.

```
/* lock-and-copy version */
char *ctime_ts(const time_t *timep,
               char *privatep)
{
    char *sharedp;

    P(&mutex);
    sharedp = ctime(timep);
    strcpy(privatep, sharedp);
    V(&mutex);
    return privatep;
}
```

Warning: Some functions like `gethostbyname` require a *deep copy*. Use reentrant `gethostbyname_r` version instead.

Thread-Unsafe Functions (Class 4)

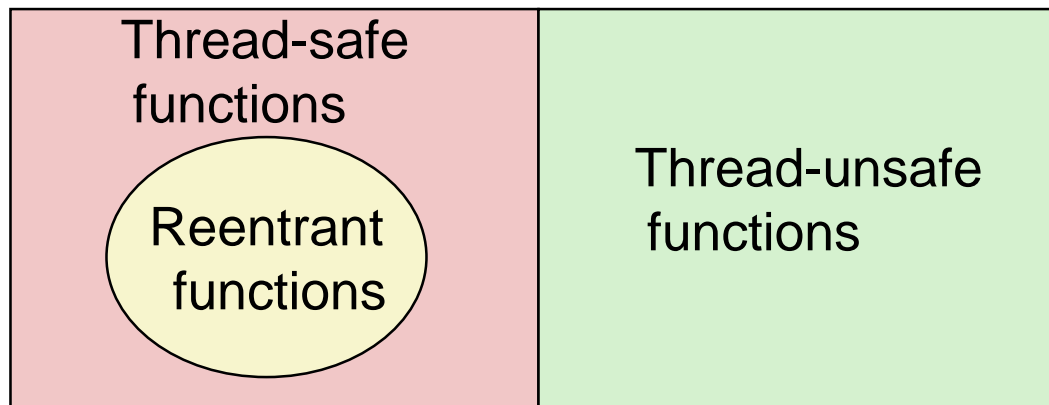
■ Calling thread-unsafe functions

- Calling one thread-unsafe function makes the entire function that calls it thread-unsafe
- Fix 1: Modify the function so it calls only thread-safe functions
- Fix 2: Protect the call site and resulting shared data with a mutex iff the callee is of class 1 or 3.

Reentrant Functions

- Def: A function is **reentrant** iff it accesses no shared variables when called by multiple threads.
 - Important subset of thread-safe functions.
 - ▶ Require no synchronization operations.
 - ▶ Only way to make a Class 2 function thread-safe is to make it reentrant (e.g., `rand_r`)

All functions



Thread-Safe Library Functions

- All functions in the Standard C Library (at the back of your K&R text) are thread-safe
 - Examples: `malloc`, `free`, `printf`, `scanf`
- Most Unix system calls are thread-safe, with a few important exceptions:

Thread-unsafe function	Class	Reentrant version
<code>asctime</code>	3	<code>asctime_r</code>
<code>ctime</code>	3	<code>ctime_r</code>
<code>gethostbyaddr</code>	3	<code>gethostbyaddr_r</code>
<code>gethostbyname</code>	3	<code>gethostbyname_r</code>
<code>inet_ntoa</code>	3	(none)
<code>localtime</code>	3	<code>localtime_r</code>
<code>rand</code>	2	<code>rand_r</code>

Synchronization in Concurrent Programs

- Threads review
- Sharing
- Mutual exclusion
- Semaphores
- **Synchronization**
 - Readers-writers problem
 - Producer-consumer problem
 - Thread safety
 - **Races**
 - Deadlocks

One Worry: Races

- A **race** occurs when correctness of the program depends on one thread reaching point x before another thread reaches point y

```
/* a threaded program with a race */
int main() {
    pthread_t tid[N];
    int i;
    for (i = 0; i < N; i++)
        Pthread_create(&tid[i], NULL, thread, &i);
    for (i = 0; i < N; i++)
        Pthread_join(tid[i], NULL);
    exit(0);
}

/* thread routine */
void *thread(void *vargp) {
    int myid = *((int *)vargp);
    printf("Hello from thread %d\n", myid);
    return NULL;
}
```


Race Elimination

- Make sure don't have unintended sharing of state

```
/* a threaded program without the race */
int main() {
    pthread_t tid[N];
    int i;
    for (i = 0; i < N; i++) {
        int *valp = malloc(sizeof(int));
        *valp = i;
        Pthread_create(&tid[i], NULL, thread, valp);
    }
    for (i = 0; i < N; i++)
        Pthread_join(tid[i], NULL);
    exit(0);
}

/* thread routine */
void *thread(void *vargp) {
    int myid = *((int *)vargp);
    free(vargp);
    printf("Hello from thread %d\n", myid);
    return NULL;
}
```

Synchronization in Concurrent Programs

- Threads review
- Sharing
- Mutual exclusion
- Semaphores
- **Synchronization**
 - Readers-writers problem
 - Producer-consumer problem
 - Thread safety
 - Races
 - **Deadlocks**

Another Worry: Deadlock

- Def: A process is **deadlocked** iff it is waiting for a condition that will never be true.
- Typical Scenario
 - Processes 1 and 2 needs two resources (A and B) to proceed
 - Process 1 acquires A, waits for B
 - Process 2 acquires B, waits for A
 - Both will wait forever!

Deadlocking With Semaphores

```
int main()
{
    pthread_t tid[2];
    Sem_init(&mutex[0], 0, 1); /* mutex[0] = 1 */
    Sem_init(&mutex[1], 0, 1); /* mutex[1] = 1 */
    Pthread_create(&tid[0], NULL, count, (void*) 0);
    Pthread_create(&tid[1], NULL, count, (void*) 1);
    Pthread_join(tid[0], NULL);
    Pthread_join(tid[1], NULL);
    printf("cnt=%d\n", cnt);
    exit(0);
}
```

```
void *count(void *vargp)
{
    int i;
    int id = (int) vargp;
    for (i = 0; i < NITERS; i++) {
        P(&mutex[id]); P(&mutex[1-id]);
        cnt++;
        V(&mutex[id]); V(&mutex[1-id]);
    }
    return NULL;
}
```

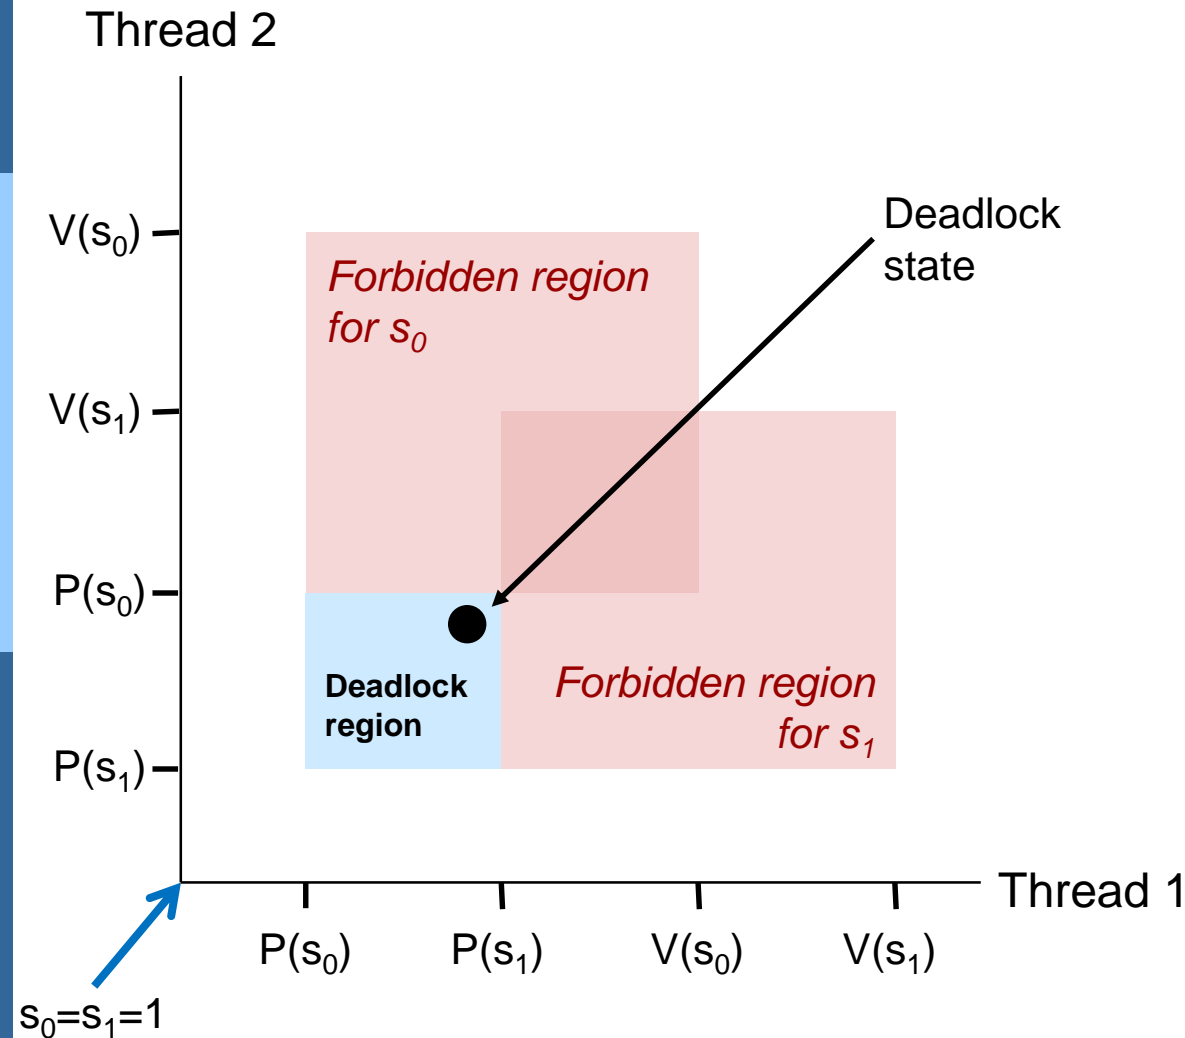
Tid[0]:

P(s₀);
P(s₁);
cnt++;
V(s₀);
V(s₁);

Tid[1]:

P(s₁);
P(s₀);
cnt++;
V(s₁);
V(s₀);

Deadlock Visualized in Progress Graph



Locking introduces the potential for *deadlock*: waiting for a condition that will never be true

Any trajectory that enters the *deadlock region* will eventually reach the *deadlock state*, waiting for either s_0 or s_1 to become nonzero

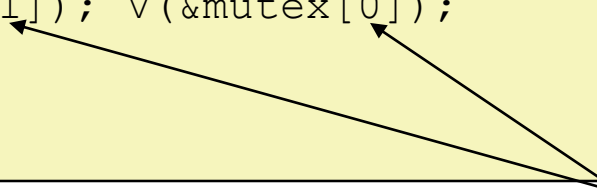
Other trajectories luck out and skirt the deadlock region

Unfortunate fact: deadlock is often nondeterministic

Avoiding Deadlocks: Acquire shared resources *in the same order*

```
int main()
{
    pthread_t tid[2];
    Sem_init(&mutex[0], 0, 1); /* mutex[0] = 1 */
    Sem_init(&mutex[1], 0, 1); /* mutex[1] = 1 */
    Pthread_create(&tid[0], NULL, count, (void*) 0);
    Pthread_create(&tid[1], NULL, count, (void*) 1);
    Pthread_join(tid[0], NULL);
    Pthread_join(tid[1], NULL);
    printf("cnt=%d\n", cnt);
    exit(0);
}
```

```
void *count(void *vargp)
{
    int i;
    int id = (int) vargp;
    for (i = 0; i < NITERS; i++) {
        P(&mutex[0]); P(&mutex[1]);
        cnt++;
        V(&mutex[1]); V(&mutex[0]);
    }
    return NULL;
}
```



Tid[0]:

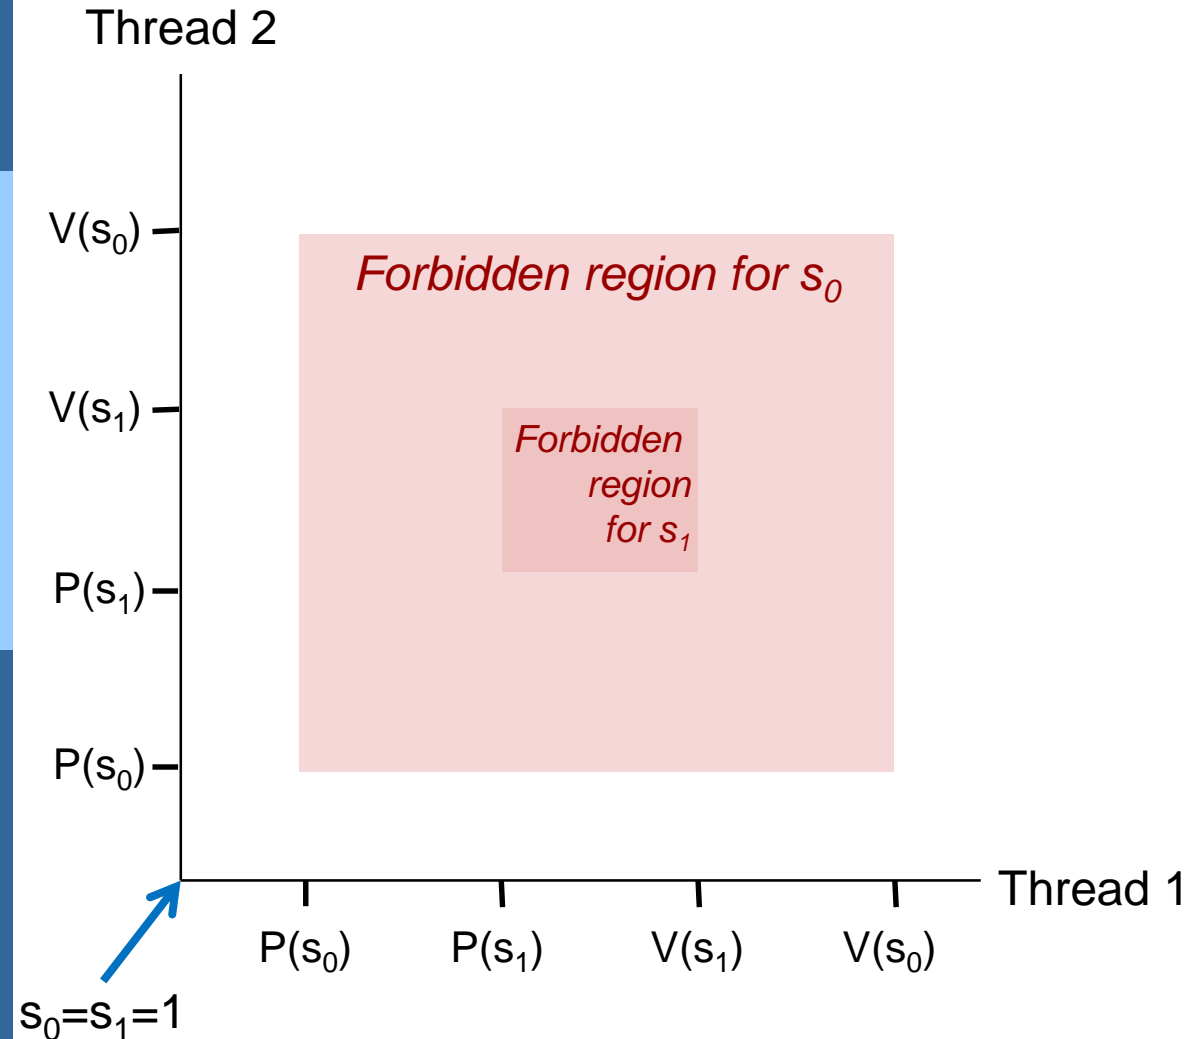
P(s₀);
P(s₁);
cnt++;
V(s₁);
V(s₀);

Tid[1]:

P(s₀);
P(s₁);
cnt++;
V(s₁);
V(s₀);

what about the
release order?

Avoiding Deadlock in Progress Graph



No way for trajectory to get stuck

Processes acquire locks in same order

Order in which locks released?

Synchronization Summary

- Threads provide another mechanism for writing concurrent programs
- Threads are popular
 - Somewhat cheaper than processes
 - Easy to share data between threads
- However, the ease of sharing has a cost
 - Easy to introduce subtle synchronization errors
- Programmers need a clear model of how variables are shared by threads.
- Variables shared by multiple threads must be protected to ensure mutually exclusive access.
- Semaphores are a fundamental mechanism for enforcing mutual exclusion.
- For more info
 - D. Butenhof, “Programming with Posix Threads”, Addison-Wesley, 1997