

---

## Homework #2

## M1522.000800 System Programming

---

Name: \_\_\_\_\_

**Due Date:** Tuesday, March 17, 2015, 23:59

**Student-Number:** \_\_\_\_\_

**Submission:** in paper form.  
There will be a drop off box in class and inside the CSAP Lab in building 301, room 419.

### Question 1

#### *Integer Arithmetic*

Consider the following code that attempts to sum the elements of an array *a*, where the number of elements is given by parameter *length*:

```
/* WARNING: This is buggy code */
float sum_elements (float a[], unsigned length)
{
    int i;
    float result = 0;

    for (i = 0; i <= length-1; i++)
        result += a[i];
    return result;
}
```

When run with argument *length* equal to 0, this code should return 0.0. Instead it encounters a memory error. Explain why this happens. Show how this code can be corrected.

## Question 2

### SYMBOLS & SYMBOL TABLES

Consider the following version of the swap.c function that counts the number of times it has been called, and main.c.

```
/* main.c */
void swap();

int buf[2] = {1, 2};

int main()
{
    swap();
    return 0;
}

/* swap.c */
extern int buf[];

int *bufp0 = &buf[0];
static int *bufp1;

static void incr()
{
    static int count = 0;
    count++;
}

void swap()
{
    int temp;

    incr();
    bufp1 = &buf[1];
    temp = *bufp0;
    *bufp0 = *bufp1;
    *bufp1 = temp;
}
```

For each symbol that is defined and referenced in swap.o, indicate if it will have a symbol table entry in the .symtab section in module swap.o. If so, indicate the module that defines the symbol (swap.o or main.o), the symbol type (local, global, or extern) and the section (.text, .data, or .bss) it occupies in that module.

Symbol	Swap.o's .symtab entry?	Symbol type	Module where define	Section
buf	Yes	extern	main.o	.data
bufp0				
bufp1				
swap				
temp	No	-	-	-
incr				
count				

### Question 3

#### *SYMBOL RESOLUTION*

Consider the following example, where *x* is defined as an *int* in one module and a *double* in another:

```
/* foo.c */
#include <stdio.h>
void f(void);

int x = 0x01234567;
int y = 0x89ABCDEF;

int main()
{
    f();
    printf("%x %x\n", x, y);
    return 0;
}

/* bar.c */
long long x;

void f()
{
    x = 0xDEADBEEFBABEFACE;
}
```

After compiling and linking with the following commands using GCC compiler driver, we can found that variable *x*, *y* is overwritten by *f()*.

```
linux > gcc -o foobar foo.c bar.c
( Warning ... )
linux > ./foobar
BABEFACE DEADBEEF
```

Explain why this happens. Show how this code can print “01234567 89ABCDEF” instead of that.

### Question 4

#### *LINKING WITH STATIC LIBRARIES*

Consider the two different approaches to use C standard libraries. One approach [A1] is to have the compiler recognize calls to the standard functions and to generate the appropriate code directly. Another [A2] is to have the compiler link archive files where standard functions are defined.

```
[A1] unix > gcc main.c /usr/lib/libc.o
[A2] unix > gcc main.c /usr/lib/libc.a
```

Explain which approach is better in common case and why.

## Question 5

### STATIC LIBRARIES & REFERENCE RESOLUTION

During the symbol resolution phase in linking with static libraries, the linker scans the relocatable object files and archives left to right in the same sequential order that they appear on the compiler driver's command line. (The driver automatically translates any .c files on the command line into .o files.) During this scan, the linker maintains a set  $E$  of relocatable object files that will be merged to form the executable, a set  $U$  of unresolved symbols (i.e., symbols referred to but not yet defined), and a set  $D$  of symbols that have been defined in previous input files. Initially,  $E$ ,  $U$ , and  $D$  are empty.

- For each input file  $f$  on the command line, the linker determines if  $f$  is an object file or an archive. If  $f$  is an object file, the linker adds  $f$  to  $E$ , updates  $U$  and  $D$  to reflect the symbol definitions and references in  $f$ , and proceeds to the next input file.
- If  $f$  is an archive, the linker attempts to match the unresolved symbols in  $U$  against the symbols defined by the members of the archive. If some archive member,  $m$ , defines a symbol that resolves a reference in  $U$ , then  $m$  is added to  $E$ , and the linker updates  $U$  and  $D$  to reflect the symbol definitions and references in  $m$ . This process iterates over the member object files in the archive until a fixed point is reached where  $U$  and  $D$  no longer change. At this point, any member object files not contained in  $E$  are simply discarded and the linker proceeds to the next input file.
- If  $U$  is nonempty when the linker finishes scanning the input files on the command line, it prints an error and terminates. Otherwise it merges and relocates the object files in  $E$  to build the output executable file.

Assume that './libfoo.a' has a definition of function 'foo()' which called in './main.c'. If we compile and link with the following command, we encounter the linking error.

```
unix > gcc -static ./libfoo.a main.c
/tmp/cc9XH6Rp.o: In function 'main':
/tmp/cc9XH6Rp.o(.text+0x18): undefined reference to 'foo'
```

Explain why this happens and how it could be fixed. (Consider linker symbol resolution scan rule.)